

The Portsmouth Papers: A Hands-On Tour of PostgreSQL Beyond the Relational Model

The Portsmouth Papers

A Hands-On Tour of PostgreSQL Beyond the Relational Model

About This Book

Author

Chapter 1 — JSONB: Semi-Structured Data Without a Schema Tax

Background

The Scenario

Exercise Goals

Installation

1 — PostgreSQL

2 — Python 3.12 and psycopg

Loading the Data

Create the database

Run the seed script

Verify the load

Exercises

Exercise 1 — The Three Extraction Operators

Exercise 2 — Filtering with Containment and Key Existence

Exercise 3 — Missing Keys vs. NULL Values

Exercise 4 — GIN Indexes

Exercise 5 — Updating Documents in Place

Exercise 6 — Expanding Arrays with `jsonb_array_elements`

Exercise 7 — Path Queries with `jsonb_path_query`

Summary — What You Should Now Know

Chapter 2 — PostGIS: Geospatial Queries on Real Geometry

Background

The Scenario

Exercise Goals

Installation

1 — PostGIS server package

2 — Enable PostGIS in the database

Loading the Data

Prerequisites

Run the Chapter 2 seed

Verify the load

Exercises

Exercise 1 — Neighbourhood Polygons and WKT

Exercise 2 — Proximity Search with `ST_DWithin`

Exercise 3 — Spatial Joins with `ST_Within` and `ST_Contains`

Exercise 4 — Computing Area in Square Kilometres

Exercise 5 — Nearest Park Using `ST_Distance` and a Lateral Join

Exercise 6 — GIST Indexes and Spatial Query Plans

Summary — What You Should Now Know

Chapter 3 — Job Queues: `FOR UPDATE SKIP LOCKED`

Background

The Scenario

Exercise Goals

Installation

Loading the Data

Run the seed script

Verify the load

Exercises

Exercise 1 — Designing the Queue Schema

Exercise 2 — The Atomic Claim Query

Exercise 3 — Simulating Concurrent Workers

Exercise 4 — Heartbeat and Stalled-Job Recovery

Exercise 5 — Dead-Lettering Exhausted Jobs

Exercise 6 — Benchmarking Claim Throughput with `pgbench`

Summary — What You Should Now Know

Chapter 4 — Full-Text Search: `tsvector`, Stopwords, and Ranking

Background

The Scenario

Exercise Goals

Installation

Loading the Data

Run the seed script

Verify the load

Exercises

Exercise 1 — Converting Text to `tsvector`

Exercise 2 — What Stopword Removal Actually Throws Away

Exercise 3 — A Maintained Column and a GIN Index

Exercise 4 — `to_tsquery` vs. `plainto_tsquery`

Exercise 5 — Ranking with `ts_rank` and `ts_rank_cd`, and `ts_headline` Snippets

Exercise 6 — A Custom Text Search Configuration for Domain Stopwords

Summary — What You Should Now Know

Chapter 5 — Fuzzy Matching: `pg_trgm`

Background

The Scenario

Exercise Goals

Installation

Loading the Data

Prerequisites

Run the Chapter 5 seed

Verify the load

Exercises

Exercise 1 — What a Trigram Is, and `similarity()`

Exercise 2 — The `%` Operator, and Why One Threshold Isn't Enough

Exercise 3 — `word_similarity()` for Partial Matches

Exercise 4 — GIN Trigram Indexes

Exercise 5 — A “Did You Mean?” Query with GiST

Exercise 6 — Trigram Matching vs. Full-Text Search, Head to Head

Summary — What You Should Now Know

Chapter 6 — Vector Search: `pgvector` for Embeddings

Background

The Scenario

Exercise Goals

Installation

1 — `pgvector`

2 — `sentence-transformers`, CPU-only

Loading the Data

Prerequisites

Run the Chapter 6 seed

Compute real embeddings for `city_documents`

Verify the load

Exercises

Exercise 1 — The vector Type

Exercise 2 — Exact Nearest Neighbors: `<->`, `<#>`, `<=>`

Exercise 3 — IVFFlat: Approximate Search with `lists` and `probes`

Exercise 4 — HNSW: A Different Trade-off

Exercise 5 — Semantic Search: Finding Meaning, Not Words

Exercise 6 — Hybrid Search: Blending Keyword and Semantic Signals

Summary — What You Should Now Know

Bonus Section — Build a Local RAG System

- What RAG actually is

- Installation — Ollama

- Ingest: chunk and embed a directory of documents

- Ask a question

- What happens when the answer isn't in the documents

- The prompt: the actual glue between the two systems

Chapter 7 — IP and Network Filtering: ip4r

- Background

- The Scenario

- Exercise Goals

- Installation

- Loading the Data

 - Run the seed script

 - Verify the load

- Exercises

 - Exercise 1 — ip4r vs. the Built-in inet/cidr

 - Exercise 2 — Containment: >> and <<=

 - Exercise 3 — GiST Indexing (and a Real Gotcha)

 - Exercise 4 — Aggregating by Subnet

 - Exercise 5 — Overlap Detection Between Allow and Block Lists

 - Exercise 6 — A Real-Time “Is This IP Blocked?” Function

 - Exercise 7 (Additional) — Rate Limiting by Host *and* /24 with a Token Bucket

- Summary — What You Should Now Know

Chapter 8 — Declarative Partitioning and BRIN Indexes

- Background

- The Scenario

- Exercise Goals

- Installation

- Loading the Data

 - Run the seed script

 - Pin the session timezone before you do anything else

 - Verify the load

- Exercises

 - Exercise 1 — Designing and Creating the Partitioned Table

 - Exercise 2 — Backfilling and Watching Automatic Partition

 - Routing

 - Exercise 3 — Partition Pruning, Seen in EXPLAIN

 - Exercise 4 — BRIN vs. B-tree, With Real Numbers


 - Exercise 5 — Dropping a Partition vs. DELETE

 - Exercise 6 — Combining LIST and RANGE: Subpartitioning

- Summary — What You Should Now Know

The Portsmouth Papers

A Hands-On Tour of PostgreSQL Beyond the Relational Model



*“PostgreSQL is not a database with extensions.
It is an extensible data platform that happens to speak SQL.”*

About This Book

Most PostgreSQL tutorials end where the interesting work begins.

Once you know how to `SELECT`, `JOIN`, and `GROUP BY`, you have mastered perhaps twenty percent of what PostgreSQL can do. The remaining eighty percent — geospatial queries, semantic search, real-time notifications, fuzzy matching, vector embeddings, distributed coordination, and more — lives in a ecosystem of extensions, index types, and language features that most practitioners never discover.

This book is a guided tour of that other eighty percent.

Each chapter is built around a concrete engineering problem faced by the fictional city of **Portsmith** and its data platform team. We store business directories as semi-structured JSON documents. We route emergency services using geospatial proximity queries. We build a job queue with no message broker. We search municipal records with fuzzy matching that survives typos and OCR errors. We expose the entire platform as a REST API with zero application code.

Every chapter follows the same structure: synthetic data is generated first, giving you a realistic dataset to work against, and then a series of exercises walks you from first principles to production-ready technique. The exercises are written to be done — not just read.

By the end, you will see PostgreSQL not as a place to store rows, but as a programmable data infrastructure layer capable of doing work that most teams reach for separate specialized systems to handle.

What you will need:

- A working PostgreSQL 16 installation (setup covered in Appendix A)
- Python 3.12+
- A Debian-based Linux environment
- Familiarity with basic SQL (SELECT, INSERT, JOIN, GROUP BY)
- Curiosity about what else is in there

Author

Chris Lee

Edition 1.0 — Portsmith, 2026

Chapter 1 — JSONB: Semi-Structured Data Without a Schema Tax

| *“A schema is a prediction about the future. JSONB lets you hedge.”*

Background

Relational databases enforce a contract: every row in a table has exactly the same columns. That contract is usually a feature — it catches mistakes and makes queries predictable. But sometimes you genuinely don't know what columns you'll need until the data arrives. A business directory is a good example. A restaurant has hours, cuisine, and a reservations policy. A hardware shop has trade accounts and parking. A hotel has star ratings and room types. Forcing them all into the same set of columns means drowning in NULLs or building a rats' nest of one-to-many extension tables.

PostgreSQL's JSONB type lets you store arbitrary JSON documents in a column while keeping the rest of your row fully relational. It is not a NoSQL escape hatch — it is an indexed, queryable, patchable document field that sits inside a SQL table. You can WHERE on it, join against it, aggregate over it, and update individual keys without rewriting the whole document.

This chapter works through all the operators and patterns you'll reach for most often. The exercises are short. Run every one, then try a variation of your own.

The Scenario

Portsmouth's city council maintains a public business directory. Every business gets a short relational record — an ID, a name, an address, and a neighbourhood — but the supplementary detail varies so widely by business type that a single fixed schema would be unworkable.

The businesses table therefore stores all of that variable metadata in a details JSONB column. Each business category contributes its own keys:

Category	JSONB keys present
restaurant	cuisine, price_range, rating, hours, tags, accepts_reservations, outdoor_seating
retail	subcategory, rating, hours, tags, payment_methods, has_parking
service	subcategory, rating, hours, tags, appointment_required, specialties
accommodation	subcategory, star_rating, rating, hours, amenities, room_types, price_per_night
entertainment	subcategory, rating, hours, tags, live_music, age_restriction

The hours value is itself a nested object keyed by day of week (mon–sun). Each day is either null (closed) or {"open": "HH:MM", "close": "HH:MM"}. Several businesses also carry optional keys that appear only for their category, like social (social media handles), happy_hour, or approved_brands.

This heterogeneity is intentional — it is what makes the data a good vehicle for learning JSONB operators.

Exercise Goals

By the end of this chapter you will be able to:

- Extract values from a JSONB document using `->`, `->>`, and `#>>`, and explain why the operator choice changes the result type.
 - Filter rows using the containment operator `@>` and the key-existence operators `?`, `?|`, and `?&`.
 - Create a GIN index on a JSONB column and confirm that PostgreSQL uses it.
 - Update a document in place with `jsonb_set` and the `||` merge operator without rewriting the whole column.
 - Expand a JSONB array into a set of rows with `jsonb_array_elements` and aggregate the results.
 - Write `jsonb_path_query` expressions to filter on nested values.
-

Installation

1 — PostgreSQL

If PostgreSQL 16 is not already installed:

```
sudo apt update
sudo apt install -y postgresql-16 postgresql-client-16
```

Confirm the server is running:

```
pg_lsclusters
```

You should see a cluster at version 16 in the online state.

2 — Python 3.12 and psycopg

```
sudo apt install -y python3.12 python3.12-venv
```

Create a virtual environment in the book's working directory and install the PostgreSQL driver:

```
python3.12 -m venv .venv
source .venv/bin/activate
pip install "psycopg[binary]"
```

Note: psycopg (version 3) is the modern Python driver for PostgreSQL. It is not available as an apt package for Python 3.12, so we install it via pip into the virtual environment. The [binary] extra bundles a pre-compiled C extension so you don't need libpq-dev.

Loading the Data

Create the database

```
# Create a role matching your OS user (skip if it already exists)
sudo -u postgres createuser --createdb "$(whoami)"

# Create the portsmith database
createdb portsmith
```

Run the seed script

From the book/ directory, with the virtual environment active:

```
python data/ch01_seed.py
```

Expected output:

```
Connecting to: dbname=portsmith
Creating schema ...
Inserting 48 businesses ...
Done – 48 rows in businesses.
```

Verify the load

Open a psql session:

```
psql portsmith
```

Run these checks. If all three pass, the data is correct.

Check 1 — Row count and column types:

```
\d businesses
```

```
Table "public.businesses"
  Column      | Type      | Collation | Nullable |
Default
-----+-----+-----+-----+
id            | integer   |           | not null |
nextval('businesses_id_seq'::regclass)
name          | text      |           | not null |
address       | text      |           | not null |
neighbourhood | text      |           | not null |
details       | jsonb     |           | not null |
```

Check 2 — Counts by neighbourhood:

```
SELECT neighbourhood, COUNT(*) AS businesses
FROM   businesses
GROUP BY neighbourhood
ORDER BY neighbourhood;
```

```
neighbourhood | businesses
-----+-----
Harbour District |          9
Industrial Port  |          7
Northgate       |          9
Old Town        |          9
Riverside       |          9
```

Check 3 — Counts by category:

```
SELECT details->>'category' AS category, COUNT(*) AS businesses
FROM businesses
GROUP BY category
ORDER BY category;
```

category	businesses
accommodation	5
entertainment	6
restaurant	15
retail	12
service	10

(5 rows)

If all three match, proceed to the exercises.

Exercises

Exercise 1 — The Three Extraction Operators

JSONB has two families of accessor. The arrow operators (-> and ->>) navigate one key at a time. The path operator (#>>) jumps straight to a deeply nested value. The difference between them is the **type they return**.

1.1 — Arrow operators

Run these two queries and compare the result types:

```
-- -> returns JSONB
SELECT name,
       details -> 'rating' AS rating_jsonb
FROM businesses
LIMIT 5;

-- ->> returns TEXT
SELECT name,
       details ->> 'rating' AS rating_text
FROM businesses
LIMIT 5;
```

Look carefully at the column headings in psql — one will show jsonb, the other text. This distinction matters the moment you try to do arithmetic:

```
-- This works: cast text to numeric
SELECT name,
       (details ->> 'rating')::numeric AS rating
FROM   businesses
ORDER BY rating DESC
LIMIT  5;
```

name	rating
Lighthouse Bookshop	4.9
Finch & Sons Barbers	4.9
Portsmouth Vet. Clinic	4.8
River Bend Bakery	4.8
Quarter Note Jazz Club	4.8

(5 rows)

Key point: Use `->` when you want to keep working with JSONB (e.g., to navigate deeper). Use `->>` when you need the final value as text for comparisons, casting, or display.

1.2 — The path operator #>>

Navigating nested keys with chained `->` quickly becomes unreadable. The path operator takes an array of keys and returns the leaf as text in one step.

```
-- Chained arrows (verbose)
SELECT name,
       details -> 'hours' -> 'fri' ->> 'close' AS friday_close
FROM   businesses
WHERE  details ->> 'category' = 'restaurant'
LIMIT  8;

-- Path operator (equivalent, cleaner)
SELECT name,
       details #>> '{hours,fri,close}' AS friday_close
FROM   businesses
WHERE  details ->> 'category' = 'restaurant'
LIMIT  8;
```

Both produce the same result. Notice that some rows return NULL — those restaurants are closed on Fridays (the `fri` key is JSON null).

name	friday_close
The Gilded Clam	22:30

Anchor & Oar Tavern		01:00
Bella Napoli		23:00
Le Petit Bistro		14:30
Dragon Palace		23:00
Spice Garden		23:00
Sol y Mar		23:00
Mango Bay Caribbean		22:30

(8 rows)

Exercise 2 — Filtering with Containment and Key Existence

2.1 — Containment: @>

The containment operator checks whether the left JSONB document contains all the key-value pairs in the right document. It is the most common way to filter on JSONB values.

Find all waterfront businesses:

```
SELECT name, neighbourhood
FROM businesses
WHERE details @> '{"tags": ["waterfront"]}'
ORDER BY neighbourhood, name;
```

name		neighbourhood
Anchor & Oar Tavern		Harbour District
Harbour Inn		Harbour District
Harbour View Theater		Harbour District
Mariners Rest B&B		Harbour District
Portsmith Fish Market		Harbour District
Saltbox Gallery		Harbour District
The Gilded Clam		Harbour District
Tidal Wave Surf Shop		Harbour District
Old Brewery Tap		Industrial Port
The Rusty Anchor		Industrial Port

(10 rows)

Why this works: The @> operator checks whether the *tags array* on the left contains the tags array ["waterfront"] on the right — array containment is supported natively.

2.2 — Key existence: ?

The ? operator tests whether a key exists in a JSONB object (or a value exists in a JSONB array). Unlike @>, it does not check the value — only presence.

Find all businesses that publish social media handles:

```
SELECT name,  
       details -> 'social' AS social_links  
FROM   businesses  
WHERE  details ? 'social'  
ORDER BY name;
```

name	social_links
Bella Napoli	{"instagram": "@bellanapoli_portsmith"}
Le Petit Bistro	{"instagram": "@lepetitbistro"}
Quarter Note Jazz Club	{"instagram": "@quarternote_portsmith"}
Spice Garden	{"instagram": "@spicegarden_portsmith"}
The Gilded Clam	{"facebook": "thegildedclam", "instagram": "@gildedclam"}
The Riverside Vegan	{"instagram": "@riverside_vegan"}

(6 rows)

2.3 — Any-key and all-key existence: ?| and ?&

?| returns true if *at least one* of the listed keys exists. ?& returns true only if *all* listed keys exist.

```
-- Businesses that offer either delivery OR takeaway  
SELECT name  
FROM   businesses  
WHERE  details ?| ARRAY['delivery', 'takeaway']  
ORDER BY name;  
  
-- Businesses that offer BOTH delivery AND takeaway  
SELECT name  
FROM   businesses  
WHERE  details ?& ARRAY['delivery', 'takeaway']  
ORDER BY name;
```

Run both and note the difference in result count. The ?& set is a strict subset of the ?| set.

Exercise 3 — Missing Keys vs. NULL Values

JSONB has two distinct ways to represent “no value”: a JSON `null` and a **missing key**. They behave differently in queries.

In our data, a business closed on Monday can be represented two ways:

- "mon": null — the key exists, the value is JSON null
- the mon key is absent entirely

The seed data uses "mon": null for closures, so we can observe this.

3.1 — The difference matters for ?

```
-- ? checks key EXISTENCE, not value
SELECT name
FROM   businesses
WHERE  (details -> 'hours') ? 'sun'    -- key exists (even if null)
ORDER BY name;
```

Run it. You should get **43 rows** — every business that uses a day-keyed hours object. The five accommodation businesses are excluded because their hours look like {"reception": "07:00-23:00"} and have no sun key.

Now try what seems like a natural follow-up:

```
-- This does NOT do what you might expect
SELECT name
FROM   businesses
WHERE  (details -> 'hours' -> 'sun') IS NOT NULL
ORDER BY name;
```

Run it. You still get **43 rows** — identical to the first query.


Why? This is one of JSONB's most important gotchas:

JSONB null is not SQL NULL.

When the seed script stores "sun": None in Python, json.dumps produces "sun": null — a JSON null *value* under an existing key. When PostgreSQL evaluates details -> 'hours' -> 'sun' on such a row, it returns the JSONB value null. That is a real value — not the absence of a value — so IS NOT NULL evaluates to TRUE.

SQL NULL only appears when the key itself is **missing** from the object. That is what happens for the five accommodation businesses: details -> 'hours' returns {"reception": "..."}, and then -> 'sun' on that object finds no such key and returns SQL NULL, so IS NOT NULL is FALSE.

The same logic, traced as a flowchart instead of read as prose:



To correctly distinguish the three states, use `jsonb_typeof()`:

State	details -> 'hours' -> 'sun'	jsonb_typeof(...)	IS NOT NULL
Key missing (accommodation)	SQL NULL	SQL NULL	FALSE
Key present, closed (null)	JSONB null	'null'	TRUE ← gotcha
Key present, open (object)	JSONB object	'object'	TRUE

The correct query for “businesses actually open on Sunday”:

```

SELECT name
FROM businesses

```



```
WHERE jsonb_typeof(details -> 'hours' -> 'sun') = 'object'
ORDER BY name;
```

This returns **25 rows** — only the businesses with a real hours object for Sunday. `jsonb_typeof` returns 'object' only for JSON objects, 'null' for JSON null, and SQL NULL for a missing key (which then fails the = test).

Rule of thumb: Never use `IS NOT NULL` to test whether a JSONB navigation result is a “real” value. Use `jsonb_typeof(...)` to check the actual JSON type, or navigate one level deeper (e.g., check `details #>> '{hours,sun,open}'`) — the `#>>` operator converts JSON null to SQL NULL, so a deeper path naturally filters it out.

3.2 — Handling accommodation differently

Hotels and hostels store hours as a single `"reception": "HH:MM-HH:MM"` string rather than a day-keyed object. This means the `? 'sun'` test returns false for them even though they may be open 24/7. Write a query that finds all businesses that are either:

- Open on Sunday (day-keyed hours with a non-null sun), **or**
- An accommodation with 24/7 reception

```
SELECT name, neighbourhood,
CASE
    WHEN jsonb_typeof(details -> 'hours' -> 'sun') =
'object'
        THEN (details #>> '{hours,sun,open}') || '-' ||
              (details #>> '{hours,sun,close}')
    WHEN details #>> '{hours,reception}' = '24/7'
        THEN '24/7'
    ELSE details #>> '{hours,reception}'
END AS sunday_availability
FROM businesses
WHERE jsonb_typeof(details -> 'hours' -> 'sun') = 'object'
      OR (details -> 'hours') ? 'reception'
ORDER BY neighbourhood, name;
```

Exercise 4 — GIN Indexes

Without an index, every JSONB query is a sequential scan — PostgreSQL reads every row and evaluates the expression. For a 48-row toy dataset that is unnoticeable, but with millions of rows it becomes a bottleneck.

A **GIN** (Generalised Inverted Index) index over a JSONB column indexes every key and value inside every document. It accelerates @>, ?, ?|, and ?& queries.

4.1 — Observe the plan before indexing

```
EXPLAIN (ANALYZE, BUFFERS)
SELECT name
FROM businesses
WHERE details @> '{"tags": ["waterfront"]}';
```

With 48 rows you will see Seq Scan (sequential scan, i.e. reading every row) — that is expected. On a larger table this would say Seq Scan with a high cost estimate.

4.2 — Create the GIN index

```
CREATE INDEX idx_businesses_details_gin
ON businesses
USING GIN (details);
```

4.3 — Observe the plan after indexing

```
EXPLAIN (ANALYZE, BUFFERS)
SELECT name
FROM businesses
WHERE details @> '{"tags": ["waterfront"]}';
```

On a small table PostgreSQL may still choose a sequential scan (the planner knows the table fits in memory). Force it to use the index to see the mechanism:

```
SET enable_seqscan = off;

EXPLAIN (ANALYZE, BUFFERS)
SELECT name
FROM businesses
WHERE details @> '{"tags": ["waterfront"]}';

SET enable_seqscan = on; -- always restore this
```

You should now see Bitmap Index Scan on idx_businesses_details_gin in the plan. The GIN index will pay for itself in production as the table grows.

What GIN indexes: The default jsonb_ops operator class indexes every key and value in the document, supporting @>, ?, ?|, ?&. A jsonb_path_ops class indexes only values (not keys), producing a

smaller index that is faster for @> but cannot accelerate ? queries.
Use jsonb_path_ops when you only ever query by value
containment and index size matters.

```
-- Alternative: value-only index (smaller, faster @>, no ? support)
CREATE INDEX idx_businesses_details_pathops
ON businesses
USING GIN (details jsonb_path_ops);
```

Exercise 5 — Updating Documents in Place

JSONB columns are immutable at the document level — you cannot change a single key without rewriting the whole value. PostgreSQL gives you two tools to do that rewrite concisely.

5.1 — jsonb_set: replace or add one value

The Lighthouse Bookshop just received a flood of five-star reviews. Update its rating to 5.0:

```
UPDATE businesses
SET details = jsonb_set(details, '{rating}', '5.0')
WHERE name = 'Lighthouse Bookshop';
```

Verify:

```
SELECT name, details ->> 'rating' AS rating
FROM businesses
WHERE name = 'Lighthouse Bookshop';
```

jsonb_set(target, path, new_value) — the path is a text array of keys. If the key already exists it is replaced; if it does not exist it is added (by default — there is a fourth parameter to suppress creation).

5.2 — jsonb_set on a nested key

The Gilded Clam is now closing an hour later on Sundays:

```
UPDATE businesses
SET details = jsonb_set(
    details,
    '{hours,sun,close}',
    '"21:00"' -- note: a JSON string must be
double-quoted
)
WHERE name = 'The Gilded Clam';
```

```

SELECT name,
       details #>> '{hours,sun,close}' AS new_sunday_close
FROM   businesses
WHERE  name = 'The Gilded Clam';

```

5.3 — || merge operator: add or overwrite multiple keys at once

The || operator merges two JSONB objects, with the right side winning on conflicts. Use it to add a verified flag and update review_count in a single statement:

```

UPDATE businesses
SET   details = details || '{"verified": true, "review_count":
999}'
WHERE name = 'Anchor & Oar Tavern';

SELECT name,
       details ->> 'verified'      AS verified,
       details ->> 'review_count' AS reviews
FROM   businesses
WHERE  name = 'Anchor & Oar Tavern';

```

Caution: || does a *shallow* merge. If details contains a nested object like hours and you merge another object that also has hours, the entire hours value is replaced — not deep-merged. Use jsonb_set for targeted nested updates.

Exercise 6 — Expanding Arrays with jsonb_array_elements

The tags key in most business documents is a JSON array. To query across tags — for example, to find the most common ones across the whole directory — you need to expand the array into individual rows.

6.1 — Expand one business’s tags

```

SELECT name,
       jsonb_array_elements_text(details -> 'tags') AS tag
FROM   businesses
WHERE  name = 'The Gilded Clam';

```

name	tag
The Gilded Clam	waterfront
The Gilded Clam	seafood
The Gilded Clam	romantic
The Gilded Clam	outdoor_seating

(4 rows)

`jsonb_array_elements_text` is a set-returning function: each element of the array becomes its own row. The result has more rows than the input.

6.2 — Count tags across all businesses

```
SELECT tag, COUNT(*) AS occurrences
FROM   businesses,
       jsonb_array_elements_text(details -> 'tags') AS tag
GROUP BY tag
ORDER BY occurrences DESC
LIMIT 15;
```

tag	occurrences
waterfront	10
takeaway	5
family_friendly	4
family_owned	4
organic	3
...	
(15 rows)	

Your exact numbers will differ — this shows the query pattern. The comma between `businesses` and the `jsonb_array_elements_text(...)` call is a **lateral join** (PostgreSQL expands the function for each input row automatically when used in the `FROM` clause this way).

6.3 — Only businesses with multiple specific tags

Find restaurants that are both `vegan_options` and `takeaway`:

```
SELECT name, details ->> 'cuisine' AS cuisine
FROM   businesses
WHERE  details @> '{"tags": ["vegan_options"]}'
AND    details @> '{"tags": ["takeaway"]}';
```

Because each `@>` call is independent, both conditions must hold. This is more efficient than expanding the array when the GIN index is in place.

Exercise 7 — Path Queries with `jsonb_path_query`

The `@?` operator and `jsonb_path_query()` function implement **SQL/JSON path language** — a mini-query language for navigating and filtering within a JSONB document. It is more expressive than chained arrow operators for conditional navigation.

7.1 — Simple path existence

Find businesses where the social links include Instagram:

```
SELECT name
FROM businesses
WHERE details @? '$.social.instagram';
```

`$.social.instagram` means: start at the root (`$`), descend to `social`, then to `instagram`. `@?` returns true if the path resolves to at least one value.

7.2 — Filter on a nested value

Find restaurants with a rating above 4.5:

```
SELECT name,
       details ->> 'rating' AS rating
FROM businesses
WHERE details @? '$ ? (@.category == "restaurant" && @.rating > 4.5)';
```

name	rating
Bella Napoli	4.6
Le Petit Bistro	4.7
River Bend Bakery	4.8
Spice Garden	4.6
The Gilded Clam	4.5

(5 rows)

The `?` (filter) syntax inside a path expression is equivalent to a `WHERE` clause inside the document.

7.3 — Find businesses open late on Friday

Businesses that close at or after 21:00 on Friday (using string comparison, which works correctly for 24-hour HH:MM strings):

```
SELECT name,
       details #>> '{hours,fri,close}' AS friday_close
FROM businesses
WHERE details @? '$.hours.fri.close ? (@ >= "21:00")'
ORDER BY friday_close DESC, name;
```

name	friday_close
Bella Napoli	23:00
Dragon Palace	23:00
Harbour View Theater	23:00
The Hungry Scholar	23:00
Spice Garden	23:00
Sol y Mar	23:00

Mango Bay Caribbean	22:30
The Gilded Clam	22:30
Thai Orchid	22:00
The Riverside Vegan	22:00
Northgate Grocers	21:00
Lotus Spa & Wellness	20:00

(12 rows)

Limitation to notice: Businesses that close *after midnight* — Anchor & Oar, The Clocktower Pub, The Rusty Anchor, and others — show a close of "01:00" or "02:00". Lexicographically, "02:00" < "21:00", so these businesses are *excluded* from the above results even though they are open very late. This is a design trade-off of storing times as plain strings. One solution is to store closing times past midnight as "25:00", "26:00", etc. — an unusual but practical convention for 24-hour time arithmetic.

7.4 — Extract values with jsonb_path_query

@? is a boolean test. jsonb_path_query() returns the actual matched values:

```

SELECT name,
       jsonb_path_query(details, '$.hours.fri.close') AS
friday_close_json
FROM   businesses
WHERE  details @? '$.hours.fri'
ORDER BY name
LIMIT 8;

```

The returned values are JSONB (note the quotes around the time strings). Use jsonb_path_query_first(...) #>> '{}' to extract a single scalar as text without the surrounding quotes.

Summary — What You Should Now Know

You have worked through the core JSONB toolkit. Here is what each operator and function you used actually does:

Tool	What it does
-> 'key'	Navigate one level; returns JSONB
->> 'key'	Navigate one level; returns text
#>> '{a,b,c}'	Navigate a path; returns text ; JSON null becomes SQL NULL

Tool	What it does
@> '{"k": "v"}'	Containment test; accelerated by GIN
? 'key'	Key exists test (regardless of value); accelerated by GIN
? / ?&	Any-key / all-keys exist; accelerated by GIN
jsonb_typeof(val)	Returns 'object', 'array', 'string', 'number', 'boolean', or 'null'; SQL NULL for a missing key
jsonb_set(col, path, val)	Replace or insert a nested value
col \ '{"k": "v"}'	Shallow-merge a document
jsonb_array_elements_text(col)	Expand a JSON array into rows
@? 'path'	SQL/JSON path existence test
jsonb_path_query(col, 'path')	SQL/JSON path — return matched values

Remember: JSONB null $\neq$ SQL NULL. Use `jsonb_typeof()` — not `IS NOT NULL` — when you need to distinguish a missing key from a key that exists but holds a JSON null value.

The key design insight from this chapter is that JSONB lets you keep a relational spine — the columns your queries always need (`id`, `name`, `neighbourhood`) — while storing everything else in a document whose shape can vary row by row. The GIN index means you do not pay a query performance penalty for that flexibility.

In the next chapter you will add point geometry to the `businesses` table and use PostGIS to answer spatial questions: which businesses are within 500 metres of the harbour, and which neighbourhood does each one belong to.

Going further: PostgreSQL 14+ supports the `jsonpath` type natively. The `jsonb_path_query_array` and `jsonb_path_query_first` variants are useful for pagination and single-value extraction. For write-heavy workloads, profile whether JSONB or a computed stored column (Chapter 16) gives better INSERT/UPDATE throughput on your hardware.

Chapter 2 — PostGIS: Geospatial Queries on Real Geometry

| “A city is not a list of rows. It is a shape on the ground.”

Background

Every interesting question about a city is ultimately a spatial question. Which businesses are near the harbour? Which neighbourhood is this address in? How large is the industrial waterfront? Relational databases answer these badly when location is stored as a text field or a pair of float columns — there is no native concept of “within”, “contains”, or “distance”.


PostGIS is a PostgreSQL extension that adds first-class geometry and geography types, plus several hundred spatial functions and operators. It turns PostgreSQL into a full spatial database: you can store points, lines, and polygons; index them with GIST; and answer proximity, containment, and area queries in SQL without an external GIS system.

This matters in practice far beyond mapping applications. Address geocoding, logistics routing, fraud detection (is this login coming from the expected region?), real estate valuation, and urban planning all reach for spatial queries. PostGIS is the standard tool for all of them in the PostgreSQL ecosystem.

The Scenario

The Portsmouth business directory from Chapter 1 stores each business’s neighbourhood as a plain text column. That works for simple filtering, but it cannot answer *where* questions: it cannot find businesses near a given coordinate, cannot verify that a business address actually falls inside its declared neighbourhood, and cannot measure distances.

This chapter adds a point geometry to every business record, then introduces three new spatial tables:



Portsmouth's six neighbourhoods, parks, and featured businesses — rendered directly from this chapter's own neighborhoods, parks, and businesses tables via `utils/render_map.py`. Every shape and dot on this map is a row you're about to query.

Table	Geometry type	What it holds
neighborhoods	POLYGON	Boundary polygons for Portsmouth's six neighbourhoods
parks	POLYGON	Six public parks and green spaces
city_infrastructure	LINESTRING	Twelve named road segments

All coordinates are in WGS-84 (SRID 4326), the same coordinate system used by GPS and most web mapping APIs.

Exercise Goals

By the end of this chapter you will be able to:

- Store and inspect POLYGON geometry using WKT (Well-Known Text) and ST_GeomFromText.
 - Run proximity searches with ST_DWithin, and understand why the ::geography cast matters for distance accuracy.
 - Perform spatial joins using ST_Within and ST_Contains, and explain the difference between them.
 - Compute polygon areas in square kilometres using ST_Area with the geography type.
 - Find the nearest feature for every row using ST_Distance and a CROSS JOIN LATERAL.
 - Create a GIST spatial index and confirm that PostgreSQL uses it.
-

Installation

1 — PostGIS server package

PostGIS is a separate package from PostgreSQL. On Debian/Ubuntu:

```
sudo apt install -y postgresql-16-postgis-3
```

2 — Enable PostGIS in the database

Connect to the portsmith database and enable the extension:

```
psql portsmith
```

```
CREATE EXTENSION IF NOT EXISTS postgis;
```

Verify it loaded:

```
SELECT postgis_full_version();
```

You should see a long string beginning with POSTGIS="3.x.x". If you see an error about the extension not existing, the server package is not installed.

Loading the Data

Prerequisites

Chapter 1's seed script must have been run first — the businesses table must exist. If it does not:

```
python data/ch01_seed.py
```

Run the Chapter 2 seed

```
python data/ch02_seed.py
```

Expected output:

```
Connecting to: dbname=portsmith
Enabling PostGIS extension ...
Applying DDL ...
Inserting 6 neighbourhoods ...
Inserting 6 parks ...
Inserting 12 road segments ...
Updating 48 business locations ...
```

Done:

```
businesses with geometry : 48
neighbourhoods           : 6
parks                     : 6
road segments             : 12
```

Verify the load

Open psql portsmith and run these four checks.

Check 1 — businesses now has a geometry column:

```
\d businesses
```

You should see a new geom column of type geometry(Point,4326).

Check 2 — neighbourhood table structure and row count:

```
SELECT name, population,
       ST_AsText(geom) AS wkt_preview
FROM   neighborhoods
ORDER BY name;
```

name	population	wkt_preview
-----+-----		
+-----		

Harbour District		4200		POLYGON((-1.805 50.69,...))
Industrial Port		2100		POLYGON((-1.77 50.69,...))
Northgate		18500		POLYGON((-1.83 50.732,...))
Old Town		6800		POLYGON((-1.805 50.71,...))
Riverside		9300		POLYGON((-1.773 50.71,...))
University Quarter		11200		POLYGON((-1.83 50.71,...))

(6 rows)

Check 3 — parks and roads:

```
SELECT COUNT(*) FROM parks;
SELECT COUNT(*) FROM city_infrastructure;
```

Both should return 6 and 12 respectively.

Check 4 — SRID on all tables:

```
SELECT f_table_name, f_geometry_column, srid, type
FROM geometry_columns
WHERE f_table_schema = 'public'
ORDER BY f_table_name;
```

f_table_name		f_geometry_column		srid		type
businesses		geom		4326		POINT
city_infrastructure		geom		4326		LINestring
neighborhoods		geom		4326		POLYGON
parks		geom		4326		POLYGON

(4 rows)

If all four pass, proceed to the exercises.

Exercises

Exercise 1 — Neighbourhood Polygons and WKT

Geometry in PostGIS is often loaded from **Well-Known Text** (WKT), a human-readable representation of shapes. Understanding WKT lets you read geometry from SQL output, write it in queries, and reason about what you stored.

1.1 — Read a polygon in WKT

```
SELECT name, ST_AsText(geom) AS wkt
FROM neighborhoods
WHERE name = 'Harbour District';
```

name	wkt
Harbour District	POLYGON((-1.805 50.69,-1.77 50.69,-1.77 50.71,...))

WKT for a polygon is `POLYGON((x1 y1, x2 y2, ...))`. In geographic coordinates the convention is *longitude latitude* (x then y), matching the X/Y convention in maths. The ring must **close** — the last coordinate must equal the first.

1.2 — Inspect the bounding box

`ST_Envelope` returns the minimum bounding rectangle for any geometry:

```
SELECT name,
       ST_XMin(ST_Envelope(geom)) AS west,
       ST_XMax(ST_Envelope(geom)) AS east,
       ST_YMin(ST_Envelope(geom)) AS south,
       ST_YMax(ST_Envelope(geom)) AS north
FROM   neighborhoods
ORDER BY name;
```

Use this output to confirm each polygon sits in the right part of the coordinate space (longitudes around -1.75 to -1.83 , latitudes around 50.69 to 50.76).

1.3 — Count vertices

```
SELECT name,
       ST_NPoints(geom) AS vertex_count
FROM   neighborhoods
ORDER BY name;
```

Each neighbourhood polygon has five vertices (four corners plus the closing duplicate). A real city boundary imported from an OS/census shapefile might have thousands.

1.4 — Understand the SRID

`ST_SRID` returns the spatial reference identifier stored with the geometry:

```
SELECT name, ST_SRID(geom) AS srid
FROM   neighborhoods
LIMIT 3;
```

SRID 4326 is the WGS-84 system used by GPS. Every geometry in these tables carries the same SRID, which means they can be compared and joined directly. If SRIDs differ, PostGIS will return an error — a deliberate safety check.

The geometry/geography distinction: In PostGIS, geometry stores coordinates in whatever units the SRS defines. For SRID 4326 that means degrees. When you ask for a distance between two geometry points in 4326, you get a value in *degrees* — nearly useless for human-scale distances. The geography type, by contrast, always works on a spheroid and returns distances in *metres*. You can cast any SRID-4326 geometry to geography with `::geography` to get metre-based calculations. The exercises use this cast throughout.

Exercise 2 — Proximity Search with ST_DWithin

`ST_DWithin(a, b, distance)` returns true when the distance between a and b is at most distance. With geography inputs the distance is in metres.

2.1 — Find businesses within 500 m of Portsmouth Pier

Portsmouth's main pier entrance sits at approximately (−1.785, 50.700). Find every business within 500 metres of it:

```
SELECT b.name,
       b.neighbourhood,
       ROUND(
         ST_Distance(b.geom::geography,
                     ST_Point(-1.785,
                             50.700)::geography)::numeric
       ) AS distance_m
FROM   businesses b
WHERE  ST_DWithin(b.geom::geography,
                  ST_Point(-1.785, 50.700)::geography,
                  500)
ORDER BY distance_m;
```

name	neighbourhood	distance_m
The Gilded Clam	Harbour District	70
Anchor & Oar Tavern	Harbour District	179
Tidal Wave Surf Shop	Harbour District	307
Harbour Inn	Harbour District	437
Portsmouth Fish Market	Harbour District	484

(5 rows)

All five results are in the Harbour District — exactly what you would expect for a search centred on the pier.

2.2 — Why ::geography matters

Try the same query without the cast. Note that `ST_Point(...)` without an explicit SRID gets SRID 0, which PostGIS refuses to compare against SRID-4326 geometry — so you must use `ST_SetSRID`:

```
SELECT b.name,
       ROUND(ST_Distance(b.geom,
                        ST_SetSRID(ST_Point(-1.785, 50.700),
4326))::numeric,
           6) AS distance_degrees
FROM   businesses b
WHERE  ST_DWithin(b.geom,
                ST_SetSRID(ST_Point(-1.785, 50.700), 4326),
                500)
ORDER BY distance_degrees;
```

The WHERE clause now has a threshold of 500 — but in *degrees*. One degree of latitude is about 111 km, so this matches businesses within roughly 55,000 km: all 48 of them, from the whole city. The `distance_degrees` values in the output are tiny fractions like 0.001000, not useful numbers.

rows returned: 48 (the entire city, not 5)

Rule: For any distance query on SRID-4326 data, always cast to `::geography` in `ST_DWithin` and `ST_Distance`. The cast has negligible performance cost at the scales used in city-level data.

2.3 — Try different radii

How many businesses fall within 1 km of the pier? 2 km?

```
SELECT radius_m,
       COUNT(*) AS business_count
FROM   (VALUES (500), (1000), (2000)) AS radii(radius_m)
CROSS JOIN LATERAL (
    SELECT 1
    FROM    businesses b
    WHERE   ST_DWithin(b.geom::geography,
                      ST_Point(-1.785, 50.700)::geography,
                      radius_m)
) AS matches
GROUP BY radius_m
ORDER BY radius_m;
```


radius_m	business_count
500	5
1000	8
2000	17

(3 rows)

The pier sits in the southern Harbour District. Doubling the radius to 2 km pulls in more of the central neighbourhoods, but the northern Northgate businesses are nearly 5 km away — a reminder that Portsmouth stretches a significant distance north from the waterfront.

Exercise 3 — Spatial Joins with ST_Within and ST_Contains

A **spatial join** links rows from two tables based on a geometric relationship rather than a key match. The most common are containment tests: does point A fall inside polygon B?

3.1 — Which neighbourhood is each business in?

```

SELECT b.name,
       b.neighbourhood          AS declared_neighbourhood,
       n.name                   AS postgis_neighbourhood
FROM   businesses b
JOIN   neighborhoods n ON ST_Within(b.geom, n.geom)
ORDER BY n.name, b.name;
```

ST_Within(a, b) returns true when geometry a lies completely inside geometry b. The query should return all 48 businesses, each matched to its neighbourhood.

Confirm the declared neighbourhood column always matches postgis_neighbourhood:

```

SELECT COUNT(*) AS mismatches
FROM   businesses b
JOIN   neighborhoods n ON ST_Within(b.geom, n.geom)
WHERE  b.neighbourhood <> n.name;
```

mismatches
0

(1 row)

Zero mismatches: the text column from Chapter 1 and the geometry are consistent.

3.2 — ST_Contains is the inverse

ST_Contains(a, b) returns true when polygon a contains geometry b — it is the mirror of ST_Within. These two queries produce identical results:

```
-- Point inside polygon
SELECT b.name FROM businesses b
JOIN  neighborhoods n ON ST_Within(b.geom, n.geom)
WHERE n.name = 'Old Town'
ORDER BY b.name;

-- Polygon contains point
SELECT b.name FROM businesses b
JOIN  neighborhoods n ON ST_Contains(n.geom, b.geom)
WHERE n.name = 'Old Town'
ORDER BY b.name;
```

```
name
-----
Bella Napoli
Finch & Sons Barbers
Le Petit Bistro
Old Town Hardware
Portsmith Accountancy Ltd.
Portsmith Arms Hotel
Portsmith Legal Group
Portsmith Tailors
The Clocktower Pub
(9 rows)
```

Both return the same 9 Old Town businesses.

The subtle difference: ST_Contains(A, B) requires that no point of B lies on the boundary of A. A point sitting exactly *on* a polygon edge would fail ST_Contains but pass ST_Covers. In practice, coordinates rarely land precisely on a boundary, so the two functions behave identically for most point-in-polygon work. Reach for ST_Covers if you need to include boundary-touching cases explicitly.

3.3 — Check for gaps: businesses in no neighbourhood

This query finds any businesses whose geometry falls outside every neighbourhood polygon — useful for catching data quality problems:

```
SELECT b.name, b.neighbourhood
FROM  businesses b
WHERE NOT EXISTS (
    SELECT 1
    FROM  neighborhoods n
```

```

        WHERE ST_Within(b.geom, n.geom)
    );

name | neighbourhood
-----+-----
(0 rows)

```

All 48 businesses are inside a neighbourhood polygon.

Exercise 4 — Computing Area in Square Kilometres

`ST_Area` returns the area of a polygon. Called on geometry (degrees), it returns a value in square degrees — meaningless to most people. Called on geography, it returns square metres.

4.1 — Incorrect: area in square degrees

```

SELECT name,
       ROUND(ST_Area(geom)::numeric, 6) AS area_sq_degrees
FROM   neighborhoods
ORDER BY area_sq_degrees DESC;

```

The numbers look tiny (around 0.0007). They are geometrically correct but impossible to interpret as real-world area because degrees are not uniform units.

4.2 — Correct: area in square metres, then kilometres

```

SELECT name,
       ROUND((ST_Area(geom::geography) / 1e6)::numeric, 2) AS
area_km2
FROM   neighborhoods
ORDER BY area_km2 DESC;

```

name	area_km2
Northgate	17.59
Old Town	5.53
Harbour District	5.50
University Quarter	4.32
Riverside	3.98
Industrial Port	3.14

(6 rows)

Northgate is by far the largest neighbourhood — it spans the entire northern width of the city. Industrial Port is the smallest, a tight strip of dockside land. Old Town and Harbour District are almost identical in area despite having very different characters.

Why the numbers differ slightly from manual estimates:

ST_Area on geography uses the WGS-84 spheroid, which accounts for the fact that the Earth is not a perfect sphere. At latitude 50.7° the correction is small (under 0.3%) but present.

4.3 — Population density

Combine the computed area with the population column:

```
SELECT name,
       population,
       ROUND((ST_Area(geom::geography) / 1e6)::numeric, 2) AS
area_km2,
       ROUND(
         (population / (ST_Area(geom::geography) /
1e6))::numeric
       ) AS pop_per_km2
FROM neighborhoods
ORDER BY pop_per_km2 DESC;
```

name	population	area_km2	pop_per_km2
University Quarter	11200	4.32	2592
Riverside	9300	3.98	2340
Old Town	6800	5.53	1230
Northgate	18500	17.59	1052
Harbour District	4200	5.50	763
Industrial Port	2100	3.14	668

(6 rows)

The University Quarter and Riverside are the most densely populated neighbourhoods — student housing and riverside apartments pack a lot of people into compact areas. The Industrial Port is the least dense despite its small size; much of it is warehouse and dock rather than residential.

Exercise 5 — Nearest Park Using ST_Distance and a Lateral Join

Finding the nearest feature from another table requires a **lateral join** — a subquery that can reference columns from the outer query row by row.

5.1 — Distance from one business to one park

ST_Distance between two geography values returns metres:

```

SELECT b.name AS business,
       p.name AS park,
       ROUND(ST_Distance(b.geom::geography,
                          p.geom::geography)::numeric) AS
distance_m
FROM   businesses b,
       parks p
WHERE  b.name = 'The Gilded Clam'
ORDER BY distance_m;

```

business	park	distance_m
The Gilded Clam	Harbourside Park	682
The Gilded Clam	Dockside Green	1315
The Gilded Clam	Market Square Gardens	2143
The Gilded Clam	Riverside Walk Park	2899
The Gilded Clam	University Grounds	3060
The Gilded Clam	Northgate Recreation Ground	5006

(6 rows)

The nearest park to The Gilded Clam is Harbourside Park, 682 m away. The ST_Distance to a polygon returns the distance from the point to the nearest point on the polygon’s boundary — zero when the point is inside the polygon.

5.2 — Nearest park for a single business (lateral pattern)

The canonical pattern for “nearest one thing” is ORDER BY ... LIMIT 1 inside a lateral subquery:

```

SELECT b.name,
       nearest.park_name,
       nearest.distance_m
FROM   businesses b
CROSS JOIN LATERAL (
    SELECT p.name AS
park_name,
           ROUND(ST_Distance(b.geom::geography,
                              p.geom::geography)::numeric) AS
distance_m
FROM     parks p
ORDER BY b.geom::geography <-> p.geom::geography
LIMIT 1
) AS nearest
WHERE  b.name = 'Quarter Note Jazz Club';

```

name	park_name	distance_m
Quarter Note Jazz Club	University Grounds	233

(1 row)

The <-> operator is the KNN (k-nearest-neighbour) distance operator. When used in an ORDER BY clause inside a lateral join, PostGIS can accelerate it with the GIST index (which you will create in Exercise 6). Using <-> in ORDER BY is preferred over ORDER BY ST_Distance(...) for this reason.

5.3 — Nearest park for every business

Remove the WHERE filter to run across all 48 businesses:

```
SELECT b.name                                AS business,
       b.neighbourhood,
       nearest.park_name,
       nearest.distance_m
FROM   businesses b
CROSS JOIN LATERAL (
    SELECT p.name                                AS
park_name,
           ROUND(ST_Distance(b.geom::geography,
                             p.geom::geography)::numeric) AS
distance_m
        FROM   parks p
        ORDER BY b.geom::geography <-> p.geom::geography
        LIMIT  1
    ) AS nearest
ORDER BY b.neighbourhood, b.name;
```

Scan the results. You should see a clean pattern: each neighbourhood's businesses point to the park in that same neighbourhood. Notice that three Riverside businesses (Portsmouth Pharmacy, Portsmouth Veterinary Clinic, Riverside Cinema) and University Bookshop all show distance_m = 0 — their coordinates fall *inside* the park polygon. ST_Distance to a polygon returns zero when the point is contained within it.

5.4 — Aggregate: average walking distance to a park, by neighbourhood

```
SELECT b.neighbourhood,
       ROUND(AVG(nearest.distance_m)::numeric) AS avg_distance_m
FROM   businesses b
CROSS JOIN LATERAL (
    SELECT ROUND(ST_Distance(b.geom::geography,
                             p.geom::geography)::numeric) AS
distance_m
        FROM   parks p
        ORDER BY b.geom::geography <-> p.geom::geography
        LIMIT  1
    ) AS nearest
```

```
GROUP BY b.neighbourhood
ORDER BY avg_distance_m;
```

This gives a rough “walkability to green space” metric per neighbourhood. Old Town ranks first — Market Square Gardens sits centrally within the neighbourhood. Northgate comes last despite having the largest park, because the park is in the far north of the district while businesses cluster near the southern edge.

Exercise 6 — GIST Indexes and Spatial Query Plans

Without an index, every spatial query is a sequential scan: PostgreSQL reads every row, applies the geometry function, and discards non-matching rows. For a 48-row dataset that is instant, but at 48 million rows it is not.

PostGIS spatial queries are accelerated by **GIST** (Generalised Search Tree) indexes. A GIST index on a geometry column builds a tree of bounding boxes, allowing the planner to quickly eliminate large portions of the table.

6.1 — Observe the plan before indexing

```
EXPLAIN (ANALYZE, BUFFERS)
SELECT name
FROM businesses
WHERE ST_DWithin(geom::geography,
                 ST_Point(-1.785, 50.700)::geography,
                 500);
```

On the 48-row businesses table you will see something like:

```
Seq Scan on businesses (cost=0.00..4.60 rows=1 width=...) ...
  Filter: (st_dwithin(...))
  Rows Removed by Filter: 43
```

A sequential scan is expected on a tiny table — the planner knows the overhead of an index lookup exceeds the cost of reading all 48 rows.

6.2 — Create GIST indexes on the geometry columns

```
CREATE INDEX idx_businesses_geom
ON businesses USING GIST (geom);

CREATE INDEX idx_neighborhoods_geom
ON neighborhoods USING GIST (geom);

CREATE INDEX idx_parks_geom
```

```

ON parks USING GIST (geom);

CREATE INDEX idx_city_infrastructure_geom
ON city_infrastructure USING GIST (geom);

```

6.3 — The geometry/geography index split

With `enable_seqscan` off, verify which plan the proximity query uses:

```

SET enable_seqscan = off;

EXPLAIN (ANALYZE, BUFFERS)
SELECT name
FROM businesses
WHERE ST_DWithin(geom::geography,
                ST_Point(-1.785, 50.700)::geography,
                500);

SET enable_seqscan = on;

```

You will see — perhaps surprisingly — that the planner *still* uses a sequential scan, even though `idx_businesses_geom` exists:

```

Seq Scan on businesses (cost=100000000000.00...) ...
  Filter: st_dwithin((geom)::geography, ...)

```

Why? The GIST index was built on `geom` (type `geometry`). The query filters on `geom::geography` (type `geography`). These are different types — the index is not usable for geography operations.

To accelerate geography-based distance queries you need a **functional index** on the cast:

```

CREATE INDEX idx_businesses_geom_geography
ON businesses USING GIST (CAST(geom AS geography));

```

Now with `enable_seqscan` off:

```

SET enable_seqscan = off;

EXPLAIN (ANALYZE, BUFFERS)
SELECT name
FROM businesses
WHERE ST_DWithin(geom::geography,
                ST_Point(-1.785, 50.700)::geography,
                500);

SET enable_seqscan = on;

```

```

Index Scan using idx_businesses_geom_geography on businesses
  Index Cond: ((geom)::geography &&
               _st_expand('...', '500'))

```



```
Filter: st_dwithin((geom)::geography, ...)
Rows Removed by Filter: 2
```

The index is now used. The index condition uses && (bounding-box overlap on the spheroid) to quickly discard most rows, then `st_dwithin` rechecks the exact distance for the survivors.

6.4 — Verify the spatial join plan

The containment join uses geometry-to-geometry operations, so `idx_businesses_geom` (the plain geometry index) is used there:

```
SET enable_seqscan = off;

EXPLAIN (ANALYZE, BUFFERS)
SELECT b.name, n.name AS neighbourhood
FROM   businesses b
JOIN   neighborhoods n ON ST_Within(b.geom, n.geom);

SET enable_seqscan = on;
```

Nested Loop ...

```
-> Seq Scan on neighborhoods n  (6 rows – tiny table)
-> Index Scan using idx_businesses_geom on businesses b
    Index Cond: (geom @ n.geom)
    Filter: st_within(geom, n.geom)
```

For each neighbourhood polygon, `idx_businesses_geom` is used to find businesses whose bounding box falls inside the neighbourhood's bounding box (`geom @ n.geom`), then `st_within` rechecks the exact containment.

The index rule: A GIST index on `geom` geometry accelerates geometry-to-geometry operations (`ST_Within`, `ST_Contains`, `ST_Intersects`, &&). A GIST index on `CAST(geom AS geography)` accelerates geography operations (`ST_DWithin(...::geography, ...)`, <-> on geography). If you query both ways, create both indexes — they coexist happily on the same table.

Summary — What You Should Now Know

You have worked through the core PostGIS toolkit for points, polygons, and linear features. Here is a reference for everything used:

Function / operator	What it does
<code>ST_GeomFromText(wkt, srid)</code>	Parse WKT into a geometry with the given SRID
<code>ST_AsText(geom)</code>	Format geometry as WKT for display
<code>ST_SetSRID(ST_MakePoint(lon, lat), srid)</code>	Construct a point geometry
<code>ST_Point(lon, lat)::geography</code>	Construct a point and cast to geography
<code>geom::geography</code>	Cast a 4326 geometry to geography (distances now in metres)
<code>ST_SRID(geom)</code>	Return the SRID stored with a geometry
<code>ST_NPoints(geom)</code>	Count vertices in a geometry
<code>ST_Envelope(geom)</code>	Return the bounding-box rectangle
<code>ST_XMin/XMax/YMin/YMax(geom)</code>	Extract bounding-box extents
<code>ST_DWithin(a, b, d)</code>	True when the distance between a and b is $\leq d$
<code>ST_Distance(a, b)</code>	Distance between two geometries (metres for geography)
<code>a::geography <-> b::geography</code>	KNN distance operator; use in ORDER BY for index acceleration
<code>ST_Within(a, b)</code>	True when a lies completely inside b
<code>ST_Contains(a, b)</code>	True when a contains b (inverse of ST_Within)
<code>ST_Area(geom::geography)</code>	Area in square metres on the spheroid
<code>CROSS JOIN LATERAL (... LIMIT 1)</code>	Nearest-neighbour join pattern
<code>USING GIST (geom)</code>	Create a GIST spatial index

Geometry vs geography in one sentence: Use geometry for storage and when working with projected coordinate systems where units are already metres or feet. Cast to geography any time you need distance, area, or proximity results in real-world units from SRID-4326 data.

The coordinates added in this chapter will carry forward. Chapter 12 extends the `city_infrastructure` road network into a graph and uses a recursive CTE to find the shortest path between two intersections.

Going further: PostGIS supports many more geometry types — `MULTIPOLYGON` for areas with holes, `GEOMETRYCOLLECTION` for mixed types, and 3D geometries with a Z coordinate for elevation data. For routing specifically, `pgRouting` builds on PostGIS to provide Dijkstra, A, and turn-restriction-aware shortest-path algorithms over road network graphs. For importing real boundary data, `shp2pgsql` converts ESRI Shapefiles directly into PostGIS-compatible `INSERT` statements, and `ogr2ogr` handles GeoJSON, KML, GeoPackage, and dozens of other formats.*

Chapter 3 — Job Queues: FOR UPDATE SKIP LOCKED

| “A queue is just a table that everyone is racing to read.”

Background

Sooner or later almost every application needs a queue: a list of work items that a pool of workers processes one at a time, safely, without two workers ever grabbing the same item. The reflexive answer is to reach for a message broker — Redis, RabbitMQ, SQS, Kafka. Those tools earn their keep at serious scale. But if your data already lives in PostgreSQL, running a second system just to hand out rows to workers is often unnecessary complexity: another service to deploy, monitor, and keep consistent with the database.

PostgreSQL can do this job itself. The `FOR UPDATE SKIP LOCKED` row-locking clause, combined with an ordinary table, gives you an atomic “claim the next item and don’t let anyone else touch it” primitive — the same guarantee a dedicated queue product sells you, built out of two SQL keywords. This chapter builds a job queue from scratch: the schema, the atomic claim query, concurrent worker behaviour, stalled-job recovery, and a dead-letter path for jobs that keep failing.

This is not a toy exercise. This exact pattern — a status column, a claim query, a heartbeat, a dead-letter table — is what libraries like `river`, `oban` (Elixir), and countless in-house job runners implement on top of PostgreSQL in production.

The Scenario


Portsmouth’s permitting office processes a steady stream of permit applications: building work, business licenses, public events, signage, and demolitions. Each application needs to move through a review pipeline, and the office wants that processing to happen

asynchronously and reliably — work should never be lost, never double-processed, and a crashed reviewer process shouldn't leave an application stuck in limbo forever.

The jobs table models this as a queue. Every row is one permit application awaiting review. A status column tracks its life cycle, a priority column lets safety-critical work (demolitions) jump ahead of routine work (sign permits), and a retry counter with a companion dead_letter_jobs table handles applications whose processing keeps failing.

Column	Purpose
status	queued → in_progress → completed (or back to queued, or dead-lettered)
priority	1 (most urgent — demolitions) to 5 (least urgent — sign permits)
payload	JSONB — the permit application details
attempts / max_attempts	Retry bookkeeping
claimed_by / claimed_at	Which worker has the job, and since when
heartbeat_at	Updated periodically by the worker while it holds the job

The status column's life cycle, drawn out — this is what Exercises 2 through 5 each implement one piece of:



Every arrow here is a specific query you'll write by hand later in this chapter — there's no hidden state machine enforcing this, just the status column, the claim query, and the reclaim sweep agreeing on what each value means.

No extensions are required for this chapter — everything here is built on core PostgreSQL locking semantics.

Exercise Goals

By the end of this chapter you will be able to:

- Design a queue table schema, including a partial index tuned for the claim query's access pattern.
- Write the atomic `UPDATE ... FROM (SELECT ... FOR UPDATE SKIP LOCKED)` claim query, and explain why naive two-step approaches race.

- Observe, with two concurrent `psql` sessions, that `SKIP LOCKED` lets workers proceed past rows another worker is holding — and contrast that with the blocking behaviour of plain `FOR UPDATE`.
 - Implement a heartbeat/timeout mechanism that reclaims jobs abandoned by a crashed worker.
 - Route jobs that exhaust their retries into a dead-letter table.
 - Benchmark claim throughput at different concurrency levels with `pgbench`.
-

Installation

This chapter needs nothing beyond what Chapter 1 already set up: PostgreSQL 16 and a Python 3.12 virtual environment with `psycopg`. If you skipped Chapter 1, see its Installation section. You will also use `pgbench` for Exercise 6, which ships with the standard PostgreSQL client tools (`postgresql-client-16` on Debian/Ubuntu).

Loading the Data

Run the seed script

From the `book/` directory, with the virtual environment active:

```
python data/ch03_seed.py
```

Expected output:

```
Connecting to: dbname=portsmith
Creating schema ...
Inserting 45 jobs ...
Done – 45 rows in jobs, all queued.
```

The seed script is self-contained — it does not depend on Chapter 1 or 2's data.

Verify the load

Open `psql portsmith` and run these checks.

Check 1 — table structure:

```
\d jobs
```

Table "public.jobs"			
Column	Type	Collation	Nullable
Default			
id	bigint		not null
nextval('jobs_id_seq'::regclass)			
job_type	text		not null
payload	jsonb		not null
status	text		not null
'queued'::text			
priority	smallint		not null
5			
attempts	integer		not null
0			
max_attempts	integer		not null
3			
created_at	timestamp with time zone		not null
clock_timestamp()			
claimed_at	timestamp with time zone		
claimed_by	text		
heartbeat_at	timestamp with time zone		
completed_at	timestamp with time zone		
last_error	text		

Indexes:

"jobs_pkey" PRIMARY KEY, btree (id)

"idx_jobs_claim_order" btree (priority, created_at, id) WHERE status = 'queued'::text

"idx_jobs_status" btree (status)

Check constraints:

"jobs_status_check" CHECK (status = ANY (ARRAY['queued'::text, 'in_progress'::text, 'completed'::text, 'failed'::text]))

Check 2 — job counts by type and priority:

```
SELECT job_type, priority, COUNT(*) AS jobs
FROM jobs
GROUP BY job_type, priority
ORDER BY priority;
```

job_type	priority	jobs
demolition_permit	1	4
business_license	2	12
building_permit	3	15
event_permit	4	8
sign_permit	5	6

(5 rows)

Check 3 — everything starts queued:

```
SELECT status, COUNT(*) FROM jobs GROUP BY status;
```



```
status | count
-----+-----
queued |    45
(1 row)
```

If all three match, proceed to the exercises.

Note: If you re-run `ch03_seed.py` at any point during the exercises to reset to a clean state, it drops and recreates both `jobs` and `dead_letter_jobs`.

Exercises

Exercise 1 — Designing the Queue Schema

1.1 — Why a partial index

The claim query (which you'll write in Exercise 2) only ever looks at rows where `status = 'queued'`, ordered by `priority` then `created_at`. As the queue runs, the vast majority of rows will end up completed — a normal B-tree index on `(priority, created_at, id)` would faithfully index every one of those settled rows even though the claim query never looks at them.

`idx_jobs_claim_order` is a **partial index** — it only indexes rows matching `WHERE status = 'queued'`:

```
CREATE INDEX idx_jobs_claim_order
ON jobs (priority, created_at, id)
WHERE status = 'queued';
```

This keeps the index small regardless of how many historical jobs pile up in completed or failed state, because settled rows are never in it.

1.2 — Why `id` is part of the sort key, not just `created_at`

You might expect `ORDER BY priority, created_at` to be enough — oldest job in the highest-priority bucket goes first. But timestamps are not always unique. If two jobs are inserted in the same transaction, both can get an identical `created_at` (more on this below), and `ORDER BY` over tied values has no defined order. Appending the primary key, `id`, as a final tiebreaker guarantees a deterministic order even when timestamps collide:

```
ORDER BY priority ASC, created_at ASC, id ASC
```

A gotcha worth knowing: `now()` returns the **transaction's** start time, not the current wall-clock time — every call to `now()` inside the same transaction returns the same value. The seed script originally used `created_at TIMESTAMPTZ DEFAULT now()`, and because all 45 rows were inserted in one transaction, every single row ended up with an *identical* `created_at`. The fix is `clock_timestamp()`, which returns the actual current time at the moment it's evaluated, differing row to row even within one transaction. `jobs.created_at` uses `clock_timestamp()` for exactly this reason. This is the same family of surprise as the JSONB `null` vs. SQL `NULL` gotcha from Chapter 1: a function name that looks interchangeable with another is not.

1.3 — Confirm the index is used

```
EXPLAIN (ANALYZE, BUFFERS)
SELECT id
FROM   jobs
WHERE  status = 'queued'
ORDER BY priority ASC, created_at ASC, id ASC
FOR UPDATE SKIP LOCKED
LIMIT 1;
```

```
Limit  (cost=0.14..6.17 rows=1 width=24) (actual
time=0.022..0.022 rows=1 loops=1)
  Buffers: shared hit=3
  -> LockRows  (cost=0.14..12.20 rows=2 width=24) (actual
time=0.021..0.021 rows=1 loops=1)
    Buffers: shared hit=3
    -> Index Scan using idx_jobs_claim_order on jobs
(cost=0.14..12.18 rows=2 width=24) (actual time=0.010..0.010
rows=1 loops=1)
      Filter: (status = 'queued'::text)
      Buffers: shared hit=2
```

Now drop the index and run the identical query again inside a transaction you roll back (so the drop doesn't stick):

```
BEGIN;
DROP INDEX idx_jobs_claim_order;

EXPLAIN (ANALYZE, BUFFERS)
SELECT id
FROM   jobs
WHERE  status = 'queued'
ORDER BY priority ASC, created_at ASC, id ASC
FOR UPDATE SKIP LOCKED
LIMIT 1;

ROLLBACK;
```

```

Limit (cost=9.51..9.53 rows=1 width=24) (actual
time=0.071..0.072 rows=1 loops=1)
  -> LockRows (cost=9.51..9.54 rows=2 width=24) (actual
time=0.070..0.071 rows=1 loops=1)
    -> Sort (cost=9.51..9.52 rows=2 width=24) (actual
time=0.067..0.068 rows=1 loops=1)
      Sort Key: priority, created_at, id
      Sort Method: quicksort Memory: 27kB
    -> Bitmap Heap Scan on jobs (cost=4.16..9.50
rows=2 width=24) (actual time=0.018..0.027 rows=45 loops=1)
      Recheck Cond: (status = 'queued'::text)
    -> Bitmap Index Scan on idx_jobs_status
(cost=0.00..4.16 rows=2 width=0) (actual time=0.013..0.013 rows=45
loops=1)
      Index Cond: (status = 'queued'::text)

```

Without `idx_jobs_claim_order`, PostgreSQL still finds the queued rows (via `idx_jobs_status`), but it must then **sort all of them** to find the one with the lowest (`priority`, `created_at`, `id`) — an extra Sort node. With the tailored partial index, the rows are already stored in claim order, so PostgreSQL walks the index and stops at the first match. On a 45-row table the difference is invisible; on a busy production queue with a deep backlog, eliminating the sort on every single claim matters a great deal.

Exercise 2 — The Atomic Claim Query

2.1 — Why a naive two-step claim races

The tempting first approach is to SELECT a candidate row, then UPDATE it in a second statement:

```

-- Do not do this – it has a race condition
SELECT id FROM jobs WHERE status = 'queued' ORDER BY priority,
created_at, id LIMIT 1;
-- ... application reads id = 7 ...
UPDATE jobs SET status = 'in_progress' WHERE id = 7;

```

Between the SELECT and the UPDATE, nothing stops a second worker from running the exact same SELECT, reading the same `id = 7`, and also issuing the UPDATE. Both workers now believe they own job 7. This is a classic **check-then-act race condition** — the gap between reading and acting is exactly where two workers can interleave.

2.2 — FOR UPDATE closes the gap, but blocks

Locking the row as part of the SELECT closes the race:

```

        SELECT id FROM jobs WHERE status = 'queued' ORDER BY priority,
created_at, id
        FOR UPDATE LIMIT 1;

```

Now a second worker running the same query, in a separate transaction, **blocks** until the first worker's transaction commits or rolls back — correct, but it means every worker but one sits idle waiting for a lock instead of moving on to a different job. Section 3.1 demonstrates this.

2.3 — SKIP LOCKED lets workers move past each other

Adding SKIP LOCKED tells PostgreSQL: if the next candidate row is already locked by another transaction, don't wait for it — skip it and consider the row after it.

```

        SELECT id FROM jobs WHERE status = 'queued' ORDER BY priority,
created_at, id
        FOR UPDATE SKIP LOCKED LIMIT 1;

```

This is the piece that makes PostgreSQL usable as a concurrent queue: N workers can run this query at the same instant and each will walk away with a *different* row, with no blocking and no double-claims.

2.4 — The full atomic claim: lock, update, and return in one statement

Locking the row is only half the job — you still need to mark it in_progress before releasing the lock, and you want the whole thing to happen as a single round trip. Combine the SELECT ... FOR UPDATE SKIP LOCKED with the UPDATE using a CTE:

```

WITH next_job AS (
    SELECT id
    FROM jobs
    WHERE status = 'queued'
    ORDER BY priority ASC, created_at ASC, id ASC
    FOR UPDATE SKIP LOCKED
    LIMIT 1
)
UPDATE jobs
SET status = 'in_progress',
    claimed_at = now(),
    claimed_by = 'demo-worker',
    heartbeat_at = now(),
    attempts = attempts + 1
FROM next_job
WHERE jobs.id = next_job.id
RETURNING jobs.id, jobs.job_type, jobs.payload ->>

```

```
'application_id' AS application_id,
                jobs.priority, jobs.attempts, jobs.status;

 id |      job_type      | application_id | priority | attempts |
-----+-----
 1 | demolition_permit | DP-2024-0001  |         |         1 |
in_progress
(1 row)
```

This single statement is the entire claim operation: find the best candidate, skip anything locked, mark it claimed, and hand back its data — atomically, with no window for a race. This is exactly the query `data/ch03_worker.py` runs (see `CLAIM_SQL`). Run it again and the demolition permit at `id = 2` comes back next — `id = 1` is now `in_progress`, so it's no longer a candidate.

Roll this back if you ran it directly in `psql` and want to keep the queue clean for later exercises: wrap it in `BEGIN; ... ROLLBACK;`.

Exercise 3 — Simulating Concurrent Workers

3.1 — `SKIP LOCKED` vs. plain `FOR UPDATE`, in two `psql` sessions

Open two terminals with `psql portsmith` in each. In **Session A**, start a transaction, claim 5 rows, and hold the transaction open (don't commit yet):

```
-- Session A
BEGIN;
SELECT id, job_type FROM jobs WHERE status = 'queued'
ORDER BY priority, created_at, id
FOR UPDATE SKIP LOCKED LIMIT 5;

 id |      job_type
-----+-----
 1 | demolition_permit
 2 | demolition_permit
 3 | demolition_permit
 4 | demolition_permit
 5 | business_license
(5 rows)
```

Leave that transaction open. In **Session B**, run the identical query:

```
-- Session B (Session A is still open, holding locks on ids 1-5)
BEGIN;
SELECT id, job_type FROM jobs WHERE status = 'queued'
```

```
ORDER BY priority, created_at, id
FOR UPDATE SKIP LOCKED LIMIT 5;
```

```
id |      job_type
----+-----
 6 | business_license
 7 | business_license
 8 | business_license
 9 | business_license
10 | business_license
(5 rows)
```

Session B returns immediately with a **different** set of rows — it silently skipped 1–5 because they were locked, and moved on to the next five unlocked candidates. Commit or roll back both sessions to release the locks:

```
COMMIT;    -- run in both sessions
```

3.2 — Now try it with plain FOR UPDATE (no SKIP LOCKED)

Repeat the same two-session experiment, but drop SKIP LOCKED:


```
-- Session A
BEGIN;
SELECT id, job_type FROM jobs WHERE status = 'queued'
ORDER BY priority, created_at, id
FOR UPDATE LIMIT 5;
-- (leave this transaction open)

-- Session B – this will hang
BEGIN;
SELECT id, job_type FROM jobs WHERE status = 'queued'
ORDER BY priority, created_at, id
FOR UPDATE LIMIT 5;
```

Session B does not return. It is blocked, waiting for Session A's row locks to be released. Only once Session A runs COMMIT (or ROLLBACK) does Session B's query complete — and when it does, it returns the *same* five rows Session A had, now that they're unlocked again:

```
-- (Session B, after Session A commits)
id |      job_type
----+-----
 1 | demolition_permit
 2 | demolition_permit
 3 | demolition_permit
 4 | demolition_permit
 5 | business_license
(5 rows)
```

This is the difference in one sentence: **plain FOR UPDATE serializes workers through the lock; SKIP LOCKED lets them fan out across the table.** For a job queue, you always want the latter — a blocked worker is a wasted worker. Both runs side by side, as a timeline:



3.3 — Real concurrent workers with `ch03_worker.py`

Reset the data (`python data/ch03_seed.py`) if you ran the manual claim in Exercise 2.4 without rolling it back. Then launch two workers in the background, each capped to 10 jobs so they finish quickly:

```
python data/ch03_worker.py --worker-id w1 --max-jobs 10 --fail-rate 0 &
python data/ch03_worker.py --worker-id w2 --max-jobs 10 --fail-rate 0 &
wait
```

Each prints a running log as it claims and completes jobs, e.g.:

```
[w1] claimed job 2 (demolition_permit, attempt 1/3): DP-2024-0002
- processing for 0.09s
[w1] job 2 completed
[w2] claimed job 1 (demolition_permit, attempt 1/3): DP-2024-0001
- processing for 0.12s
[w2] job 1 completed
```

Once both finish, confirm the split was clean — 20 jobs completed total, no job claimed by both workers:

```
SELECT claimed_by, COUNT(*) FROM jobs WHERE status = 'completed'
GROUP BY claimed_by ORDER BY claimed_by;
```

claimed_by	count
w1	10
w2	10

(2 rows)

```
SELECT status, COUNT(*) FROM jobs GROUP BY status;
```

status	count
completed	20
queued	25

(2 rows)

10 + 10 = 20, matching exactly the `--max-jobs 10` cap on each worker — no overlaps, no double-processing, no lost jobs.

Exercise 4 — Heartbeat and Stalled-Job Recovery

Claiming a job is only half the reliability story. What happens if the worker that claimed a job crashes, gets OOM-killed, or loses its network connection halfway through? The row is stuck at `status = 'in_progress'` forever unless something notices and puts it back.

4.1 — The heartbeat column

`heartbeat_at` exists for exactly this. A well-behaved worker updates it periodically while it holds a job (see the `while` loop in `process_job()` in `ch03_worker.py`, which sends a heartbeat roughly every 2 seconds during simulated processing). If a job has been `in_progress` for a long time *and* its heartbeat has gone stale, that's a strong signal the worker holding it is dead — a live worker would have updated it.

4.2 — Simulate a crashed worker

Pick any queued job and manually walk it through what a real claim would do, then simulate a crash by never sending a heartbeat again:

```
UPDATE jobs
SET status = 'in_progress', claimed_at = now(),
    claimed_by = 'worker-crashed', heartbeat_at = now(),
    attempts = attempts + 1
```



```

WHERE id = 21;

-- Simulate 5 minutes of silence from the "crashed" worker
UPDATE jobs SET heartbeat_at = now() - interval '5 minutes' WHERE
id = 21;

```

4.3 — Run the reclaim sweep

```
python data/ch03_reclaim.py --timeout 30
```

```

Connecting to: dbname=portsmith
  job 21 (building_permit) stalled – requeued (attempt 1)

```

```
Done: 1 requeued, 0 dead-lettered.
```

The sweep looks for `in_progress` rows whose `heartbeat_at` is older than `--timeout` seconds (`FIND_STALLED_SQL` in `ch03_reclaim.py`), and since this job's attempts (1) is still below `max_attempts` (3), it goes back to `queued`:

```

SELECT id, status, attempts, claimed_by, heartbeat_at, last_error
FROM   jobs WHERE id = 21;

```

```

id | status | attempts | claimed_by | heartbeat_at
|               last_error
----+-----+-----+-----+-----+-----
+-----+-----+-----+-----+-----+-----
 21 | queued |         1 |           |           | stalled: no
heartbeat since ... (last claimed by worker-crashed)

```

Notice attempts stayed at 1 — the reclaim did not reset it. The attempt the crashed worker used is still counted; the job doesn't get a free retry just because the worker that lost it never got to report failure honestly.

4.4 — Why the timeout, not an outright deadline

A fixed timeout on the *last heartbeat* (rather than on total processing time) lets jobs run arbitrarily long as long as they keep proving they're alive. A 10-minute report-generation job with a 30-second heartbeat interval will never be mistaken for stalled, while a worker that dies mid-task is detected within one missed heartbeat window. Run `ch03_reclaim.py` again immediately — it correctly finds nothing to do, since the reclaimed job now has a fresh `queued` state and no stale `in_progress` row exists:

```
No jobs stalled beyond 30s – nothing to do.
```

Exercise 5 — Dead-Lettering Exhausted Jobs

Some jobs will never succeed no matter how many times you retry them — a malformed application, a permanently invalid address, a bug that only triggers on one particular payload. Retrying forever wastes worker time and can mask a real problem. Once a job has used its last attempt, it belongs in `dead_letter_jobs`: out of the active queue, but preserved for a human to inspect.

5.1 — The `dead_letter_jobs` schema

```
\d dead_letter_jobs
```

Column Default	Type	Collation	Nullable
id	bigint		not null
job_type	text		not null
payload	jsonb		not null
priority	smallint		not null
attempts	integer		not null
max_attempts	integer		not null
created_at	timestamp with time zone		not null
last_error	text		
failed_at	timestamp with time zone		not null

```
now()
```

It mirrors `jobs` but adds `failed_at` and drops the in-flight columns (`status`, `claimed_by`, `heartbeat_at`) that no longer mean anything once a job is out of the active queue.

5.2 — Push a job past its retry limit

Continuing with job 21 from Exercise 4 (now back to `queued` with `attempts = 1`), fast-forward it to its last attempt and let it stall again:

```
UPDATE jobs
SET    status = 'in_progress', claimed_at = now(),
      claimed_by = 'worker-crashed-2', heartbeat_at = now() -
interval '5 minutes',
      attempts = max_attempts
WHERE  id = 21;
```

```
python data/ch03_reclaim.py --timeout 30
```

Connecting to: dbname=portsmith
job 21 (building_permit) exhausted retries — dead-lettered

Done: 0 queued, 1 dead-lettered.

This time attempts (3) is no longer less than max_attempts (3), so the sweep's DEAD_LETTER_SQL runs instead: it DELETES the row from jobs and INSERTs it into dead_letter_jobs in one CTE, so the row is never visible in neither-place or both-places, even under concurrent access.

```
SELECT id FROM jobs WHERE id = 21;
```

```
id
----
(0 rows)
```

```
SELECT id, job_type, payload ->> 'application_id' AS
application_id,
      attempts, max_attempts, last_error, failed_at
FROM   dead_letter_jobs WHERE id = 21;
```

```
id |      job_type      | application_id | attempts | max_attempts
|                    | last_error    |          |          |
failed_at
-----+-----+-----+-----+-----
+-----+-----+-----+-----+-----
+-----+-----+-----+-----+-----
21 | building_permit   | BP-2024-0005  |         3 |         3 |
stalled: no heartbeat since ... (last claimed ...) | 2026-07-12
22:15:45.613583-04
(1 row)
```

The exact same dead-lettering logic runs inline inside `ch03_worker.py` when a job fails organically (not via a stall) on its last attempt — see `DEAD_LETTER_SQL` in that script. Whether a job dies from an explicit failure or from going silent, it ends up in the same place, with a record of why.

Exercise 6 — Benchmarking Claim Throughput with `pgbench`

6.1 — A benchmark script that doesn't drain the queue

`pgbench` repeatedly runs a SQL script against the database for a fixed duration, from any number of concurrent client connections — exactly the tool for measuring how the claim query holds up under contention. The catch: if the script just claims a job and leaves it claimed, 45 rows disappear from the pool in well under a second and every worker after

that finds nothing to do. `data/ch03_claim_bench.sql` works around this by claiming a job and then **immediately releasing it back to queued** in the same transaction:

```
BEGIN;

WITH next_job AS (
    SELECT id
    FROM jobs
    WHERE status = 'queued'
    ORDER BY priority ASC, created_at ASC, id ASC
    FOR UPDATE SKIP LOCKED
    LIMIT 1
)
UPDATE jobs
SET     status      = 'in_progress',
        claimed_at  = now(),
        claimed_by  = 'bench-' || :client_id,
        heartbeat_at = now()
FROM   next_job
WHERE  jobs.id = next_job.id
RETURNING jobs.id AS claimed_id \gset

UPDATE jobs
SET     status = 'queued', claimed_at = NULL, claimed_by = NULL,
        heartbeat_at = NULL
WHERE  id = :claimed_id;

COMMIT;
```

`:client_id` is a `pgbench` built-in variable — the number of the simulated client running this script. `\gset` captures the claimed row's `id` into a `pgbench` variable (`:claimed_id`) so the release step can target it. This keeps the pool at a constant 45 claimable rows for the whole benchmark run, so throughput reflects lock contention rather than the queue running dry.

6.2 — Run at increasing concurrency

```
pgbench -n -c 1 -j 1 -T 5 -f data/ch03_claim_bench.sql portsmith
pgbench -n -c 4 -j 4 -T 5 -f data/ch03_claim_bench.sql portsmith
pgbench -n -c 16 -j 16 -T 5 -f data/ch03_claim_bench.sql portsmith
```

`-c` is the number of simulated concurrent workers (clients), `-j` the number of `pgbench` threads driving them, and `-T 5` runs each test for 5 seconds. Results on the machine used to write this chapter:

Clients (-c)	tps	Avg. latency
1	799	1.25 ms

Clients (-c)	tps	Avg. latency
4	1,202	3.33 ms
16	2,570	6.23 ms

Your numbers will differ with your hardware — the shape of the result is what matters. Throughput climbs with concurrency because SKIP LOCKED lets additional workers keep finding unlocked rows to claim instead of queueing up behind each other; latency per transaction also climbs because more workers are contending for the same 45-row table and for CPU. Try dropping `idx_jobs_claim_order` (inside a `BEGIN; ... ROLLBACK; block`, as in Exercise 1.3) and rerunning the `-c 16` case — you should see tps fall, since every claim now pays for a Sort over all queued rows instead of an index walk.

6.3 — Where this stops being enough

pgbench here is exercising claim contention on 45 rows — a stand-in for “the working set the claim query actually scans.” A real permitting-office queue might hold a very different shape of data (thousands of queued rows, heavy skew toward one priority bucket), and the honest way to benchmark your own workload is to seed a table that resembles it, not a 45-row toy. Chapter 20 (pg_stat_statements and Query Performance) picks this thread back up: rather than a synthetic benchmark, it measures and fixes real slow queries observed in production traffic.

Summary — What You Should Now Know

You built a working, concurrency-safe job queue entirely out of core PostgreSQL features. Here is a reference for the pieces:

Tool	What it does
FOR UPDATE	Locks selected rows; other transactions selecting the same rows block until released
FOR UPDATE SKIP LOCKED	Locks selected rows; other transactions skip already-locked rows instead of blocking
WITH next_job AS (... FOR UPDATE SKIP LOCKED LIMIT 1) UPDATE ... FROM next_job	The atomic claim pattern: find, lock, and mark a row in one round trip

Tool	What it does
Partial index WHERE status = 'queued'	Keeps the claim-path index small regardless of how many settled rows accumulate
clock_timestamp() vs. now()	now() freezes at transaction start; clock_timestamp() reflects real wall-clock time on every call
heartbeat_at + timeout sweep	Detects and recovers jobs abandoned by a crashed worker
dead_letter_jobs	Removes permanently-failing jobs from the active queue while preserving them for inspection
pgbench -f script.sql	Benchmarks a custom SQL workload at a chosen concurrency level

The key design insight from this chapter is that a reliable queue is not one clever query — it’s the combination of an atomic claim, a way to detect workers that go silent, and a place for work that genuinely cannot succeed. PostgreSQL’s row-locking primitives give you the first two almost for free; the dead-letter table is just a normal table.

The jobs table you built here is reused directly in later chapters: Chapter 13 adds a trigger that fires NOTIFY on status changes so listeners can react to queue activity in real time, Chapter 14 uses advisory locks alongside it for leader-election patterns, and Chapter 19 schedules the reclaim sweep you wrote by hand in Exercise 4 to run automatically with pg_cron.

Going further: the pattern in this chapter — status column, atomic claim, heartbeat, dead-letter — is exactly what PostgreSQL-backed job-queue libraries like river (Go) and oban (Elixir) implement, with more polish around scheduling, uniqueness constraints, and observability. If you outgrow a single-database queue — cross-database work distribution, guaranteed ordering at very high throughput, or consumer groups — that’s the point at which a dedicated broker starts to earn its operational cost. For most applications below that scale, the table you just built is enough.

Chapter 4 — Full-Text Search: `tsvector`, Stopwords, and Ranking

“Grep finds characters. Full-text search finds meaning — or at least gets a lot closer.”

Background

`LIKE '%harbour%'` finds the literal substring “harbour”. It will not find “Harbour”, “harbours”, or a document about “the harbor” (different spelling) unless you handle every variation yourself, and it cannot tell you whether a match in the title is more relevant than a match buried in paragraph six. As soon as an application needs to search prose — meeting minutes, support tickets, product descriptions, articles — substring matching stops being enough. The usual reflex is to reach for Elasticsearch or a similar dedicated search engine.

PostgreSQL has had a full-text search engine built in since version 8.3. It tokenizes text into words, reduces those words to normalized root forms (*stemming*), discards low-information words like “the” and “of” (*stopwords*), and stores the result in a specialized `tsvector` type that a GIN index can search in milliseconds. A companion `tsquery` type lets you express boolean and phrase searches, and a family of ranking functions score how well each match fits the query. None of this requires an extension — it is core PostgreSQL, the same engine `websearch_to_tsquery`-powered search boxes on production sites are built on.

This chapter builds a search feature over a small document archive: how text becomes a `tsvector`, what stopwords removal actually throws away, how to keep the vector in sync as rows change, how query syntax differs between a developer-facing boolean query language and a plain-text search box, and how to rank and highlight results. The final exercise customizes the stopwords list itself — because in a *city* government’s document archive, the words “city” and “council” appear so often they stop being useful signals, exactly the kind of domain-specific tuning full-text search is designed to support.

The Scenario

Portsmouth’s city clerk publishes an archive of public records: council meeting minutes, zoning ordinances, and public notices. Residents need to search this archive by keyword — “what happened with the Riverside dog park?”, “which ordinances mention Canal Road?” — and get back the most relevant documents first, with a highlighted snippet showing why each one matched.

The `city_documents` table holds all three document types in one place, with a `body` column of plain prose text and a few relational columns for filtering:

Column	Purpose
<code>doc_type</code>	<code>council_minutes</code> , <code>zoning_ordinance</code> , or <code>public_notice</code>
<code>department</code>	Which city department published it (City Council, Public Works, ...)
<code>title</code>	Short document title
<code>body</code>	Full plain-text document body — the field full-text search runs over
<code>published_date</code>	When the document was published

The documents share recurring topics on purpose — the Riverside dog park, the Harbour District waterfront rezoning, the Canal Road bike lane — so that searches in the exercises return more than one hit, and ranking has something real to sort between. No extensions are required for this chapter: `tsvector`, `tsquery`, and every function used below are core PostgreSQL.

Exercise Goals

By the end of this chapter you will be able to:

- Convert text to a `tsvector` and read the `lexeme:position` notation it produces.
- Explain exactly what stopword removal and stemming do to a document, using `ts_debug` to see PostgreSQL’s tokenizer classify each word.
- Maintain a `tsvector` column automatically with a trigger, and index it with GIN for sub-millisecond searches.

- Write boolean queries with `to_tsquery` and understand why `plainto_tsquery` is the safer choice for a plain search box.
 - Rank results with `ts_rank` and `ts_rank_cd`, and generate highlighted snippets with `ts_headline`.
 - Build a custom text search configuration that treats domain-specific words as stopwords, and see it change real search results.
-

Installation

This chapter needs nothing beyond what Chapter 1 already set up: PostgreSQL 16 and a Python 3.12 virtual environment with `psycopg`. If you skipped Chapter 1, see its Installation section. Full-text search is core PostgreSQL — there is no extension to enable.

Loading the Data

Run the seed script

From the `book/` directory, with the virtual environment active:

```
python data/ch04_seed.py
```

Expected output:

```
Connecting to: dbname=portsmith
Creating schema ...
Inserting 30 documents ...
Done – 30 rows in city_documents.
```

The seed script is self-contained — it does not depend on any earlier chapter’s data. Note that it deliberately does **not** create a `tsvector` column or index; you build both by hand in Exercise 3, which is the point of the chapter.

Verify the load

Open `psql portsmith` and run these checks.

Check 1 — table structure:

```
\d city_documents
```

Table "public.city_documents"

Column	Type	Collation	Nullable
id	integer		not null
nextval('city_documents_id_seq'::regclass)			
doc_type	text		not null
department	text		not null
title	text		not null
body	text		not null
published_date	date		not null

Indexes:

- "city_documents_pkey" PRIMARY KEY, btree (id)
- "idx_city_documents_doc_type" btree (doc_type)
- "idx_city_documents_published_date" btree (published_date)

Check constraints:

- "city_documents_doc_type_check" CHECK (doc_type = ANY (ARRAY['council_minutes'::text, 'zoning_ordinance'::text, 'public_notice'::text]))

Check 2 — counts by document type:

```
SELECT doc_type, COUNT(*) AS documents
FROM city_documents
GROUP BY doc_type
ORDER BY doc_type;
```

doc_type	documents
council_minutes	10
public_notice	10
zoning_ordinance	10

(3 rows)

Check 3 — counts by department:

```
SELECT department, COUNT(*) AS documents
FROM city_documents
GROUP BY department
ORDER BY department;
```

department	documents
City Council	9
Finance	1
Parks & Recreation	3
Planning & Zoning	12
Public Works	5

(5 rows)

If all three match, proceed to the exercises.

Exercises

Exercise 1 — Converting Text to tsvector

1.1 — A basic conversion

`to_tsvector` is the function that turns plain text into PostgreSQL's searchable representation. Try it on one document's body:

```
SELECT title FROM city_documents WHERE id = 1;
```

```
title
```

```
-----
```

```
Council Minutes – Harbour District Waterfront Renovation Budget
```

```
SELECT to_tsvector('english', body) FROM city_documents WHERE id = 1;
```

```
'along':38 'approv':67 'begin':79 'boardwalk':34 'bond':53  
'budget':10 'citi':3  
'construct':76 'conven':5 'council':4,41,61 'cover':28 'debat':56  
'director':18  
'discuss':43 'district':14 'expect':77 'fall':82 'first':69  
'fund':44,74  
'grant':50 'harbour':13,40,72 'includ':46 'infrastructur':49  
'light':36  
'measur':54 'member':42 'new':32 'one':65 'pedestrian':33 'phase':  
26,70  
'plan':27 'portsmith':2 'present':22 'propos':9 'public':20  
'renov':16,73  
'repair':30 'review':7 'seawal':29 'six':63 'sourc':45 'state':48  
'three':25 'three-phas':24 'timelin':59 'upgrad':37 'vote':62  
'waterfront':15 'work':21
```

1.2 — Reading the format

Each entry is `'lexeme':position,position,...`. A **lexeme** is a normalized word form, not the literal text — notice `'vote':62` even though the source text says “voted”, `'approv':67` for “approve”, and `'renov':16,73` for both “renovation” (position 16) and “renovation” again later (position 73, from “harbour **renovation** funding”). PostgreSQL stems words to their root so that a search for “vote” also matches documents containing “voted”, “votes”, or “voting” — you don’t have to enumerate every inflection yourself.

The **positions** are the word’s ordinal location in the original text (1st word, 2nd word, ...). They exist for two reasons: `ts_rank` uses them to compute how spread out or clustered a document’s matches are, and phrase search (`<->`, used in Exercise 4) uses them to require words to appear adjacent to each other, not just anywhere in the document.

1.3 — Case and punctuation are already handled

Note that “Council” (capitalized, position 4) and “council” (lowercase, positions 41 and 61) both collapsed into the single lexeme ‘`council`’ with three positions — `to_tsvector` lowercases everything and strips punctuation as part of tokenization, before stemming ever runs.

Exercise 2 — What Stopword Removal Actually Throws Away

2.1 — Compare `simple` and `english` configurations

PostgreSQL ships several **text search configurations** — named bundles of tokenizer + dictionary rules. `simple` lowercases and tokenizes but does *no* stemming and drops *no* stopwords; `english` does both. Run the same body through each and count the resulting lexemes:

```
SELECT array_length(regexp_split_to_array(body, '\s+'), 1) AS
raw_words
FROM city_documents WHERE id = 1;

raw_words
-----
      80

SELECT array_length(tsvector_to_array(to_tsvector('simple', body)),
1) AS lexemes_simple,
      array_length(tsvector_to_array(to_tsvector('english', body)),
1) AS lexemes_english
FROM city_documents WHERE id = 1;

lexemes_simple | lexemes_english
-----+-----
              59 |                49
```

80 raw whitespace-separated words collapse to 59 distinct lexemes under `simple` (repeated words like “the” and “council” count once each, but nothing is removed or stemmed) and to 49 under `english` — ten fewer, because `english` additionally discards stopwords like “the”, “a”,

“of”, “to”, “in”, and “and” entirely. They carry no search-relevant meaning on their own, and indexing them would only bloat the index with entries that match nearly every document.

2.2 — Watch the tokenizer classify each word with `ts_debug`

`ts_debug` is the diagnostic function for seeing exactly what a configuration does to a piece of text, word by word — useful whenever a search isn’t matching what you expect and you need to know why. Run it on a short sentence:

```
SELECT alias, token, dictionaries, lexemes
FROM   ts_debug('english', 'The council voted to approve the
budget for the harbour renovation.')
WHERE  alias <> 'blank';
```

alias	token	dictionaries	lexemes
asciiword	The	{english_stem}	{}
asciiword	council	{english_stem}	{council}
asciiword	voted	{english_stem}	{vote}
asciiword	to	{english_stem}	{}
asciiword	approve	{english_stem}	{approv}
asciiword	the	{english_stem}	{}
asciiword	budget	{english_stem}	{budget}
asciiword	for	{english_stem}	{}
asciiword	the	{english_stem}	{}
asciiword	harbour	{english_stem}	{harbour}
asciiword	renovation	{english_stem}	{renov}

Every token gets classified (`asciiword` here — a plain word) and routed to the `english_stem` dictionary. Content words come back with a stemmed lexeme; stopwords (“The”, “to”, “the”, “for”, “the”) come back with an **empty** `lexemes` array — recognized, looked up, and explicitly discarded. That empty array is the entire mechanism: a word is a stopwords precisely when its dictionary entry maps it to nothing.

2.3 — When you’d reach for `simple` instead

`english` is the right default for prose. `simple` is useful for columns where stemming would be actively wrong — product SKUs, tag lists, or anything where “waterfront” and “waterfronts” should *not* be treated as the same token. `city_documents.body` is prose, so the rest of this chapter uses `english`.

Exercise 3 — A Maintained Column and a GIN Index

Computing `to_tsvector(body)` at query time works, but it means recomputing the same tokenization on every single search, over every row, every time. The standard pattern is to store the vector in its own column, keep it current as rows change, and index it.

3.1 — Add the column and backfill it

```
ALTER TABLE city_documents ADD COLUMN search_vector tsvector;

UPDATE city_documents
SET    search_vector = to_tsvector('english', title || ' ' ||
body);
```

Indexing `title` alongside `body` means a search for a word that only appears in the title (not repeated in the body text) still matches.

3.2 — Keep it current with a trigger

A column populated once goes stale the moment anyone edits `title` or `body`. A `BEFORE` trigger recomputes it on every write:

```
CREATE FUNCTION city_documents_search_vector_update() RETURNS
trigger AS $$
BEGIN
    NEW.search_vector := to_tsvector('english', NEW.title || ' '
|| NEW.body);
    RETURN NEW;
END;
$$ LANGUAGE plpgsql;

CREATE TRIGGER trg_city_documents_search_vector
BEFORE INSERT OR UPDATE OF title, body ON city_documents
FOR EACH ROW
EXECUTE FUNCTION city_documents_search_vector_update();
```

`OF title, body` scopes the trigger so it only fires when one of the two source columns actually changes — an `UPDATE` that only touches `published_date` doesn't pay for a needless re-tokenization.

Note — the manual way vs. the modern way: This trigger pattern is how every PostgreSQL version has supported derived `tsvector` columns, and it's still exactly right when the derived value depends on more than one column, as it does here (`title` and `body`). PostgreSQL 12 added **generated columns** (`GENERATED ALWAYS AS (...) STORED`), which handle the common case — a `tsvector` derived from a single column — with less boilerplate and no trigger function to maintain by hand. Chapter 16 revisits

this exact table and replaces this trigger with a generated column once body alone is the source. Both approaches produce an identical tsvector; the trigger is simply the general-purpose tool that works whenever the derivation is more than one column deep.

3.3 — Index it with GIN

```
CREATE INDEX idx_city_documents_search_vector
ON city_documents USING GIN (search_vector);
```

```
\d city_documents
```

Table "public.city_documents"			
Column	Type	Collation	Nullable
	Default		
-----+-----+-----+-----			
+-----+-----+-----+-----			
id	integer		not null
nextval('city_documents_id_seq'::regclass)			
doc_type	text		not null
department	text		not null
title	text		not null
body	text		not null
published_date	date		not null
search_vector	tsvector		

Indexes:

```
"city_documents_pkey" PRIMARY KEY, btree (id)
"idx_city_documents_doc_type" btree (doc_type)
"idx_city_documents_published_date" btree (published_date)
"idx_city_documents_search_vector" gin (search_vector)
```

Check constraints:

```
"city_documents_doc_type_check" CHECK (doc_type = ANY
(ARRAY['council_minutes'::text, 'zoning_ordinance'::text,
'public_notice'::text]))
```

Triggers:

```
trg_city_documents_search_vector BEFORE INSERT OR UPDATE OF
title, body ON city_documents FOR EACH ROW EXECUTE FUNCTION
city_documents_search_vector_update()
```

3.4 — Confirm the index is used

The @@ operator tests whether a tsvector matches a tsquery:

```
EXPLAIN (ANALYZE, BUFFERS)
SELECT id, title
FROM city_documents
WHERE search_vector @@ to_tsquery('english', 'harbour &
waterfront');
```

QUERY PLAN

```
-----  
Seq Scan on city_documents (cost=0.00..8.38 rows=1 width=36)  
(actual time=0.025..0.047 rows=5 loops=1)  
  Filter: (search_vector @@ '''harbour' &  
  'waterfront''':tsquery)  
  Rows Removed by Filter: 25  
  Buffers: shared hit=8
```

On 30 rows the planner reasonably decides a sequential scan is cheaper than an index lookup — same behaviour you saw with the JSONB GIN index in Chapter 1. Force the index to confirm it works, exactly as in that chapter:

```
SET enable_seqscan = off;  
  
EXPLAIN (ANALYZE, BUFFERS)  
SELECT id, title  
FROM   city_documents  
WHERE  search_vector @@ to_tsquery('english', 'harbour &  
waterfront');
```

```
SET enable_seqscan = on;  -- always restore this
```

Query Plan

```
-----  
Bitmap Heap Scan on city_documents (cost=12.98..16.99 rows=1  
width=36) (actual time=0.024..0.033 rows=5 loops=1)  
  Recheck Cond: (search_vector @@ '''harbour' &  
  'waterfront''':tsquery)  
  Heap Blocks: exact=5  
  Buffers: shared hit=10  
  -> Bitmap Index Scan on idx_city_documents_search_vector  
  (cost=0.00..12.98 rows=1 width=0) (actual time=0.014..0.014 rows=5  
  loops=1)  
    Index Cond: (search_vector @@ '''harbour' &  
    'waterfront''':tsquery)  
    Buffers: shared hit=5
```

Bitmap Index Scan on `idx_city_documents_search_vector` confirms the GIN index is doing the work. This is what makes full-text search viable at scale: instead of tokenizing every row's text on every query, PostgreSQL looks up each query lexeme directly in the index.

Exercise 4 — `to_tsquery` vs. `plainto_tsquery`

4.1 — `to_tsquery`: a boolean query language

to_tsquery expects an expression using explicit operators: & (AND), | (OR), ! (NOT), and <-> (phrase — “immediately followed by”):

```
SELECT id, title
FROM   city_documents
WHERE  search_vector @@ to_tsquery('english', 'harbour &
waterfront')
ORDER BY id;
```

id	title
----	-----
1	Council Minutes – Harbour District Waterfront Renovation Budget
8	Council Minutes – Harbour District Food Truck Permits
12	Zoning Ordinance – Harbour District Waterfront Height Variance
22	Public Notice – Public Hearing on Harbour District Waterfront Rezoning
30	Public Notice – Annual Fireworks Display and Street Closures

(5 rows)

! excludes a term. This finds documents about flooding that are *not* about the flood-mitigation infrastructure program specifically:

```
SELECT id, title
FROM   city_documents
WHERE  search_vector @@ to_tsquery('english', 'flood & !
mitigation')
ORDER BY id;
```

id	title
-----	-----
5	Council Minutes – Riverside Road Resurfacing Program

(1 row)

That single hit mentions “flooding” (which stems to flood) while discussing drainage damage, but never says “mitigation” — exactly what the query asked for.

4.2 — to_tsquery has no forgiveness for plain text

Feed it a raw phrase with no operator between the words, and it errors instead of guessing what you meant:

```
SELECT id, title
FROM   city_documents
WHERE  search_vector @@ to_tsquery('english', 'flood mitigation');
```

ERROR: syntax error in tsquery: "flood mitigation"

This is `to_tsquery`'s defining trade-off: it is a precise query language for code that constructs queries deliberately, and it is the wrong function to hand raw user input from a search box — a user typing “flood mitigation” with no operators will get an error page instead of results.

4.3 — `plainto_tsquery`: safe for a search box

`plainto_tsquery` takes plain text, tokenizes and stems it exactly like `to_tsvector` does, drops stopwords, and **ANDs** everything that survives. Feed it the same two words that made `to_tsquery` error out in 4.2, but typed the way an actual user would type them — with a stray “the” in front:

```
SELECT plainto_tsquery('english', 'the flood mitigation');

plainto_tsquery
-----
'flood' & 'mitig'
```

“the” never makes it into the query — `plainto_tsquery` drops it as a stopword, exactly the way `to_tsvector` would drop it from a document, and **ANDs** together whatever content words are left. Run it against the table:

```
SELECT id, title
FROM   city_documents
WHERE  search_vector @@ plainto_tsquery('english', 'the flood
mitigation')
ORDER BY id;
```

```
id | title
----+-----
10 | Council Minutes – Special Session on Flood Mitigation
Infrastructure
25 | Public Notice – City Council Special Session on Flood
Mitigation
(2 rows)
```

No error this time, and no operators to get wrong — `plainto_tsquery` took a sentence fragment with a stopword in it and produced exactly the query 4.2 had to write by hand. The trade-off runs the other way from `to_tsquery`: you get safety and stopword handling for free, but you lose the ability to express OR, NOT, or phrase search — `plainto_tsquery` always **ANDs** every surviving word together.

4.4 — | for OR searches

```
SELECT to_tsquery('english', 'dog | bike');
```

```

to_tsquery
-----
'dog' | 'bike'

SELECT id, title
FROM   city_documents
WHERE  search_vector @@ to_tsquery('english', 'dog | bike')
ORDER BY id;

id | title
----
+-----
5 | Council Minutes – Riverside Road Resurfacing Program
6 | Council Minutes – Riverside Dog Park Funding
9 | Council Minutes – Canal Road Bike Lane Expansion
15 | Zoning Ordinance – Reduced Parking Minimums Near Transit
Corridors
24 | Public Notice – Riverside Dog Park Ribbon-Cutting Event
28 | Public Notice – Canal Road Bike Lane Construction Schedule
30 | Public Notice – Annual Fireworks Display and Street Closures
(7 rows)

```

Document 5 (road resurfacing) and 15 (parking minimums) show up because each mentions “bike” in passing — the resurfacing minutes note the work being coordinated with “the planned Canal Road bike lane construction,” and the parking ordinance references “the city’s ongoing bike lane expansion.” OR genuinely means *either*, which is exactly why it’s useful and why it returns more, looser matches than AND.

Exercise 5 — Ranking with `ts_rank` and `ts_rank_cd`, and `ts_headline` Snippets

`@@` tells you whether a document matches — it says nothing about *how well*. A user searching a 30-document archive can eyeball every hit, but a user searching 30,000 documents needs the best matches first.

5.1 — `ts_rank`: weight by term frequency

```

SELECT id, title, round(ts_rank(search_vector, query)::numeric, 4)
AS rank
FROM   city_documents, to_tsquery('english', 'harbour &
waterfront') AS query
WHERE  search_vector @@ query
ORDER BY rank DESC;

id | title | rank
----

```

```

+-----+
+-----+
12 | Zoning Ordinance – Harbour District Waterfront Height
Variance | 0.1986
1 | Council Minutes – Harbour District Waterfront Renovation
Budget | 0.1959
22 | Public Notice – Public Hearing on Harbour District
Waterfront Rezoning | 0.1895
8 | Council Minutes – Harbour District Food Truck
Permits | 0.0999
30 | Public Notice – Annual Fireworks Display and Street
Closures | 0.0992
(5 rows)

```

ts_rank scores primarily on how often the query terms appear relative to the document’s overall length — it does not care where in the document they appear or how close together they are. Documents 12, 1, and 22 rank highest because “harbour” and “waterfront” are central, repeated themes in short documents; 8 and 30 rank lower because each mentions the harbour only in passing within a longer document about something else.

5.2 — ts_rank_cd: weight by proximity too (“cover density”)

```

SELECT id, title,
       round(ts_rank(search_vector, query)::numeric, 4) AS rank,
       round(ts_rank_cd(search_vector, query)::numeric, 4) AS
rank_cd
FROM   city_documents, to_tsquery('english', 'harbour &
waterfront') AS query
WHERE  search_vector @@ query
ORDER BY rank_cd DESC;

```

```

id |
title | rank | rank_cd
----+-----+
1 | Council Minutes – Harbour District Waterfront Renovation
Budget | 0.1959 | 0.1107
22 | Public Notice – Public Hearing on Harbour District
Waterfront Rezoning | 0.1895 | 0.1053
12 | Zoning Ordinance – Harbour District Waterfront Height
Variance | 0.1986 | 0.0700
30 | Public Notice – Annual Fireworks Display and Street
Closures | 0.0992 | 0.0545
8 | Council Minutes – Harbour District Food Truck
Permits | 0.0999 | 0.0542
(5 rows)

```

The order changes: document 12 had the *highest* `ts_rank` but drops to third under `ts_rank_cd`. `ts_rank_cd` additionally rewards matched terms that cluster close together in the text (“cover density”) — in documents 1 and 22, “harbour” and “waterfront” appear right next to each other (“Harbour District **Waterfront** Renovation...”); in document 12 they occur further apart. Neither function is universally “more correct” — `ts_rank` is the standard choice; `ts_rank_cd` is worth trying when word proximity itself signals relevance, as it often does for short phrase-like queries.

5.3 — Highlighted snippets with `ts_headline`

A ranked list of titles is useful; showing *why* a document matched, with the query terms highlighted in context, is what users actually expect from a search results page:

```
SELECT ts_headline('english', body, to_tsquery('english',
'harbour & waterfront'),
                                'StartSel=**, StopSel=**, MaxWords=25,
MinWords=10')
FROM   city_documents
WHERE  id = 1;
```

```
**Harbour** District **waterfront** renovation. The Director of
Public Works presented
```

`ts_headline` re-scans the original text (not the `tsvector`) looking for the query terms, extracts a window around the best-matching fragment bounded by `MinWords/MaxWords`, and wraps each match in `StartSel/StopSel` markers — `**...**` here, but in a web application you’d use `<mark>...</mark>` or similar. This is genuinely expensive compared to a `tsvector` lookup (it re-tokenizes the source text on every call), so use it only on the page of results you’re actually displaying, never inside a `WHERE` clause.

Exercise 6 — A Custom Text Search Configuration for Domain Stopwords

6.1 — The problem: some “content” words carry no information here

This is a city document archive. The words “Portsmouth”, “city”, and “council” appear constantly — they are structurally present in almost every record, the way “the” is present in almost every English sentence. Standard english stopwords removal has no way to know that, because it is tuned for general English, not this specific corpus:

```
SELECT count(*) FROM city_documents
WHERE search_vector @@ to_tsquery('english', 'city & council');
```

```
count
-----
      8
```

Eight of thirty documents — over a quarter of the entire archive — “match” a query on words that describe almost nothing about what makes any one of them relevant. Left alone, these words dilute ranking scores and inflate plain-text queries with noise terms the searcher didn’t mean to require.

6.2 — Build a stopwords file that extends the standard list

A text search configuration’s stopwords list is a plain text file, one word per line, that PostgreSQL reads from its `tsearch_data` directory. Start from the existing `english.stop` file so you keep every standard English stopwords, then append the domain-specific ones:

```
PG_SHAREDIR=$(pg_config --sharedir)
sudo bash -c "cat '$PG_SHAREDIR/tsearch_data/english.stop' >
'$PG_SHAREDIR/tsearch_data/portsmith_english.stop'"
sudo bash -c "printf 'portsmith\ncity\ncouncil\n' >>
'$PG_SHAREDIR/tsearch_data/portsmith_english.stop'"
```

Placing files here requires root, since `tsearch_data` lives under PostgreSQL’s shared install directory rather than anything database-owned. If your platform’s PostgreSQL package stores it elsewhere, `pg_config --sharedir` will always point at the right location.

6.3 — Wire the file into a dictionary and a configuration

A stopwords file alone does nothing — it has to be attached to a **text search dictionary**, and that dictionary mapped into a **text search configuration** that queries can actually reference:

```
CREATE TEXT SEARCH DICTIONARY portsmith_stem (
    TEMPLATE = snowball,
    LANGUAGE = english,
    STOPWORDS = portsmith_english
);

CREATE TEXT SEARCH CONFIGURATION public.portsmith_english (COPY =
pg_catalog.english);

ALTER TEXT SEARCH CONFIGURATION portsmith_english
ALTER MAPPING FOR asciiword, asciihword, hword_asciipart,
```

```
word, hword, hword_part
    WITH portsmith_stem;
```

TEMPLATE = snowball reuses PostgreSQL’s standard English stemming algorithm — only the stopword list changes, not the stemming rules. COPY = pg_catalog.english starts the new configuration as an exact clone of english (same tokenizer, same handling of numbers, emails, URLs, ...), and the ALTER MAPPING line is the one substitution: word tokens now route through portsmith_stem instead of the stock english_stem dictionary.

6.4 — Confirm the noise words are gone

```
SELECT to_tsvector('portsmith_english', body) FROM city_documents
WHERE id = 25;
```

```
'affect':37 'area':27 'attend':44 'basin':59 'comment':47
'conven':11 'develop':61
'discuss':16 'drain':56 'encourag':42 'engin':53 'flood':17,40,68
'given':4
'herebi':3 'infrastructur':19 'last':65 'mitig':18 'northgat':24
'notic':1
'open':31 'present':50 'propos':60 'provid':46 'public':34
'resid':36
'respons':63 'retain':25 'retent':58 'riversid':21 'session':14,29
'special':13
'spring':39 'staff':48 'storm':55 'wall':26 'year':66
```

Compare to to_tsvector('english', body) on the same row, which still carries 'citi':8,52, 'council':9, and 'portsmith':7 — three entries present under english and absent under portsmith_english. Everything else is untouched, because only the stopword list changed.

6.5 — Watch a noise-only query collapse to nothing

```
SELECT to_tsquery('portsmith_english', 'city & council');
```

NOTICE: text-search query contains only stop words or doesn't contain lexemes, ignored

```
to_tsquery
-----
```

```
(1 row)
```

Under portsmith_english, “city” and “council” are *both* stopwords, so a query built from nothing else has nothing left to search for — PostgreSQL says so explicitly instead of silently matching everything (or nothing) for an unclear reason.

6.6 — See it fix a real plain-text search

This is the exercise’s payoff: a resident searching for “Portsmouth dog park council” — a completely natural way to phrase it — under english:

```
SELECT plainto_tsquery('english', 'Portsmouth dog park council');

plainto_tsquery
-----
'portsmith' & 'dog' & 'park' & 'council'

SELECT id, title FROM city_documents
WHERE search_vector @@ plainto_tsquery('english', 'Portsmouth dog
park council')
ORDER BY id;

id | title
---+-----
24 | Public Notice – Riverside Dog Park Ribbon-Cutting Event
(1 row)
```

Only one hit. The council minutes that actually approved funding for the same dog park (document 6) never happens to use the literal word “Portsmouth,” so requiring it as an AND term silently excludes the single most relevant record. Under portsmouth_english:

```
SELECT plainto_tsquery('portsmouth_english', 'Portsmouth dog park
council');

plainto_tsquery
-----
'dog' & 'park'

SELECT id, title FROM city_documents
WHERE search_vector @@ plainto_tsquery('portsmouth_english',
'Portsmouth dog park council')
ORDER BY id;

id | title
---+-----
6  | Council Minutes – Riverside Dog Park Funding
24 | Public Notice – Riverside Dog Park Ribbon-Cutting Event
(2 rows)
```

Stripped down to the two words that actually distinguish this search — dog and park — both relevant documents come back. This is the concrete argument for a custom configuration: it is not a cosmetic tweak, it changes which documents a real user actually finds.

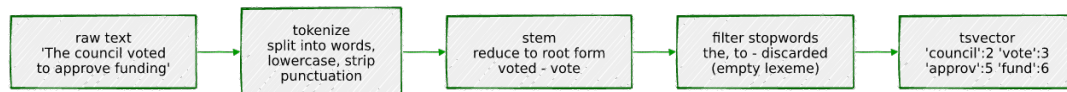
Note: the search_vector column built in Exercise 3 was populated with the stock english configuration, so it still contains 'citi', 'council', and 'portsmith' lexemes. The comparisons

above work because @@ only needs the *query*'s surviving lexemes to be a subset of the document's — extra lexemes in the document that the query no longer asks for are simply irrelevant. To fully adopt portsmith_english as the table's search configuration going forward, you would update the trigger from Exercise 3 to call `to_tsvector('portsmith_english', ...)` and rebuild `search_vector` for existing rows.

Summary — What You Should Now Know

Tool	What it does
<code>to_tsvector(config, text)</code>	Tokenizes, lowercases, stems, and removes stopwords, producing a searchable <code>lexeme:position</code> vector
<code>to_tsquery(config, expr)</code>	Parses a boolean query language (<code>&</code> , <code>\ </code> , <code>!</code> , <code><-></code>); errors on plain text with no operators
<code>plainto_tsquery(config, text)</code>	Tokenizes plain text the same way as <code>to_tsvector</code> , then ANDs every surviving lexeme — safe for a search box
<code>@@</code>	Tests whether a <code>tsvector</code> matches a <code>tsquery</code>
<code>ts_debug(config, text)</code>	Shows exactly how each token was classified and which lexeme (if any) it produced — the tool for “why didn't this match?”
GIN index on a <code>tsvector</code> column	Turns @@ from a sequential scan into an index lookup
<code>ts_rank / ts_rank_cd</code>	Score match quality by term frequency, or by term frequency and proximity (“cover density”)
<code>ts_headline(config, text, query, options)</code>	Re-scans the source text and returns a highlighted snippet — expensive; use only on displayed results
<code>CREATE TEXT SEARCH DICTIONARY ... (STOPWORDS = ...) + CREATE TEXT SEARCH CONFIGURATION</code>	Build a custom configuration with domain-specific stopwords

The key design insight from this chapter is that full-text search is not one function call — it's a pipeline:



You can inspect every stage of that pipeline with `ts_debug`, store its output with a maintained column and a GIN index, query against it with two very different trade-offs (`to_tsquery` for precision, `plainto_tsquery` for safety), and tune for your specific corpus by swapping the stopwords list — exactly the kind of tuning a generic search engine bolted on top of your database can't do without shipping your schema knowledge to it.

The `city_documents` table you built here is reused directly in later chapters: Chapter 6 (`pgvector`) adds embeddings to the same table for semantic search and builds a hybrid search combining `ts_rank` with cosine distance, and Chapter 16 (Generated Columns) replaces the hand-written trigger from Exercise 3 with a `GENERATED ALWAYS AS (...) STORED` column.

Going further: `websearch_to_tsquery` (PostgreSQL 11+) is worth knowing about even though this chapter didn't use it — it accepts the same `quoted-phrase` and `-exclude` syntax as a typical web search box (e.g. "`flood mitigation`" `-infrastructure`) while remaining as forgiving as `plainto_tsquery`, making it the best default for a real user-facing search field. For very large document sets, look into PostgreSQL's built-in support for weighted vectors (`setweight()`, ranking title matches above body matches) and consider whether a dedicated search engine becomes worthwhile once ranking quality and query latency requirements outgrow what a GIN index over a single table can deliver — the honest answer, for most applications, is later than you'd expect.

Chapter 5 — Fuzzy Matching:

pg_trgm

“Every registry that outlives a year of real data entry ends up with the same person in it twice, spelled two different ways.”

Background

Full-text search, from the last chapter, is built for a specific kind of imprecision: the same *word*, inflected differently (“vote”, “voted”, “voting”). It does nothing for a different, equally common kind of imprecision — the same word, **spelled** differently. “Portsmith” typed as “Portsmith” versus “Portsmyth”. “McAllister” versus “MacAllister”. A resident registry filled in by hand, twice, by two different clerks, on two different days. Stemming cannot fix a typo, because a typo isn’t a grammatical variant of the correct word — it’s a different string that a human reader recognizes as “close enough” and a computer, by default, does not.

pg_trgm is PostgreSQL’s answer: it breaks strings into overlapping three-character sequences (**trigrams**) and measures how many trigrams two strings share. Two spellings of the same name share most of their trigrams even when several letters differ; two unrelated strings usually share almost none. That single idea — measured, thresholded, and indexed — is enough to power “did you mean?” search boxes, deduplicate messy records, and accelerate substring LIKE queries that would otherwise force a sequential scan.

This chapter builds a small deduplication and search tool around that idea: what a trigram actually is, how `similarity()` scores a pair of strings, how to turn that into an index instead of an $O(n^2)$ scan, and — the harder, more honest part — where trigram matching’s judgment calls are, because no similarity threshold gets every case right.

The Scenario

Portsmouth’s resident registry has been filled in over years by different staff at different counters, and it shows: the same person appears twice under two spellings of their name often enough that nobody trusts a raw `COUNT(*)` on it anymore. Separately, the city’s 311 line fields phone searches for local businesses, and callers rarely spell a business name correctly on the first try.

Two tables model this:

Table	Purpose
residents	Synthetic residents, including intentional near-duplicate entries — the same person, typed twice, with a typo, transposition, or variant spelling the second time
business_names	A flat (<code>business_id</code> , <code>name</code>) lookup extending the businesses table from Chapter 1, used for “did you mean?” search

`residents` includes a `true_duplicate_of` column that a real registry would **not** have — you would not know in advance which rows are duplicates; that is the entire problem fuzzy matching exists to solve. It’s here purely so the exercises can check whether your queries found the right answers. The dataset also includes two pairs that are deliberately *not* marked as duplicates, on purpose — more on those in Exercise 2.

Exercise Goals

By the end of this chapter you will be able to:

- Explain what a trigram is and compute `similarity()` between two strings by hand for a short example.
- Use the `%` operator to find near-matches at a configurable similarity threshold, and explain why no single threshold is perfectly correct.
- Use `word_similarity()` to find a short query as a strong partial match inside a longer string, and explain how it differs from `similarity()`.
- Create a GIN trigram index and confirm it accelerates both `%` and `LIKE '%term%'` queries that would otherwise force a sequential scan.

- Build a “did you mean?” query using a GiST trigram index and `ORDER BY ... <->`, and explain why GiST — not GIN — is the right index for that access pattern.
 - Compare trigram matching against full-text search on the same kind of short, keyword-style query, and know which one to reach for.
-

Installation

`pg_trgm` ships as part of the standard `postgresql-16` package on Debian/Ubuntu — unlike PostGIS in Chapter 2, there is no separate package to install. Enable it in the database:

```
CREATE EXTENSION IF NOT EXISTS pg_trgm;
```

Loading the Data

Prerequisites

Chapter 1’s seed script must have been run first — `business_names` is populated directly from the `businesses` table:

```
python data/ch01_seed.py
```

Run the Chapter 5 seed

```
python data/ch05_seed.py
```

Expected output:

```
Connecting to: dbname=portsmith
Creating schema ...
Inserting 58 residents ...
Populating business_names from businesses ...
Done – 58 rows in residents, 48 rows in business_names.
```

Verify the load

Open `psql portsmith` and run these checks.

Check 1 — table structure:

```
\d residents
```

```

                                Table "public.residents"
      Column      | Type      | Collation | Nullable |
      |            | Default   |           |          |
      +-----+-----+-----+-----+
id               | integer   |           | not null |
nextval('residents_id_seq'::regclass)
full_name        | text      |           | not null |
neighbourhood    | text      |           | not null |
true_duplicate_of | integer   |           |          |
Indexes:
    "residents_pkey" PRIMARY KEY, btree (id)
Foreign-key constraints:
    "residents_true_duplicate_of_fkey" FOREIGN KEY
    (true_duplicate_of) REFERENCES residents(id)
Referenced by:
    TABLE "residents" CONSTRAINT
    "residents_true_duplicate_of_fkey" FOREIGN KEY (true_duplicate_of)
    REFERENCES residents(id)

```

Check 2 — every marked duplicate points at a real canonical row:

```

SELECT COUNT(*) AS duplicate_pairs
FROM   residents
WHERE  true_duplicate_of IS NOT NULL;

duplicate_pairs
-----
12

```

Check 3 — business_names mirrors businesses 1:1:

```

SELECT COUNT(*) FROM business_names;

count
-----
48

```

If all three match, proceed to the exercises.

Exercises

Exercise 1 — What a Trigram Is, and `similarity()`

1.1 — Break a string into trigrams

`show_trgm()` returns the actual set of trigrams PostgreSQL generates for a string:


```
SELECT show_trgm('Eleanor');
```

```
show_trgm
```

```
-----
{" e"," el","ano","ean","ele","lea","nor","or "}
```

Before splitting, PostgreSQL pads the string with two leading spaces and one trailing space — " eleanor " — then takes every overlapping run of three characters: " e", " el", "ele", "lea", "ean", "ano", "nor", "or ". The padding matters: it means the first and last letters of a word each get their own distinguishing trigram (" e" marks “starts with e”; "or " marks “ends with or”), so two strings that only differ at the very start or end still register as different rather than accidentally looking identical in the middle.

padded string: two leading spaces, one trailing space



10 characters, 8 overlapping windows of 3 (10 - 3 + 1)

1.2 — `similarity()`: how much overlap, as a fraction

```
SELECT similarity('Eleanor Whitmore', 'Elenor Whitmore');
```

```
similarity
```

```
-----
0.7368421
```

`similarity()` compares the trigram sets of both strings and returns roughly the fraction that overlap — mostly shared trigrams (both names are 16-17 characters, differing by one dropped letter) means a

high score close to 1. Two unrelated strings — “Eleanor Whitmore” and, say, “Ironside Auto” — share almost no trigrams and score close to 0. The scale is intuitive by construction: 1.0 is identical, 0.0 is nothing alike.

1.3 — All twelve duplicate pairs, scored

```

SELECT a.full_name AS canonical, b.full_name AS duplicate_entry,
       round(similarity(a.full_name, b.full_name)::numeric, 3) AS
sim
FROM   residents a
JOIN   residents b ON b.true_duplicate_of = a.id
ORDER BY a.id;
```

canonical	duplicate_entry	sim
Eleanor Whitmore	Elenor Whitmore	0.737
Jonathan Castellano	Jonathon Castellano	0.739
Priyanka Deshmukh	Priyanka Deshmuk	0.842
Bartholomew Okonkwo	Bartholemew Okonkwo	0.739
Marguerite Delacroix	Marguerite Delacroiux	0.792
Siobhan McAllister	Siobhan MacAllister	0.773
Theodore Vance	Theodor Vance	0.813
Anastasia Volkov	Anastassia Volkov	0.842
Desmond Okafor	Desmund Okafor	0.667
Genevieve Laurent	Genevieve Lorent	0.667
Mikhail Petrenko	Mikail Petrenko	0.737
Fitzgerald Osei	Fitzgerld Osei	0.722

Every genuine duplicate in this dataset scores well above 0.6, regardless of whether the typo was a dropped letter, a swapped letter, or a spelling variant. That range is the working intuition Exercise 2 turns into an actual threshold.

Exercise 2 — The % Operator, and Why One Threshold Isn’t Enough

2.1 — %: “similar enough,” using a session-wide threshold

Computing `similarity()` for every pair by hand doesn’t scale. The `%` operator wraps it into a boolean test against a configurable cutoff:

```

SHOW pg_trgm.similarity_threshold;

pg_trgm.similarity_threshold
-----
0.3
```


0.3 is the default. Self-join residents against itself to find every pair of *different* rows that clears it — this is the actual “find likely duplicates” query:

```

SELECT a.id, a.full_name, b.id, b.full_name,
       round(similarity(a.full_name, b.full_name)::numeric, 3) AS
sim
FROM   residents a
JOIN   residents b ON a.id < b.id
WHERE  a.full_name % b.full_name
ORDER BY sim DESC, a.id;

```

id	full_name	id	full_name	sim
-----+-----		-----+-----		-----+-----
35	Priyanka Deshmukh	36	Priyanka Deshmuk	0.842
45	Anastasia Volkov	46	Anastassia Volkov	0.842
43	Theodore Vance	44	Theodor Vance	0.813
39	Marguerite Delacroix	40	Marguerite Delacroix	0.792
41	Siobhan McAllister	42	Siobhan MacAllister	0.773
57	Nadia Kowalski	58	Nadia Kowalska	0.765
33	Jonathan Castellano	34	Jonathon Castellano	0.739
37	Bartholomew Okonkwo	38	Bartholemew Okonkwo	0.739
31	Eleanor Whitmore	32	Elenor Whitmore	0.737
51	Mikhail Petrenko	52	Mikail Petrenko	0.737
53	Fitzgerald Osei	54	Fitzgerld Osei	0.722
47	Desmond Okafor	48	Desmund Okafor	0.667
49	Genevieve Laurent	50	Genevieve Lorent	0.667
55	Robert Ashworth	56	Bobby Ashworth	0.409

(14 rows)

`a.id < b.id` keeps each pair once instead of twice (A-vs-B and B-vs-A are the same comparison). Notice: **fourteen** rows came back, not twelve. Check `true_duplicate_of` on every row involved and you’ll find twelve of these pairs marked as genuine duplicates and two that are not — “Nadia Kowalski” / “Nadia Kowalska” and “Robert Ashworth” / “Bobby Ashworth”. They’re in this dataset on purpose, and — this is the point of the exercise — look at *where* they landed in the ranking.

2.2 — A false positive sitting in the middle of the true positives

“Nadia Kowalski” and “Nadia Kowalska” are two unrelated Riverside residents who happen to share a first name and a Polish surname that differs only in its masculine/feminine ending. Their similarity, 0.765, isn’t an edge case at the bottom of the list — it sits **between** “Siobhan McAllister”/“Siobhan MacAllister” (0.773) and “Bartholomew Okonkwo”/“Bartholemew Okonkwo” (0.739), two genuine duplicates. `similarity()` is doing exactly what it’s designed to do: these two strings really are that close. Whether that means “probably the same person” is a judgment call the function cannot make on its own — the central

limitation of fuzzy matching. It finds **candidates**, not verified duplicates. A production dedup pipeline treats a % match as “flag for review” or “compare a second field too” (a shared phone number or address), never as “merge automatically.”

2.3 — Prove no threshold fixes it

Try to set a threshold that excludes the Kowalski pair:

```
SET pg_trgm.similarity_threshold = 0.77;
```

Re-run the query from 2.1. “Nadia Kowalski”/“Nadia Kowalska” (0.765) is gone — but so are seven of the twelve genuine duplicates, everything scoring below 0.77. There is no number you can put in that SET statement that keeps all twelve real duplicates and excludes the Kowalski pair, because a real duplicate (“Bartholomew Okonkwo”/“Bartholemew Okonkwo”, 0.739) scores *lower* than the false positive you’re trying to exclude. Put the threshold back before continuing:

```
SET pg_trgm.similarity_threshold = 0.3;
```

2.4 — A different failure: the false negative you’d never even see

“Robert Ashworth” and “Bobby Ashworth” are, in this dataset’s backstory, the same person — entered once with a nickname. Notice they only made the results list at all (0.409) because of the shared surname “Ashworth”; as given names alone, “Robert” and “Bobby” share nothing:

```
SELECT similarity('Robert', 'Bobby');
```

```
similarity
-----
0
```

A resident with a *less* distinctive shared surname and a nicknamed given name would score near zero overall and never appear in a % results list at any threshold — not a borderline case to review, just silently absent. `similarity()` measures string closeness, not identity; it has no way to know “Bobby” is short for “Robert” unless something else tells it, such as a synonym table consulted alongside it.

2.5 — The honest takeaway

Between the Kowalski pair and the Ashworth pair, this dataset has one false positive that no threshold can cleanly exclude without also losing real duplicates, and one true duplicate that scores so low it would

never surface as a candidate in the first place. Tune `pg_trgm.similarity_threshold` to trade recall against precision for your own data — but go in expecting a trade-off, not a number that gets both sides to zero.

Exercise 3 — `word_similarity()` for Partial Matches

`similarity()` compares two whole strings — it penalizes a short query against a long target just for being different lengths, which makes it a poor fit for “does this short input appear as a strong match somewhere inside this longer string?”

3.1 — Compare the two functions on the same query

```
SELECT name, round(similarity('bake', name)::numeric, 3) AS
full_sim
FROM business_names
ORDER BY full_sim DESC, name
LIMIT 3;
```

name	full_sim
River Bend Bakery	0.222
Campus Bike & Sports	0.091
Finch & Sons Barbers	0.091

```
SELECT name, round(word_similarity('bake', name)::numeric, 3) AS
word_sim
FROM business_names
ORDER BY word_sim DESC, name
LIMIT 3;
```

name	word_sim
River Bend Bakery	0.800
Bay Street Electronics	0.400
Finch & Sons Barbers	0.400

Same query, same top result, very different score: 0.222 versus 0.800. `word_similarity()` finds the best-matching *substring extent* inside the target — effectively “if I could crop this longer string down to the part that best matches my query, how similar would that piece be?” — rather than scoring the query against the target’s full length. “bake” against “River Bend **Bak**ery” scores high because “Bake” is a near-perfect fragment match, even though “bake” is a small fraction of the full business name.

3.2 — When to reach for which

Use `similarity()` when you're comparing two values that *should* represent the same whole thing — two spellings of one name, as in Exercises 1 and 2. Use `word_similarity()` when a short query is expected to be a fragment of a longer field — autocomplete-style search-as-you-type, or matching a partial business name a caller remembers.

Exercise 4 — GIN Trigram Indexes

4.1 — Without an index

```
EXPLAIN (ANALYZE, BUFFERS)
SELECT name FROM business_names WHERE name LIKE '%Bakery%';
```

QUERY PLAN

```
-----
Seq Scan on business_names  (cost=0.00..25.88 rows=1 width=32)
(actual time=0.008..0.009 rows=1 loops=1)
  Filter: (name ~~ '%Bakery% '::text)
  Rows Removed by Filter: 47
  Buffers: shared hit=1
```

A plain B-tree index cannot help here — `%Bakery%` has no fixed prefix, so there's nothing for a B-tree to seek to. This is exactly the case `pg_trgm` was built for: it indexes every trigram in every row, so a `LIKE` pattern (itself broken into trigrams) can be looked up directly.

4.2 — Create the index

```
CREATE INDEX idx_business_names_trgm
ON business_names
USING GIN (name gin_trgm_ops);
```

4.3 — Confirm it's used

As in earlier chapters, force the index on a table this small to see the mechanism:

```
SET enable_seqscan = off;

EXPLAIN (ANALYZE, BUFFERS)
SELECT name FROM business_names WHERE name LIKE '%Bakery%';

SET enable_seqscan = on;  -- always restore this
```

QUERY

PLAN

```
-----
Bitmap Heap Scan on business_names  (cost=21.82..25.83 rows=1
```

```

width=32) (actual time=0.024..0.025 rows=1 loops=1)
  Recheck Cond: (name ~~ '%Bakery% '::text)
  Heap Blocks: exact=1
  Buffers: shared hit=10
  -> Bitmap Index Scan on idx_business_names_trgm
(cost=0.00..21.82 rows=1 width=0) (actual time=0.016..0.017 rows=1
loops=1)
    Index Cond: (name ~~ '%Bakery% '::text)
    Buffers: shared hit=9

```

Bitmap Index Scan on idx_business_names_trgm — the same GIN index also accelerates %, word_similarity() %>, and both LIKE and ILIKE with leading wildcards, all from one index.

Exercise 5 — A “Did You Mean?” Query with GiST

5.1 — The pattern: order by trigram distance, take the top few

pg_trgm defines a distance operator, <-> — the complement of similarity (smaller means more alike) — which makes “closest matches first” a plain ORDER BY ... LIMIT:

```

AS sim
SELECT name, round(similarity(name, 'Ironsyde Auto')::numeric, 3)
FROM business_names
ORDER BY name <-> 'Ironsyde Auto', name
LIMIT 5;

```

name	sim
Ironside Auto	0.647
AutoFix Portsmith	0.143
Harbour Inn	0.040
The Art Depot	0.037
Riverside Cinema	0.033

A caller who typed “Ironsyde Auto” (transposed letters) gets “Ironside Auto” back as the clear top match, well clear of the noise below it — this is the entire “did you mean?” feature.

5.2 — Why the GIN index from Exercise 4 doesn’t help here

```

EXPLAIN (ANALYZE)
SELECT name FROM business_names
ORDER BY name <-> 'Ironsyde Auto'
LIMIT 5;

```

QUERY PLAN

```

Limit (cost=2.40..2.41 rows=5 width=36) (actual
time=0.182..0.183 rows=5 loops=1)
  -> Sort (cost=2.40..2.52 rows=48 width=36) (actual
time=0.181..0.182 rows=5 loops=1)
    Sort Key: ((name <-> 'Ironsyde Auto')::text)
    Sort Method: top-N heapsort  Memory: 25kB
    -> Seq Scan on business_names (cost=0.00..1.60 rows=48
width=36) (actual time=0.040..0.157 rows=48 loops=1)

```

Even with `idx_business_names_trgm` in place, this plan is a full scan plus a sort — GIN accelerates *lookups* (“which rows contain trigrams matching this pattern”) but has no notion of a distance ordering. Nearest-neighbor queries need an index that natively understands “closest,” which is what **GiST** provides:

```

CREATE INDEX idx_business_names_trgm_gist
ON business_names
USING GIST (name gist_trgm_ops);

EXPLAIN (ANALYZE)
SELECT name FROM business_names
ORDER BY name <-> 'Ironsyde Auto'
LIMIT 5;

```

QUERY PLAN

```

-----
Limit (cost=0.14..1.07 rows=5 width=36) (actual
time=0.130..0.143 rows=5 loops=1)
  -> Index Scan using idx_business_names_trgm_gist on
business_names (cost=0.14..9.09 rows=48 width=36) (actual
time=0.129..0.141 rows=5 loops=1)
    Order By: (name <-> 'Ironsyde Auto')::text)

```

Order By: `(name <-> 'Ironsyde Auto')::text)` inside an Index Scan — the GiST index walks straight to the nearest rows instead of scoring and sorting every row in the table. **Rule of thumb: GIN for %/LIKE membership lookups, GiST for ORDER BY ... <-> ... LIMIT n nearest-match queries.** It’s common to keep both, as this table now does, if an application needs both access patterns.

5.3 — One more example

```

SELECT name, round(similarity(name, 'Portsmouth Veterinary
Clinic')::numeric, 3) AS sim
FROM business_names
ORDER BY name <-> 'Portsmouth Veterinary Clinic', name
LIMIT 5;

```

name	sim
-----+-----	

Portsmouth Veterinary Clinic	0.833
AutoFix Portsmouth	0.286
Portsmouth Pharmacy	0.286
Portsmouth Tailors	0.286
Portsmouth Arms Hotel	0.263

Exercise 6 — Trigram Matching vs. Full-Text Search, Head to Head

Chapter 4 built full-text search over `city_documents`; this chapter built trigram matching over `business_names`. Both can answer “find rows matching this word” — they are good at it for different, almost opposite, reasons.

6.1 — A spelling variant: trigram wins

Portsmouth’s own documents consistently use the British spelling “harbour.” A caller who types the American spelling “harbor” gets nothing from full-text search — stemming normalizes *inflection* (“harbors” → “harbor”), not *spelling*:

```
SELECT to_tsvector('english', 'Harbour View Theater') @@
to_tsquery('english', 'harbor');

?column?
-----
f
```

Trigram similarity doesn’t care that they’re “different words” — it only sees how many three-letter fragments overlap, and “harbor”/“harbour” share almost all of theirs:

```
SELECT similarity('harbor', 'harbour');

similarity
-----
0.5
```

6.2 — A short, correctly-spelled keyword: full-text search wins

Run a short, exact keyword against `business_names` both ways:

```
SELECT name, round(similarity(name, 'bay')::numeric, 3) AS sim
FROM business_names
ORDER BY sim DESC, name
LIMIT 6;

name | sim
-----+-----
```

Mango Bay Caribbean	0.200
Bay Street Electronics	0.174
River Bend Bakery	0.105
Finch & Sons Barbers	0.095
Bella Napoli	0.063
Le Petit Bistro	0.053

A three-letter query like “bay” barely produces any trigrams of its own, so the scores are all low and close together — “River Bend Bakery” and “Finch & Sons Barbers” show up with real, if small, similarity despite having nothing to do with “bay.” There’s no clean gap between signal and noise. Full-text search, which matches on whole tokens rather than character fragments, has no such problem:

```
SELECT name FROM business_names
WHERE to_tsvector('english', name) @@ plainto_tsquery('english',
'bay')
ORDER BY name;
```

```
name
-----
Bay Street Electronics
Mango Bay Caribbean
(2 rows)
```

Exactly the two relevant matches, no noise, no threshold to tune.

6.3 — The rule of thumb

Short queries made of correctly-spelled whole words belong to full-text search — it was built to match tokens precisely and cheaply via a GIN index over lexemes. Queries that might be misspelled, transposed, or OCR-damaged belong to trigram matching — it was built to tolerate exactly that kind of noise, at the cost of weaker signal on very short inputs. Many real search boxes run both: try full-text search first, and fall back to a trigram “did you mean?” only when it returns nothing.

Summary — What You Should Now Know

Tool	What it does
<code>show_trgm(text)</code>	Shows the padded, overlapping 3-character sequences a string breaks into
<code>similarity(a, b)</code>	Fraction of shared trigrams between two whole strings — best for “are

Tool	What it does
	these two values the same thing, spelled differently?”
<code>word_similarity(a, b)</code>	Best-matching substring extent of a inside b — best for “is a a fragment somewhere inside b?”
<code>a % b</code>	Boolean test: does <code>similarity(a, b)</code> clear <code>pg_trgm.similarity_threshold</code> (default 0.3)?
<code>a <-> b</code>	Trigram distance (<code>1 - similarity</code>) — sort by this for nearest-match-first results
<code>GIN (col gin_trgm_ops)</code>	Accelerates <code>%</code> and <code>LIKE/ILIKE</code> membership-style lookups
<code>GIST (col gist_trgm_ops)</code>	Accelerates <code>ORDER BY col <-> query LIMIT n</code> nearest-neighbor lookups
<code>pg_trgm.similarity_threshold</code>	Session-level cutoff for <code>%</code> ; tune it, but expect trade-offs, not a perfect split

The key design insight from this chapter is that fuzzy matching produces *candidates*, not verdicts. The `%` operator surfaced all twelve real duplicate pairs in the `residents` table — and two pairs of genuinely different, or differently-named, people right alongside them, at every threshold tried. That isn’t a shortcoming to engineer away; it’s the nature of measuring string similarity instead of identity. Build fuzzy matching into a review queue or a secondary confirmation step, not an automatic merge.

The `business_names` table you built here is reused in Chapter 10 (PostgREST), which exposes the “did you mean?” query from Exercise 5 as a public RPC endpoint.

*Going further: `pg_trgm` also provides `%>` and `<->` variants tuned for `word_similarity()` rather than `similarity()`, useful for indexing substring-style “did you mean?” search the same way Exercise 5 indexed whole-string search. For very large tables, GiST trigram indexes are typically smaller and faster to build than GIN but slower for pure membership lookups — benchmark both on your actual data distribution rather than assuming one is strictly better. And if deduplication is a recurring, business-critical process rather than an occasional query, look at dedicated record-linkage tools (e.g. the *dedupe* Python library), which combine trigram-style string similarity with other fields — address,*

phone number, date of birth — and a trained classifier, instead of a single threshold on a single column.

Chapter 6 — Vector Search:

pgvector for Embeddings

“Full-text search finds documents that use your words. Vector search finds documents that share your meaning, whether or not they share a single word with you.”

Background

Chapter 4 matched documents by shared vocabulary. Chapter 5 matched strings by shared spelling. Neither technique can find a document about “the city’s dog park” when a resident searches “somewhere my pet can run off-leash” — there isn’t a misspelling to correct or a stemmed word in common; the two phrases are simply *about the same thing*, expressed in unrelated words. No amount of stemming, stopword tuning, or trigram threshold-adjusting closes that gap, because all three earlier techniques operate on the text itself. Closing it requires a representation of what the text *means*.

That representation is an **embedding**: a machine learning model reads a piece of text and outputs a vector — literally an array of a few hundred floating-point numbers — positioned in a high-dimensional space such that texts with similar meaning end up as nearby points and unrelated texts end up far apart. “Somewhere my pet can run off-leash” and “the city’s dog park” land close together in that space even though they share zero words. This chapter never trains such a model — it uses a small, well-established, freely available one (all-MiniLM-L6-v2) to turn text into 384-number vectors, and PostgreSQL’s pgvector extension to store, index, and search them.

The two hard problems pgvector solves are storage (a native vector type, instead of smuggling floats through a JSONB array or a comma-separated text column) and **search at scale**. Finding the closest vector to a query by brute force means comparing it against every single row — fine for the thirty documents in Chapter 4’s table, hopeless for a million-row photo library. This chapter builds up to two different **approximate nearest neighbor (ANN)** index types, IVFFlat and HNSW, that trade a small amount of accuracy for a large amount of

speed, and it’s honest about exactly how much accuracy each one actually costs you at its defaults, because the answer is more surprising than most introductions let on.

The Scenario

Two tables carry this chapter, and they get their vectors two very different ways, on purpose.

`city_documents` — the council minutes, zoning ordinances, and public notices from Chapter 4 — gets a real embedding column computed by actually running `all-MiniLM-L6-v2` over each document’s title and body. These are genuine embeddings: semantic search over them in Exercise 5 finds real conceptual matches, not a scripted demo.

`city_photos` is different, and says so in its own name: it’s a synthetic stand-in for a photo library’s image embeddings (the kind a real vision model like CLIP would produce for actual photographs), used purely to give Exercises 3 and 4 a table large enough — 5,000 rows across ten categories — for approximate indexing to be worth doing at all. Thirty rows, like `city_documents` has, will never be large enough to make an ANN index pay for itself; a sequential scan over thirty rows is already about as fast as computation gets. There is no real photo behind any row in `city_photos` — each one is a random point clustered around one of ten category “anchors” in 384-dimensional space, which reproduces the *shape* of a real embedding space (similar things cluster; different things don’t) without needing a single actual image.

Table	Vectors	Used for
<code>city_documents</code>	Real — <code>all-MiniLM-L6-v2</code> embeddings of each document’s title + body	Semantic search, hybrid search (Exercises 5-6)
<code>city_photos</code>	Synthetic — 5,000 rows, 10 clustered categories, no real images	ANN indexing at scale (Exercises 3-4)

Exercise Goals

By the end of this chapter you will be able to:

- Store embeddings in a native vector column and understand what the numbers in it actually represent.
 - Compute exact nearest neighbors with $\langle - \rangle$ (L2), $\langle \# \rangle$ (negative inner product), and $\langle = \rangle$ (cosine distance), and know when the choice between them actually changes your results.
 - Build an IVFFlat index, and correctly interpret its `lists` and `probes` parameters instead of trusting the defaults.
 - Build an HNSW index, and make an informed build-time-vs-recall trade-off between it and IVFFlat instead of guessing.
 - Run real semantic search: find documents by meaning even when they share no vocabulary with the query.
 - Build hybrid search that blends a keyword score (`ts_rank`, from Chapter 4) with a semantic score into one ranking.
-

Installation

This chapter has two dependencies, and they're independent of each other — you need both, but you're installing two unrelated things for two unrelated reasons.

1 — pgvector

```
sudo apt install -y postgresql-16-pgvector
```

Enable the extension. **Unlike `pg_trgm` in Chapter 5, `pgvector 0.6.x` is not a “trusted” extension** — its control file doesn't set `trusted = true`, so a regular database-owning role cannot enable it with a plain `CREATE EXTENSION`. It has to be done once, by a superuser:

```
sudo -u postgres psql portsmith -c "CREATE EXTENSION vector;"
```

Confirm it's there:

```
SELECT extversion FROM pg_extension WHERE extname = 'vector';

extversion
-----
0.6.0
```

2 — sentence-transformers, CPU-only

The embedding model runs locally, in Python, via the sentence-transformers library — no API key, no network calls at query time. It depends on PyTorch, and installing PyTorch’s default package pulls in a full NVIDIA CUDA toolkit (several gigabytes) even on a machine with no GPU, which is pure waste for embedding a few dozen short documents on CPU. Install the CPU-only build explicitly, then sentence-transformers on top of it:

```
python3.12 -m venv .venv    # if you haven't already, from earlier chapters
source .venv/bin/activate
pip install torch --index-url https://download.pytorch.org/whl/cpu
pip install sentence-transformers pgvector
```

Note: this installs roughly 3GB into your virtual environment (PyTorch and its numerical dependencies are not small, even CPU-only) and the first time any script loads all-MiniLM-L6-v2, it downloads about 90MB of model weights from Hugging Face and caches them in ~/.cache/huggingface. Every run after that is fully offline. If you see a Warning: You are sending unauthenticated requests to the HF Hub line the first time, that’s expected and harmless — it’s just Hugging Face reminding you that an account would raise your download rate limit.

Loading the Data

Prerequisites

Chapter 4’s seed script must have been run first:

```
python data/ch04_seed.py
```

ch04_seed.py alone is **not** enough, though — this chapter’s hybrid search (Exercise 6) needs city_documents.search_vector, and that column only gets created by hand-running Chapter 4’s **Exercise 3** SQL; it is not part of ch04_seed.py’s schema. (If you re-run ch04_seed.py after already having done Exercise 3 — say, to reset the document data — it drops city_documents and rebuilds it *without* search_vector, silently undoing that exercise.) ch06_seed.py (below) checks for the column and will stop with the exact fix if it’s missing; if you’d rather do it now, here it is verbatim from Chapter 4:

```

ALTER TABLE city_documents ADD COLUMN search_vector tsvector;

UPDATE city_documents
SET    search_vector = to_tsvector('english', title || ' ' ||
body);

CREATE OR REPLACE FUNCTION city_documents_search_vector_update()
RETURNS trigger AS $$
BEGIN
    NEW.search_vector := to_tsvector('english', NEW.title || ' '
|| NEW.body);
    RETURN NEW;
END;
$$ LANGUAGE plpgsql;

CREATE TRIGGER trg_city_documents_search_vector
BEFORE INSERT OR UPDATE OF title, body ON city_documents
FOR EACH ROW
EXECUTE FUNCTION city_documents_search_vector_update();

CREATE INDEX idx_city_documents_search_vector
ON city_documents USING GIN (search_vector);

```

Run the Chapter 6 seed

This adds the (initially empty) embedding column to `city_documents` and creates and populates `city_photos`. It needs only `psycopg2`, `numpy`, and `pgvector` — not `sentence-transformers` — since `city_photos`'s vectors are synthetic:

```
python data/ch06_seed.py
```

Expected output:

```

Connecting to: dbname=portsmith
Applying DDL ...
Generating 5000 synthetic photo embeddings ...
Done — 5000 rows in city_photos across 10 categories.
city_documents.embedding added but left NULL — run `python data/
ch06_embed_documents.py` next.

```

Compute real embeddings for city_documents

This is the step that actually loads and runs the model:

```
python data/ch06_embed_documents.py
```

Expected output:

```

Connecting to: dbname=portsmith
Loading all-MiniLM-L6-v2 (first run downloads ~90MB of model
weights) ...
Embedding 30 documents ...
Done – 30 rows in city_documents now have an embedding.

```

Verify the load

Open psql portsmith and run these checks.

Check 1 — city_photos structure:

```

\d city_photos

Table "public.city_photos"
  Column      |      Type      | Collation | Nullable |
  |            |      Default   |           |          |
  -----+-----+-----+-----+
 id            | integer        |           | not null |
 nextval('city_photos_id_seq'::regclass)
 category      | text           |           | not null |
 neighbourhood | text           |           | not null |
 caption       | text           |           | not null |
 embedding     | vector(384)    |           | not null |
Indexes:
    "city_photos_pkey" PRIMARY KEY, btree (id)

```

Check 2 — 500 photos per category:

```

SELECT category, COUNT(*) FROM city_photos GROUP BY category ORDER
BY category;

```

category	count
community_event	500
harbour_waterfront	500
historic_architecture	500
industrial_dock	500
infrastructure_construction	500
municipal_building	500
public_park	500
residential_street	500
street_market	500
wildlife_nature	500

(10 rows)

Check 3 — every city_documents row has an embedding:

```

SELECT COUNT(*) AS total, COUNT(embedding) AS with_embedding FROM
city_documents;

```


total	with_embedding
30	30

If all three match, proceed to the exercises.

Exercises

Exercise 1 — The vector Type

1.1 — Literals and dimensions

A vector literal is a bracketed list of numbers:

```
SELECT '[1,2,3]':vector;
```

vector
[1,2,3]

A `vector(384)` column, like `city_documents.embedding`, enforces its dimension count at write time — every row’s vector must have exactly 384 numbers. There is no meaningful way to compare a 384-dimensional embedding to a 3-dimensional one, so PostgreSQL simply won’t let a mismatched value in:

```
SELECT '[1,2,3]':vector(384);
```

ERROR: expected 384 dimensions, not 3

1.2 — What’s actually in the column

`city_documents.embedding` isn’t hand-written — it came from `ch06_embed_documents.py` calling `all-MiniLM-L6-v2` on each document’s title and body. Look at a real one:

```
SELECT id, title, embedding FROM city_documents WHERE id = 1;
```

The embedding value prints as 384 comma-separated floating-point numbers between roughly -1 and 1 — not something a human reads directly, but exactly what every operator and index in this chapter operates on. What those particular 384 numbers *mean* isn’t individually interpretable (no single dimension corresponds to a

concept like “is about harbours”); what matters is only their *position relative to other documents’ vectors*, which is what the distance operators in Exercise 2 measure.

Exercise 2 — Exact Nearest Neighbors: <->, <#>, <=>

pgvector gives you three distance operators. They agree more often than you’d expect, and the case where they disagree is worth understanding before it surprises you in production.

2.1 — Three tiny vectors, three operators

```
SELECT '[1,0,0]':vector <-> '[0,1,0]':vector AS l2_distance;
```

```
l2_distance
-----
1.4142135623730951
```

<-> is **Euclidean (L2) distance** — straight-line distance between the two points, exactly like distance in ordinary geometry. Two orthogonal unit vectors are $\sqrt{2}$ apart.

```
SELECT '[1,0,0]':vector <#> '[0,1,0]':vector AS
neg_inner_product;
```

```
neg_inner_product
-----
-0
```

<#> is the **negative inner product** — the dot product of the two vectors, negated (pgvector negates it so that, consistent with the other two operators, *smaller means closer*). Orthogonal vectors have a dot product of zero.

```
SELECT '[1,0,0]':vector <=> '[0,1,0]':vector AS cosine_distance;
```

```
cosine_distance
-----
1
```

<=> is **cosine distance** ($1 - \text{cosine similarity}$) — the angle between the two vectors, completely ignoring their length. Orthogonal vectors have maximum cosine distance, 1.

2.2 — Where they disagree: magnitude


The three operators can rank the *same* pair of candidates in a *different* order the moment the vectors involved aren't the same length. Compare a query vector against two candidates — one that points in exactly the same direction but is twice as long, and one that points in a slightly different direction but is almost the same length:

```
SELECT
  '[1,0]':vector <-> '[2,0]':vector AS l2_to_a,
  '[1,0]':vector <-> '[1,0.5]':vector AS l2_to_b,
  round(('[1,0]':vector <=> '[2,0]':vector)::numeric, 4) AS
cos_to_a,
  round(('[1,0]':vector <=> '[1,0.5]':vector)::numeric, 4) AS
cos_to_b;

 l2_to_a | l2_to_b | cos_to_a | cos_to_b
-----+-----+-----+-----
      1 |    0.5 | 0.0000 | 0.1056
```

The ranking flips. By $\leftarrow$ (L2), candidate B ([1,0.5], distance 0.5) is closer than candidate A ([2,0], distance 1). By $\leftarrow$ (cosine), it's the reverse — A is distance 0 (identical direction) while B is 0.1056 away. Neither operator is “wrong”; they're answering different questions. L2 asks “how far apart are these points.” Cosine asks “how similarly are these vectors oriented, regardless of scale.”

L2 distance says B is closer. Cosine distance says A is closer.



A sits exactly on the query’s own line (the overlap in the diagram above is “same direction, twice the length”), which is why its cosine distance is zero. The dashed circle is centered on the query with radius equal to B’s L2 distance — B sits right on that boundary, while A sits well outside it, which is the straight-line sense in which B is “closer.”

2.3 — Why this chapter normalizes, and what that buys you

`ch06_embed_documents.py` calls `model.encode(texts, normalize_embeddings=True)` — every stored embedding is scaled to unit length before it’s written. Once every vector in a comparison has the same length, the magnitude difference that caused Exercise 2.2’s disagreement simply can’t occur, and all three operators produce the *same ranking* (just on different numeric scales). Embedding the query “waterfront redevelopment funding” (Exercise 5 shows how) and checking the top 3 `city_documents` rows by each operator confirms it:

	<-> (L2)	<#> (neg. inner product)	<=> (cosine)
Rank 1	doc 1 — 0.8605	doc 1 — -0.6298	doc 1 — 0.3702
Rank 2	doc 22 — 1.0546	doc 22 — -0.4439	doc 22 — 0.5561
Rank 3	doc 5 — 1.0646	doc 5 — -0.4333	doc 5 — 0.5667

Same three documents, same order, every time — only the numbers’ scale differs between columns. The practical rule: **normalize your embeddings, and pick whichever operator has an index you can build** — for normalized vectors it stops being a search-quality decision and becomes a performance one, covered next.

Exercise 3 — IVFFlat: Approximate Search with lists and probes

3.1 — Exact search doesn’t scale

`city_photos` has 5,000 rows — small by real standards, but already enough to feel the cost of brute-force search. Find the 10 nearest neighbors of photo `id = 1` by comparing it against every other row:

```
EXPLAIN (ANALYZE, BUFFERS)
SELECT id, category, embedding <-> (SELECT embedding FROM
city_photos WHERE id = 1) AS dist
```

```

FROM city_photos
WHERE id != 1
ORDER BY embedding <-> (SELECT embedding FROM city_photos WHERE
id = 1)
LIMIT 10;

```

QUERY PLAN

```

-----
Limit (cost=1439.32..1439.35 rows=10 width=30) (actual
time=4.164..4.167 rows=10 loops=1)
  -> Sort (cost=1431.02..1443.52 rows=4999 width=30) (actual
time=4.163..4.164 rows=10 loops=1)
    Sort Key: ((city_photos.embedding <-> $0))
    Sort Method: top-N heapsort  Memory: 26kB
    -> Seq Scan on city_photos (cost=0.00..1323.00
rows=4999 width=30) (actual time=0.024..3.381 rows=4999 loops=1)

```

Every one of the 4,999 other rows gets its distance computed and sorted — about 4.6ms here. That cost grows linearly with table size; at real photo library scale (millions of rows) it stops being viable.

3.2 — Build the index

```

CREATE INDEX idx_city_photos_ivfflat
ON city_photos
USING ivfflat (embedding vector_l2_ops)
WITH (lists = 50);

```

```

CREATE INDEX
Time: 152.129 ms

```

`lists = 50` partitions the 5,000 vectors into 50 clusters (via k-means at index-build time) and stores each vector under its nearest cluster centroid. A rule of thumb from the pgvector docs: `lists ≈ rows / 1000` for up to roughly a million rows.

3.3 — The planner picks it up on its own

```

EXPLAIN (ANALYZE, BUFFERS)
SELECT id, category, embedding <-> (SELECT embedding FROM
city_photos WHERE id = 1) AS dist
FROM city_photos
WHERE id != 1
ORDER BY embedding <-> (SELECT embedding FROM city_photos WHERE
id = 1)
LIMIT 10;

```

QUERY

PLAN

```

-----
Limit (cost=61.55..62.35 rows=10 width=30) (actual

```

```
time=0.145..0.168 rows=10 loops=1)
-> Index Scan using idx_city_photos_ivfflat on city_photos
(cost=53.25..455.00 rows=4999 width=30) (actual time=0.145..0.166
rows=10 loops=1)
    Order By: (embedding <-> $0)
```

No `enable_seqscan = off` trick needed this time, unlike the small tables in earlier chapters — at 5,000 rows the planner’s own cost estimate already favors the index: **4.6ms down to about 0.2ms.**

3.4 — The catch: default probes gives up a lot of accuracy

An IVFFlat search only examines the clusters nearest the query vector, not all 50 — controlled by `ivfflat.probes`, which **defaults to 1**. Compare the approximate top-10 against the true exact top-10:

```
-- Exact (force a sequential scan to get true ground truth):
SET enable_indexscan = off;
SELECT array_agg(id) FROM (
    SELECT id FROM city_photos WHERE id != 1
    ORDER BY embedding <-> (SELECT embedding FROM city_photos
WHERE id = 1) LIMIT 10
) t;

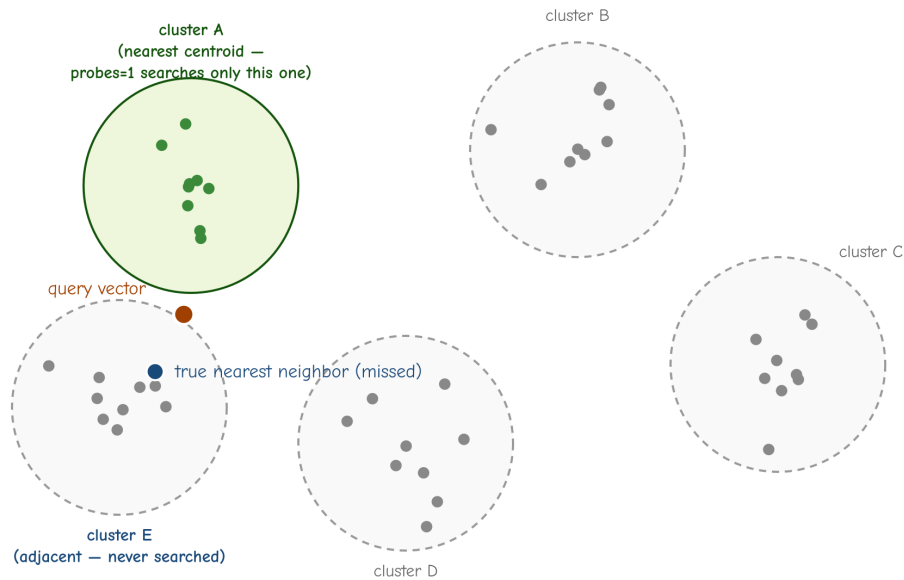
          array_agg
-----
{146,357,485,468,500,261,384,397,185,226}

-- Approximate, default probes = 1:
RESET enable_indexscan;
SELECT array_agg(id) FROM (
    SELECT id FROM city_photos WHERE id != 1
    ORDER BY embedding <-> (SELECT embedding FROM city_photos
WHERE id = 1) LIMIT 10
) t;

          array_agg
-----
{485,261,185,226,385,22,274,433,56,327}
```

Only **4 of the 10** results match (485, 261, 185, 226) — **40% recall**, at the default setting, on a real query. This is not a contrived worst case; it’s what `ivfflat.probes = 1` actually does: it searches only the single cluster closest to the query and simply never looks at the other 49, even though some of the true nearest neighbors live in a different cluster.

probes=1 only searches the nearest cluster; even when a true neighbor sits just across the boundary



IVFFlat clusters vectors at index-build time. probes=1 checks only the nearest centroid (A) — a real neighbor one cluster over is invisible to the search.

A note on reproducing these exact numbers: IVFFlat’s clustering step uses k-means with a random initialization at *index build* time — even against `city_photos`’s fixed, seeded synthetic data, rebuilding the same `CREATE INDEX ... USING ivfflat` statement can shift which photos land in which of the 50 clusters, and therefore which specific IDs and exact recall percentage `probes = 1` produces. Rerunning this exercise on your own machine will very likely show different IDs and a different percentage than the ones printed here. The pattern that reproduces reliably is the *shape*: `probes = 1` is measurably, often substantially, incomplete.

3.5 — Raise probes, watch recall recover

```
SET ivfflat.probes = 2;
```

```
array_agg
```

```
-----  
{146,357,485,468,500,261,384,397,185,226}
```

At `probes = 2` the result already matches the exact set exactly — **100% recall** on this query. Timing at a slightly more conservative `probes = 5`, and at `probes = 50` (every list — the most exhaustive an IVFFlat search can be):

probes	Recall (this query)	Query time
1 (default)	40%	~0.4ms
5	100%	~0.6ms
50 (all lists)	100%	~0.6ms

Timings vary run to run more than the recall figures do (expect noise in the tenths of a millisecond on a table this size) — the shape that matters is that probes = 1 is consistently fastest *and* the least reliable, while anything from probes = 5 up delivers full recall on this query without a further, proportional cost from probing still more lists; the IVFFlat index structure keeps even an “examine every list” search meaningfully faster than the 4.6ms plain sequential scan from Exercise 3.1. Reset before continuing:

```
SET ivfflat.probes = 1;
```

Two probes happened to be enough for this particular query; it isn’t a guarantee for every query, which is why a small safety margin (a starting point of probes $\approx \sqrt{\text{lists}}$, so $\sqrt{50} \approx 7$) is the honest recommendation rather than the bare minimum that worked once. The overall lesson: **IVFFlat’s out-of-the-box defaults are fast and unreliable at the same time** — lists and probes are not optional tuning, they’re the whole feature.

Exercise 4 — HNSW: A Different Trade-off

4.1 — Build it, and time the build

```
CREATE INDEX idx_city_photos_hnsw
ON city_photos
USING hnsw (embedding vector_l2_ops);
```

```
CREATE INDEX
Time: 513.963 ms
```

Compare against IVFFlat’s build time from Exercise 3.2: **152ms**. HNSW (Hierarchical Navigable Small World) builds a multi-layer graph structure connecting each vector to its approximate neighbors, which costs roughly **3x longer to build** than IVFFlat’s k-means clustering, on this table.

4.2 — Recall at HNSW’s default settings

Drop (or temporarily disable) the IVFFlat index so the planner has no choice but to use HNSW, then repeat the exact same top-10 query from Exercise 3.4:

```
BEGIN;
DROP INDEX idx_city_photos_ivfflat;

SELECT array_agg(id) FROM (
    SELECT id FROM city_photos WHERE id != 1
    ORDER BY embedding <-> (SELECT embedding FROM city_photos
WHERE id = 1) LIMIT 10
) t;

ROLLBACK;  -- put idx_city_photos_ivfflat back

array_agg
-----
{146,357,485,468,500,261,384,397,185,226}
```

Exact match with the true top-10 — 100% recall, at HNSW’s default settings, no tuning required. Query time for this was about 0.63ms. Repeating the comparison against a second, independent query point (`id = 1001`, a `public_park` photo) confirms it isn’t a fluke: HNSW again returned the exact top-10, while IVFFlat at default `probes = 1` again recovered only 3 of 10 (30% recall — as Exercise 3.4 noted, the exact percentage shifts between index rebuilds; what doesn’t shift is that it’s consistently well short of complete). HNSW’s default, across every test in this chapter, has not missed a single true neighbor.

4.3 — The trade-off, stated plainly

	IVFFlat	HNSW
Build time (5,000 rows)	152ms	514ms (~3x)
Recall at defaults	30-40% in this chapter’s tests, varies by build (needs probes tuned up)	~100% (no tuning needed)
Query time (well-tuned)	Comparable	Comparable
Tuning burden	On you, every session (probes is a session-level SET)	Minimal

This matches the general guidance in the pgvector documentation: **HNSW is the better default choice for most workloads** — better recall out of the box, no per-session parameter to remember to set. IVFFlat’s advantages are real but narrower: faster to build (relevant if you rebuild the index often, e.g. after large bulk loads) and somewhat lower memory overhead at very large scale. Choose IVFFlat deliberately for those reasons, not by default.

Exercise 5 — Semantic Search: Finding Meaning, Not Words

5.1 — A query with no vocabulary overlap

`data/ch06_semantic_search.py` embeds a query string with the same model used to build the column, and searches `city_documents` by cosine similarity:

```
python data/ch06_semantic_search.py "stray dogs and pet-friendly spaces"
```

id	sim	title
6	0.4447	Council Minutes – Riverside Dog Park Funding
24	0.3927	Public Notice – Riverside Dog Park Ribbon-Cutting Event
16	0.2683	Zoning Ordinance – Riverside Short-Term Rental Restrictions
...		

Not one word of the query — “stray,” “dogs,” “pet-friendly,” “spaces” — appears in document 6’s title, and full-text search (Chapter 4) would return zero rows for it; there’s no shared lexeme to match on, misspelled or otherwise, so Chapter 5’s trigram matching wouldn’t help either. The embedding model has no notion of exact words at all — it read “the city’s dog park” and “somewhere my pet can run off-leash” as *about the same thing*, which is precisely the gap described in this chapter’s Background.

5.2 — A stronger example

```
python data/ch06_semantic_search.py "waterfront redevelopment funding"
```

id	sim	title
1	0.6298	Council Minutes – Harbour District Waterfront Renovation Budget
22	0.4439	Public Notice – Public Hearing on Harbour District Waterfront Rezoning

```

5    0.4333  Council Minutes – Riverside Road Resurfacing
Program
12   0.4207  Zoning Ordinance – Harbour District Waterfront
Height Variance
9    0.4101  Council Minutes – Canal Road Bike Lane Expansion
...

```

5.3 — Reading the gap, not just the ranking

Look at the jump between rank 1 (0.6298) and rank 2 (0.4439) — a much bigger drop than between ranks 2 through 5. That gap is meaningful: it says document 1 isn’t merely the *best available* match, it’s a *genuinely strong* one, while the rest are progressively looser associations. Semantic search has a property full-text search doesn’t: **it never returns zero results**. Every document has *some* cosine similarity to any query — the model will confidently rank all thirty documents even for a nonsense query, it just won’t find any of them close. A real application needs a similarity-score cutoff (or a “confidence” gap check like this one) to distinguish “found it” from “here’s our least-bad guess,” a problem keyword search’s honest zero-rows sidesteps entirely.

Exercise 6 — Hybrid Search: Blending Keyword and Semantic Signals

Semantic search and keyword search fail differently. Semantic search misses exact terms that matter — a specific ordinance number, a street name — treating them as just more words to average into a general impression of meaning. Keyword search misses paraphrases entirely, as Exercise 5 just showed. Hybrid search runs both and blends the scores.

6.1 — Keyword-only ranking

```

SELECT id, title, round(ts_rank(search_vector,
plainto_tsquery('english','bike lane'))::numeric,4) AS kw_rank
FROM   city_documents
WHERE  search_vector @@ plainto_tsquery('english','bike lane')
ORDER BY kw_rank DESC;

```

```

id |
title                                | kw_rank
----+-----
9 | Council Minutes – Canal Road Bike Lane
Expansion                            | 0.2928

```

```

28 | Public Notice – Canal Road Bike Lane Construction
Schedule | 0.1960
30 | Public Notice – Annual Fireworks Display and Street
Closures | 0.1759
5 | Council Minutes – Riverside Road Resurfacing
Program | 0.0992
15 | Zoning Ordinance – Reduced Parking Minimums Near Transit
Corridors | 0.0991

```

6.2 — Semantic-only ranking, same query

```

python data/ch06_semantic_search.py "bike lane" --top 8

id      sim  title
9      0.5513 Council Minutes – Canal Road Bike Lane Expansion
28     0.4433 Public Notice – Canal Road Bike Lane Construction
Schedule
15     0.4117 Zoning Ordinance – Reduced Parking Minimums Near
Transit Corridors
5      0.3965 Council Minutes – Riverside Road Resurfacing
Program
21     0.2754 Public Notice – Bay Street Water Main Repair Road
Closure
6      0.2688 Council Minutes – Riverside Dog Park Funding
30     0.2653 Public Notice – Annual Fireworks Display and Street
Closures
10     0.2359 Council Minutes – Special Session on Flood
Mitigation Infrastructure

```

Both agree on the top 2 (documents 9 and 28 — a real, strong signal either way finds them). They disagree further down: document 30 ranks **3rd** by keyword (it happens to use the phrase “biking along the Canal Road bike lane” once, in passing) but only **7th** by semantic similarity, while document 15 ranks **5th** by keyword but **3rd** semantically (it’s substantively about bike-lane-adjacent parking policy, using different phrasing throughout).

6.3 — Blend both into one score

```

python data/ch06_semantic_search.py "bike lane" --hybrid --top 8

id      kw      sem  hybrid  title
9      0.2928  0.5513  0.4479  Council Minutes – Canal Road Bike
Lane Expansion
28     0.1960  0.4433  0.3444  Public Notice – Canal Road Bike
Lane Construction Schedule
15     0.0991  0.4117  0.2867  Zoning Ordinance – Reduced
Parking Minimums Near Transit Corridors
5      0.0992  0.3965  0.2776  Council Minutes – Riverside Road
Resurfacing Program
30     0.1759  0.2653  0.2295  Public Notice – Annual Fireworks

```

Display and Street Closures

21	0.0000	0.2754	0.1652	Public Notice – Bay Street Water Main Repair Road Closure
6	0.0000	0.2688	0.1613	Council Minutes – Riverside Dog Park Funding
10	0.0000	0.2359	0.1415	Council Minutes – Special Session on Flood Mitigation Infrastructure

The underlying SQL (HYBRID_SQL in `ch06_semantic_search.py`) is a LEFT JOIN between a keyword-matched candidate set and a semantic top-10 candidate set, so documents 21, 6, and 10 — never matched by the keyword query at all (`kw = 0.0000`) — still appear, carried in purely by semantic similarity:

```
hybrid_score = 0.4 * COALESCE(kw_score, 0) + 0.6 * COALESCE(sem_score, 0)
```

Document 15 moves from 5th (keyword) up to 3rd (hybrid); document 30 moves from 3rd (keyword) down to 5th (hybrid) — the blend genuinely changes the ranking, not just the score.

6.4 — The honest caveat about the weights

0.4 and 0.6 here are fixed constants chosen because `ts_rank` and cosine similarity happen to land in roughly comparable numeric ranges on *this* dataset — they are not principled, and they are not portable to a different corpus or a different embedding model without re-checking that assumption. A production hybrid-search system typically normalizes each signal (min-max or z-score, over the actual candidate set returned) before blending, rather than trusting that two arbitrary scoring functions happen to land in the same numeric neighborhood. Treat the weights here as a teaching simplification, not a formula to copy into production untested.

Summary — What You Should Now Know

Tool	What it does
<code>vector(N)</code>	A fixed-dimension floating-point array column type; enforces N at write time
<code><-></code>	L2 (Euclidean) distance — straight-line distance between points
<code><#></code>	Negative inner product — dot product, negated so smaller means closer

Tool	What it does
<code><=></code>	Cosine distance — angle between vectors, ignoring magnitude
<code>normalize_embeddings=True</code>	Makes all three operators agree on ranking; without it, they can disagree
<code>USING ivfflat (col vector_l2_ops) WITH (lists = N)</code>	Clusters vectors into N buckets; fast to build, needs probes tuned up from its default to get usable recall
<code>SET ivfflat.probes = N</code>	How many clusters an IVFFlat search actually examines — defaults to 1, often too low
<code>USING hnsw (col vector_l2_ops)</code>	Graph-based ANN index; ~3x slower to build than IVFFlat here, but near-exact recall by default
<code>ts_rank(...) + (1 - cosine_distance) weighted blend</code>	Hybrid search — keyword precision plus semantic recall in one ranking

The key design insight from this chapter is that “approximate” in “approximate nearest neighbor” is not a rounding error — at IVFFlat’s default settings, this chapter repeatedly measured 30-40% recall on real queries, which means 6-7 of the top 10 *true* nearest neighbors were silently missing from the results, with no error, no warning, nothing to indicate anything was wrong. Every ANN index trades recall for speed on a dial you control (probes, ef_search); the failure mode to avoid is not knowing where that dial is set.

The `city_documents` table’s embedding column is reused in Chapter 10 (PostgREST), which exposes the semantic search query from Exercise 5 as a public RPC endpoint alongside the fuzzy “did you mean?” search from Chapter 5.

Going further: this chapter used all-MiniLM-L6-v2 (384 dimensions) because it’s small, fast on CPU, and good enough to teach with — production systems often use larger models (OpenAI’s text-embedding-3-small at 1536 dimensions, or similar) for better semantic quality, at proportionally higher storage and compute cost per vector. For tables much larger than city_photos’s 5,000 rows, revisit lists (scale roughly with row count) and consider hnsw’s m and ef_construction build-time parameters, which trade index size and build time for recall the same way probes trades query time for it. And if your embeddings change

meaning over time — a newer, better model version — remember that a vector column has no built-in versioning; re-embedding the whole table and rebuilding the index is a full migration, not an in-place update.

Bonus Section — Build a Local RAG System

*This section is optional and sits outside the numbered exercises — nothing later in the book depends on it. It exists because semantic search (Exercise 5) is one step short of the thing most people actually mean when they say “AI search”: a system that doesn’t just return matching documents, but reads them and answers the question directly. That’s **Retrieval-Augmented Generation (RAG)**, and you already built the hard half of it — retrieval — in Exercises 1 through 5. This section adds the other half.*

What RAG actually is

Every large language model has a knowledge cutoff and no idea what’s in your `city_documents` table. Ask one directly “what are the rules for Portsmouth’s dog park?” and it will either say it doesn’t know or, worse, confidently invent something plausible-sounding and wrong — a **hallucination**. RAG sidesteps this without retraining or fine-tuning anything: before asking the model the question, retrieve the most relevant real text from your own database and paste it directly into the prompt, instructing the model to answer *only* from that text. The model isn’t recalling facts from training anymore; it’s reading comprehension over text you handed it seconds ago.

This means RAG is two independent systems wired together, and it’s worth keeping them mentally separate:

1. **Retrieval** — PostgreSQL, pgvector, and everything from Exercises 1-5 of this chapter. Given a question, find the most relevant chunks of text.
2. **Generation** — a large language model that turns “here are some relevant facts, here is a question” into a fluent, direct answer.

Nothing about retrieval changes when you add generation on top of it — the `<=>` cosine-distance query from Exercise 5 is exactly what runs here too. What’s new is chunking (documents are usually too long and too topically mixed to embed as a single vector) and the LLM call itself, which this section runs locally through **Ollama** rather than a paid API, for the same reason Chapter 6 as a whole runs its embedding model locally: no API key, no per-query cost, nothing leaves your machine.

Installation — Ollama

```
curl -fsSL https://ollama.com/install.sh | sh
```

This installs Ollama as a systemd service listening on localhost:11434. Pull a chat model — this section uses llama3.1:8b (about 4.9GB):

```
ollama pull llama3.1:8b
```

If disk space or RAM is tight, `ollama pull phi` gets you a much smaller (~1.6GB) model at a real quality cost; every example below works with either, since the model name is a plain argument to the RAG script, not something hardcoded.

Note: the *first* request to any given model incurs a one-time load delay (10-15 seconds isn't unusual) while Ollama loads its weights into memory. Every request after that, to the same model, is fast — Ollama keeps recently used models resident for a while. Don't judge a model's real latency by its very first response.

Ingest: chunk and embed a directory of documents

`data/rag_docs/` contains four new, longer, informational documents — a business licensing guide, a parks and recreation guide, a utilities guide, and a newcomer's guide — distinct from `city_documents`'s formal council records, written in the plain-prose register an actual city help-desk handout would use. `ch06_rag_ingest.py` splits each file into overlapping word-count chunks, embeds every chunk with `all-MiniLM-L6-v2` (the same model used everywhere else in this chapter — retrieval only works if the same model embeds both the documents and the questions), and loads them into a table you name on the command line:

```
python data/ch06_rag_ingest.py data/rag_docs --table  
portsmith_rag --recreate
```

```
Connecting to: dbname=portsmith
```

```
Loading all-MiniLM-L6-v2 ...
```

```
getting_a_business_license.txt: 4 chunks
```

```
newcomer_guide.txt: 4 chunks
```

```
parks_and_recreation.txt: 4 chunks
```

```
utilities_and_public_works.txt: 4 chunks
```

```
Done – 16 chunks from 4 files loaded into 'portsmith_rag'.
```


Each ~700-word document became 4 chunks of 180 words with a 40-word overlap between consecutive chunks (both configurable via `--chunk-size` and `--chunk-overlap`). The overlap matters: without it, a sentence that happens to fall right at a chunk boundary gets split across two chunks, and neither half alone may retrieve well for a question about it.

The table it creates is deliberately the same generic shape regardless of what you point it at:

```
\d portsmith_rag
```

Column	Type	Collation	Nullable
	Default		
id	integer		not null
nextval('portsmith_rag_id_seq'::regclass)			
source	text		not null
chunk_index	integer		not null
content	text		not null
embedding	vector(384)		not null

Indexes:

- "portsmith_rag_pkey" PRIMARY KEY, btree (id)
- "idx_portsmith_rag_embedding" hns w (embedding
- vector_cosine_ops)

`--table` is a plain identifier you choose — point the same script at a different directory with `--table` set to something else, and you have a second, independent knowledge base in the same database. Run it against your own notes, a project’s documentation, anything in `.txt` or `.md` files.

Why the table name is validated, not just interpolated: `--`

`table` comes from the command line and flows into `CREATE TABLE`, `INSERT`, and `CREATE INDEX` statements. SQL identifiers (table names) can’t be passed as query parameters the way values can — `cur.execute("... WHERE id = %s", (table_name,))` only works for *values*. `ch06_rag_ingest.py` and `ch06_rag_chat.py` both reject anything that isn’t a plain lowercase identifier before it touches SQL, and use `psycopg.sql.Identifier` to quote it correctly rather than dropping it into an f-string. For a script you run yourself against your own database this is a small concern; the habit is worth keeping anyway, because the same shortcut in a script that takes a table name from a web form is a real SQL injection hole.

Ask a question

ch06_rag_chat.py takes three required arguments — model, table, question, in that order — plus --host and --port for Ollama (both default to localhost and 11434, so you only need them if Ollama runs elsewhere):

```
python data/ch06_rag_chat.py llama3.1:8b portsmith_rag \
    "How do I get a business license in Portsmouth?" --show-context

--- retrieved context ---
[getting_a_business_license.txt#0 dist=0.2743] Getting a Business
License in Portsmouth Anyone opening a business within city limits
needs a business...
[getting_a_business_license.txt#3 dist=0.5796] to change their
registered category – a retail shop adding a small cafe counter,
for example – need ...
[getting_a_business_license.txt#2 dist=0.6149] scratch rather
than simply renewing, which means a new round of inspections for
categories that require...
[newcomer_guide.txt#0 dist=0.6176] A Newcomer's Guide to
Portsmouth Welcome to Portsmouth. This guide covers the basics every
new resident...
```

To get a business license in Portsmouth, you need to start online through the city's permitting portal, where you choose a business category since it determines which inspections and follow-up permits apply. Most applications are processed within two to three weeks, but those involving food service or alcohol take longer because they require scheduled in-person inspections.

--show-context prints exactly what got retrieved before the model ever saw the question — the same top-K cosine-distance query from Exercise 5, just with content instead of title as the payload. The answer isn't copied from any single chunk; it's the model synthesizing across the top few, which is the actual value RAG adds over raw semantic search: Exercise 5 would have handed a person four ranked chunks to read themselves, this hands them one sentence that already did the reading.

Reproducibility, honestly: the retrieved chunks and their dist values above are fully deterministic — rerun the exact query and you'll get the exact same four chunks in the exact same order, because embedding and cosine distance are pure math. The generated *answer* is not — Ollama samples from the model's output distribution, so wording will differ run to run (sometimes

prose, sometimes a bulleted list) even against identical context. That's a property of the generation half of RAG, not the retrieval half; don't expect to reproduce this exact paragraph verbatim.

What happens when the answer isn't in the documents

This is the case that actually matters for trusting a RAG system — ask something the ingested documents have no way of answering:

```
python data/ch06_rag_chat.py llama3.1:8b portsmouth_rag \
    "What is the capital of France?" --show-context

--- retrieved context ---
[newcomer_guide.txt#1 dist=0.9228] river on the city's western
edge and is primarily residential, with the new dog park, several
parks ...
[utilities_and_public_works.txt#1 dist=0.9272] requires
temporarily shutting off service to a section of a neighbourhood,
Public Works issues a boi...
[newcomer_guide.txt#0 dist=0.9361] A Newcomer's Guide to
Portsmouth Welcome to Portsmouth. This guide covers the basics every
new residen...
[newcomer_guide.txt#2 dist=0.9475] alike. Noise ordinance
enforcement is more actively watched here than in other
neighbourhoods given ...
```

I don't know. The context only provides information about the city of Portsmouth, its neighbourhoods, and various services offered by the city, but does not mention the capital of France.

Two things to notice. First, retrieval **still returned four chunks** — exactly as Exercise 5 described, cosine similarity search never returns “no results,” it returns the *least bad* matches available, and here they're genuinely irrelevant. Second, look at the distances: 0.92-0.95, versus 0.27-0.61 for the business license question. That gap is the signal a production system would actually act on — this script's prompt template happens to lean on the model's own judgment to say “I don't know,” but a more defensive design would check the top distance against a threshold *before* even calling the LLM, and skip generation entirely for a question this poorly matched. Relying on the model to notice the context is irrelevant works here because llama3.1:8b is reasonably well-behaved about it — it is not a guarantee every model or every prompt will get right, and it's the single biggest reliability risk in any RAG system: nothing stops a language model from answering confidently from weak or irrelevant context if the prompt doesn't insist otherwise.

The prompt: the actual glue between the two systems

`ask_ollama()` in `ch06_rag_chat.py` sends the model one plain-text prompt via Ollama's `/api/generate` endpoint — no special “RAG mode” exists in Ollama or in the model itself, it's just a string:

You are a help-desk assistant for the city of Portsmouth. Answer the question using ONLY the context below. If the context doesn't contain the answer, say you don't know rather than guessing.


Context:

```
(getting_a_business_license.txt) Getting a Business License in Portsmouth Anyone opening a business ...  
(getting_a_business_license.txt) to change their registered category – a retail shop adding a small ...  
...
```

Question: How do I get a business license in Portsmouth?

Answer:

This is the entire mechanism. There is no fine-tuning, no special API, nothing model-specific — “retrieval-augmented generation” is a prompt template plus a database query, which is exactly why it was worth showing you the whole thing at this level rather than reaching for a framework. The whole chapter, end to end:



*Going further: this section's retrieval is intentionally the simplest version that works — top-K by raw cosine distance, no re-ranking, no distance-threshold cutoff before generation, no limit on how many tokens of context get sent to the model regardless of length. Production RAG systems typically add a **re-ranking** step (retrieve a wider candidate set cheaply, then re-score the top candidates with a slower, more accurate model before picking the final context), **hybrid retrieval** (blend `ts_rank` back in, exactly as Exercise 6 did, since keyword precision and semantic recall are just as complementary here as they were there), and explicit **citations** in the generated answer so a user can verify a claim against the source chunk rather than trusting the model's summary outright. None of that changes the shape of what you built here — it's still retrieval, then generation, with more care taken at each step.*

Chapter 7 — IP and Network

Filtering: `ip4r`

“An IP address is a number. A CIDR block is a range. PostgreSQL’s built-in types treat the first well and the second as an afterthought — `ip4r` treats the range as the whole point.”

Background

PostgreSQL already ships `inet` and `cidr` — you don’t strictly need an extension to store an IP address. So why does one exist? Because `inet` and `cidr` were designed for storing addresses and networks, not for answering the question a security team actually asks all day: “*is this address inside any of these ranges?*” — over and over, at high volume, fast. `cidr` is also stricter than it looks: it rejects a value with any host bits set relative to its mask (192.168.1.5/24 is refused outright — more on that in Exercise 1), and neither built-in type ships a purpose-built index structure for “find every range containing this point” or “find every point inside this range” at scale. Every containment check against a large blocklist falls back to a sequential scan unless you reach for something else.

`ip4r` is that something else: a dedicated `ip4` type for single IPv4 addresses, an `ip4r` type for arbitrary IPv4 ranges (CIDR-aligned or not — also unlike `cidr`), a GiST opclass built specifically for fast range containment, and — despite the name — full IPv6 equivalents (`ip6`, `ip6r`) plus polymorphic `ipaddress`/`iprange` types that work with either family. This chapter builds a small network security monitoring setup around it: logging access attempts, maintaining block and allow lists, indexing them properly, and — because real blocklists and allowlists eventually disagree with each other — finding and resolving the conflicts that creates.

The Scenario

Portsmouth's online services — the resident portal, the business licensing API from Chapter 5, the permit dashboard — all sit behind the same logging and access-control layer. Every login attempt and API call gets recorded, and a security team maintains two lists: a **blocklist** of ranges observed doing something hostile (brute-force login attempts, known Tor exit nodes, botnet activity), and an **allowlist** of ranges that should never be blocked regardless — the city's own VPN, a trusted vendor's office network.

Table	Purpose
<code>network_events</code>	One row per login attempt or API call — type, source IP, timestamp
<code>blocklists</code>	CIDR ranges flagged as malicious, with a category and description
<code>allowlists</code>	CIDR ranges that should never be blocked, regardless of blocklist entries

Every range in this chapter's data uses IANA-reserved, non-routable address space (the RFC 5737 documentation blocks `192.0.2.0/24`, `198.51.100.0/24`, `203.0.113.0/24`, and the reserved `240.0.0.0/4` block) — synthetic security data, not a real blocklist, using addresses that were never going to belong to anyone.

Exercise Goals

By the end of this chapter you will be able to:

- Explain what `ip4r` gets you over the built-in `inet/cidr` types, in storage and in what values each type actually accepts.
- Use the containment operators `>>` and `<=<` to check whether an address falls inside a range, and handle the case where it matches more than one.
- Build a GiST index for network containment queries — and know why it can't go directly on an `ip4` column, only an `ip4r` one.
- Aggregate events by subnet using both `ip4r`-native functions and the built-in `network()/masklen()` functions via a cast.
- Detect real conflicts between an allowlist and a blocklist using the overlap operator `&&`.

- Build a GiST-indexed function that resolves those conflicts and answers “is this IP blocked?” in one call.
 - *(Additional)* Implement a token-bucket rate limiter keyed on ip4r, applied at both the host and /24 netblock level, and explain why distributed abuse needs the second bucket to be caught at all.
-

Installation

```
sudo apt install -y postgresql-16-ip4r
```

Enable it. Like pgvector in Chapter 6, ip4r’s control file does not set `trusted = true`, so a regular database-owning role can’t self-serve this one either — it has to be done once, by a superuser:

```
sudo -u postgres psql portsmith -c "CREATE EXTENSION ip4r;"
```

Confirm it:

```
SELECT extversion FROM pg_extension WHERE extname = 'ip4r';
```

```
extversion
```

```
-----
```

```
2.4
```

If it doesn’t show up: double-check which database you ran the `CREATE EXTENSION` against. `psql -l` lists every database on the cluster — it’s worth a glance if an extension you just enabled doesn’t appear where you expect it.

Loading the Data

Run the seed script

```
python data/ch07_seed.py
```

Expected output:

```
Connecting to: dbname=portsmith
Creating schema ...
Inserting 4 blocklist entries ...
Inserting 3 allowlist entries ...
Inserting 116 network events ...
Done – 116 rows in network_events, 4 blocklist entries, 3
allowlist entries.
```

Verify the load

Open `psql` `portsmith` and run these checks.

Check 1 — table structure:

```
\d network_events
```

Table			
Column	Type	Collation	Nullable
Default			
id	bigint		not null
nextval('network_events_id_seq'::regclass)			
event_type	text		not null
source_ip	ip4		not null
occurred_at	timestamp with time zone		not null
detail	text		not null

Indexes:

```
"network_events_pkey" PRIMARY KEY, btree (id)
```

Check constraints:

```
"network_events_event_type_check" CHECK (event_type = ANY
(ARRAY['login_success'::text, 'login_failure'::text,
'api_call'::text, 'api_error'::text]))
```

Check 2 — event counts by type:

```
SELECT event_type, COUNT(*) FROM network_events GROUP BY
event_type ORDER BY event_type;
```

event_type	count
api_call	47
api_error	9
login_failure	41
login_success	19

(4 rows)

Check 3 — list counts:

```
SELECT COUNT(*) FROM blocklists; -- 4
SELECT COUNT(*) FROM allowlists; -- 3
```

If all three match, proceed to the exercises.

Exercises

Exercise 1 — ip4r vs. the Built-in inet/cidr

1.1 — Storage

```
SELECT pg_column_size('203.0.113.5'::ip4)      AS ip4_size,
       pg_column_size('203.0.113.5'::inet)      AS inet_size,
       pg_column_size('203.0.113.0/24'::ip4r)   AS ip4r_size,
       pg_column_size('203.0.113.0/24'::cidr)   AS cidr_size;
```

ip4_size	inet_size	ip4r_size	cidr_size
4	10	8	10

ip4 is a fixed 4-byte integer — nothing but the address. inet costs 10 bytes because it's a varlena type carrying an address family byte and a netmask alongside the 4 address bytes, general enough to also hold IPv6. ip4r stores a range as two 4-byte bounds (8 bytes) rather than a network+prefix pair, and cidr costs the same 10 bytes as inet for the same reason. At one row this is noise; at hundreds of millions of rows in a real event log, 4 bytes versus 10 is not.

1.2 — cidr is stricter than you might expect

```
SELECT '192.168.1.5/24'::cidr;
```

```
ERROR:  invalid cidr value: "192.168.1.5/24"
DETAIL:  Value has bits set to right of mask.
```

cidr refuses any value where the host portion isn't all zero for the given mask — it only accepts genuine network addresses. inet is more permissive and keeps the host bits as part of the value:

```
SELECT '192.168.1.5/24'::inet;
```

inet
192.168.1.5/24

1.3 — ip4r doesn't require CIDR alignment at all

```
SELECT '203.0.113.5-203.0.113.20'::ip4r;
```

ip4r
203.0.113.5-203.0.113.20

```
SELECT '203.0.113.5-203.0.113.20'::cidr;
```

```
ERROR:  invalid input syntax for type cidr:
"203.0.113.5-203.0.113.20"
```

This is the real semantic difference, not just performance: `cidr` and `inet` can only represent power-of-two, mask-aligned networks. `ip4r` is a genuine range type — a start and an end bound, no requirement that they correspond to any CIDR block at all. A DHCP pool spanning `.5` through `.20` is a completely ordinary `ip4r` value and simply not expressible as a single `cidr`.

1.4 — The wider family

`ip4r` (the extension) is not limited to IPv4 despite the name:

```
\dx+ ip4r
```

lists casts and functions for `ip4`, `ip4r`, `ip6`, `ip6r`, and two polymorphic types, `ipaddress` and `iprange`, that accept either address family. This chapter sticks to IPv4 to match `network_events` and keep the examples concrete, but everything here has a direct IPv6 equivalent.

Exercise 2 — Containment: `>>` and `<=<`


2.1 — “Does any blocklist entry contain this address?”

```
SELECT id, category, description
FROM   blocklists
WHERE  cidr >> '203.0.113.150'::ip4;
```

```
id | category |
description
-----+-----
+-----
 1 | known_malicious | Repeated brute-force login attempts
against the resident portal, flagged by the security team
 2 | botnet | Subrange within 203.0.113.0/24 attributed
to a specific credential-stuffing botnet
(2 rows)
```

`>>` means “left range contains right value.” Two rows come back, not one — `203.0.113.150` falls inside both the broad `203.0.113.0/24` entry *and* the narrower `203.0.113.128/26` botnet subrange nested inside it. Real blocklists routinely have this shape: a wide, low-confidence range alongside a narrow, high-confidence one carved out of it.

203.0.113.150 falls inside both ranges at once — a real "multi-match"



2.2 — The same check, written the other way around

```
SELECT id, category, description
FROM   blocklists
WHERE  '203.0.113.150'::ip4 <= cidr;
```

Identical result set. `<=` means “left value is contained by (or equal to) right range” — `a > b` and `b <= a` are the same test from opposite sides. Which one reads more naturally depends on which value you think of as the “subject” of the query; both compile to the same containment check.

2.3 — Handle the multi-match case deliberately

Since Exercise 2.1 showed a single address can match more than one blocklist entry, any real lookup needs to decide what to do with that — return every match for an audit trail, or the most specific one for a quick yes/no. Exercise 6 builds the second.

Exercise 3 — GiST Indexing (and a Real Gotcha)

3.1 — The unindexed cost

```
EXPLAIN (ANALYZE, BUFFERS)
SELECT id, event_type, source_ip
FROM   network_events
WHERE  source_ip <= '203.0.113.0/24'::ip4r;
```

QUERY PLAN

```
-----
Seq Scan on network_events (cost=0.00..3.74 rows=1 width=44)
(actual time=0.006..0.016 rows=38 loops=1)
  Filter: ((source_ip)::ip4r <= '203.0.113.0/24'::ip4r)
  Rows Removed by Filter: 78
```

38 of 116 events fall inside that /24. At 116 rows this costs nothing; at production log volume, a sequential scan per lookup is exactly the bottleneck ip4r exists to remove.

3.2 — The naive index fails

```
CREATE INDEX idx_network_events_source_ip
ON network_events USING GIST (source_ip);
```

ERROR: data type ip4 has no default operator class for access method "gist"
HINT: You must specify an operator class for the index or define a default operator class for the data type.

This is worth sitting with rather than working around blindly: ip4r (the extension) ships a GiST operator class for ip4r (the **range** type) only.

```
SELECT opcname, amname
FROM pg_opclass oc JOIN pg_am am ON am.oid = oc.opcmethod
WHERE opcname ILIKE '%ip4%';
```

opcname		amname
-----+-----		
btree_ip4_ops		btree
btree_ip4r_ops		btree
hash_ip4_ops		hash
hash_ip4r_ops		hash
gist_ip4r_ops		gist

ip4 (a single address) has B-tree and hash opclasses for equality and ordering, but no GiST opclass at all — GiST is for indexing *containment and overlap*, which only makes sense for a range.

network_events.source_ip is declared ip4, a plain address column, so there's nothing there for GiST to build against directly.

3.3 — The fix: index the range-cast expression

```
CREATE INDEX idx_network_events_source_ip
ON network_events USING GIST ((source_ip::ip4r));
```

Every ip4 value casts losslessly to a single-address ip4r (a range whose start and end are the same address), which does have a GiST opclass. This is an **expression index** — it indexes the result of source_ip::ip4r, not the raw column — so queries have to use the same cast for the planner to recognize a match:

```
SET enable_seqscan = off;

EXPLAIN (ANALYZE, BUFFERS)
SELECT id, event_type, source_ip
```

```

FROM network_events
WHERE source_ip::ip4r <= '203.0.113.0/24'::ip4r;

SET enable_seqscan = on;

```

QUERY PLAN

```

-----
Index Scan using idx_network_events_source_ip on network_events
(cost=0.14..8.15 rows=1 width=44) (actual time=0.105..0.110
rows=38 loops=1)
  Index Cond: ((source_ip)::ip4r <= '203.0.113.0/24'::ip4r)

```

Same 38 rows, now via Index Scan. The lesson generalizes beyond `ip4r`: when a GiST opclass exists for a *range* type but your column stores *points*, an expression index bridges the gap — cast to the range type in both the index definition and every query that should use it.

Exercise 4 — Aggregating by Subnet

4.1 — The `ip4r`-native way

```

SELECT ip4r_net_prefix(source_ip, 24) AS subnet_24, COUNT(*) AS
events
FROM network_events
GROUP BY subnet_24
ORDER BY events DESC, subnet_24;

```

subnet_24	events
203.0.113.0/24	38
192.0.2.0/24	31
198.51.100.0/24	28
240.1.2.0/24	11
100.64.5.0/24	4
100.64.9.0/24	4

(6 rows)

`ip4r_net_prefix(address, prefix_length)` computes the containing network for an address directly, staying entirely within `ip4r`'s own types — no cast round-trip needed.

4.2 — The same result via the built-in `network()`/`set_masklen()`

`ip4` casts directly to `cidr`, so the standard PostgreSQL network functions work too, if you'd rather not learn `ip4r`-specific function names:

```

SELECT network(set_masklen(source_ip::cidr, 24)) AS subnet_24,
COUNT(*) AS events
FROM network_events
GROUP BY subnet_24
ORDER BY events DESC, subnet_24;

```

Identical results. `set_masklen()` overrides the prefix length on a `cidr/` `inet` value, and `network()` zeroes out the host bits to return the network address — two built-in functions doing in two steps what `ip4r_net_prefix()` does in one.

4.3 — masklen() on the blocklist itself

```

SELECT cidr, masklen(cidr::cidr) AS prefix_length, category
FROM blocklists ORDER BY id;

```

cidr	prefix_length	category
203.0.113.0/24	24	known_malicious
203.0.113.128/26	26	botnet
240.1.2.0/25	25	tor_exit_node
198.51.100.0/28	28	brute_force_source

Worth noticing as a pattern, not just a query result: the narrower the prefix (higher number, smaller range), the more specific and higher-confidence the category tends to be here — /24 for a broad “something’s wrong in this range” flag, /26 and /28 for a credential-stuffing botnet and an automated feed’s specific finding. Range size is itself a signal about how much to trust an entry.

Exercise 5 — Overlap Detection Between Allow and Block Lists

5.1 — Find every conflict with &&

```

SELECT b.id AS block_id, b.cidr AS blocked_range, b.category,
a.id AS allow_id, a.cidr AS allowed_range, a.description
FROM blocklists b
JOIN allowlists a ON b.cidr && a.cidr
ORDER BY b.id, a.id;

```

block_id	blocked_range	category	allow_id	description
4	198.51.100.0/28	brute_force_source	1	198.51.100.0/24 Portsmouth City Hall internal network and staff

VPN
(1 row)

&& is the general overlap test — true if the two ranges share *any* address, regardless of which contains which. One real conflict: an automated brute-force detection feed flagged 198.51.100.0/28, a sub-range that sits entirely inside the city’s own allowlisted VPN block. This is exactly the failure mode automated blocklist feeds produce in practice — someone’s VPN concentrator fails logins at a rate that looks like an attack from outside, and gets flagged from inside a range that was explicitly trusted.

5.2 — Confirm it isn’t hypothetical

```
SELECT DISTINCT source_ip
FROM   network_events
WHERE  source_ip::ip4r <= '198.51.100.0/28'::ip4r
ORDER BY source_ip;
```

```
source_ip
-----
198.51.100.5
198.51.100.7
198.51.100.12
```

Three real addresses in `network_events` sit inside the disputed range — this isn’t an edge case sitting unused in the blocklist table, it’s actively affecting real logged traffic from the city’s own network.

Exercise 6 — A Real-Time “Is This IP Blocked?” Function

6.1 — Index the lookup tables themselves

Exercise 3 needed an expression index because `network_events.source_ip` is `ip4` (a point). `blocklists.cidr` and `allowlists.cidr` are already `ip4r`, so they index directly, no casting required:

```
CREATE INDEX idx_blocklists_cidr ON blocklists USING GIST (cidr);
CREATE INDEX idx_allowlists_cidr ON allowlists USING GIST (cidr);

SET enable_seqscan = off;

EXPLAIN (ANALYZE, BUFFERS)
SELECT category, description FROM blocklists WHERE cidr >>
'203.0.113.150'::ip4;

SET enable_seqscan = on;
```

PLAN

```
-----
Index Scan using idx_blocklists_cidr on blocklists
(cost=0.13..8.15 rows=1 width=64) (actual time=0.073..0.074 rows=2
loops=1)
  Index Cond: (cidr >> '203.0.113.150'::ip4r)
```

6.2 — The function: allowlist wins

Exercise 5 found a real conflict. A production “is this blocked?” check has to resolve it one way, consistently — this function checks the allowlist *first* and short-circuits if it matches, so a trusted range is never blocked no matter what an automated feed says about a sub-range of it:

```
CREATE OR REPLACE FUNCTION is_blocked(check_ip ip4)
RETURNS TABLE (blocked BOOLEAN, reason TEXT) AS $$
BEGIN
  IF EXISTS (SELECT 1 FROM allowlists WHERE cidr >> check_ip)
THEN
    RETURN QUERY SELECT FALSE, 'allowlisted'::TEXT;
    RETURN;
  END IF;

  RETURN QUERY
  SELECT TRUE, b.category || ': ' || b.description
  FROM   blocklists b
  WHERE  b.cidr >> check_ip
  LIMIT 1;

  IF NOT FOUND THEN
    RETURN QUERY SELECT FALSE, 'not listed'::TEXT;
  END IF;
END;
$$ LANGUAGE plpgsql STABLE;
```

LIMIT 1 on the blocklist branch is deliberate — Exercise 2 showed an address can match multiple blocklist entries, and for a yes/no gate one matching reason is enough; an audit tool would drop the LIMIT and return every match instead.

6.3 — Test it against every case this chapter built

```
SELECT * FROM is_blocked('203.0.113.150'); -- plain blocklist
match

blocked |
reason
-----
+-----
```


t | known_malicious: Repeated brute-force login attempts against the resident portal, flagged by the security team

```
SELECT * FROM is_blocked('198.51.100.7'); -- the Exercise 5 conflict
```

blocked	reason
f	allowlisted

Blocklisted *and* allowlisted, and the function correctly refuses to block it — exactly the resolution Exercise 5’s conflict needed.

```
SELECT * FROM is_blocked('192.0.2.10'); -- vendor allowlist, no conflict
SELECT * FROM is_blocked('8.8.8.8'); -- not on any list
```

blocked	reason
f	allowlisted

blocked	reason
f	not listed

Four calls, four distinct real outcomes, all backed by GiST-indexed lookups against tables that would scale to millions of blocklist entries without changing a line of this function.

Exercise 7 (Additional) — Rate Limiting by Host *and* /24 with a Token Bucket

Blocklists are a permanent, deliberate verdict — someone reviewed a range and decided it’s hostile. Rate limiting is different: it’s a *temporary* throttle applied to traffic that isn’t necessarily malicious, just too frequent, and it needs to make that decision on every single request, fast. This exercise builds one of the most common shapes for it — a **token bucket** — and applies it at two levels at once: per individual host, and per /24 netblock, because some abuse only shows up when you stop looking at hosts one at a time.


7.1 — The token bucket model

Each bucket has a **capacity** (the size of a burst it can absorb all at once) and a **refill rate** (tokens added per second, up to capacity). Every request tries to consume one token: if at least one is available, the request is allowed and a token is spent; if not, it’s denied. Tokens refill

continuously based on real elapsed time, not on a fixed clock tick — a bucket that's been idle for ten seconds has ten seconds' worth of refill waiting, whether or not anything asked.

7.2 — Why two buckets, not one

A per-host limit alone has a blind spot: ten different addresses in the same /24, each individually staying just under the limit, add up to ten times the traffic that netblock was ever supposed to send — a classic shape for a botnet spread across one compromised network, or NAT'd traffic from behind a single gateway. A per-host bucket alone never trips. Checking a *second*, looser bucket keyed to the containing /24 catches exactly that pattern, without needing to lower the per-host limit enough to hurt legitimate single users.



`check_rate_limit()`, built below, is exactly this diagram: the host bucket is checked — and consumed — first, and only a request that clears it goes on to spend a netblock token too.

7.3 — Schema: bucket keyed by ip4r, not ip4

```
CREATE TABLE rate_limit_buckets (  
  bucket      ip4r PRIMARY KEY,  
  capacity    NUMERIC NOT NULL,
```

```

        refill_rate NUMERIC NOT NULL,
        tokens      NUMERIC NOT NULL,
        updated_at  TIMESTAMPTZ NOT NULL DEFAULT clock_timestamp()
    );

```

Using ip4r as the key, rather than ip4 plus a separate “is this a host or a netblock” flag, means a host bucket and a netblock bucket are just two different-sized ranges sharing one table and one lookup path — a host is stored as its /32, a netblock as its /24.

7.4 — The atomic refill-and-consume function

The refill math and the allow/deny decision have to happen as one atomic operation — two concurrent requests against the same bucket must not both read “4 tokens available” and both proceed, the same race Chapter 3 built FOR UPDATE SKIP LOCKED to avoid in the job queue. Here the fix is simpler — plain FOR UPDATE, no SKIP LOCKED — because a rate limit *should* make the second concurrent request wait its turn for the lock rather than skip ahead to a different row; there’s only one row per bucket, not a queue of interchangeable ones:

```

CREATE OR REPLACE FUNCTION try_consume_bucket(
    target_bucket ip4r,
    p_capacity     NUMERIC,
    p_refill_rate  NUMERIC
) RETURNS TABLE (allowed BOOLEAN, tokens_remaining NUMERIC) AS $$
BEGIN
    INSERT INTO rate_limit_buckets (bucket, capacity, refill_rate,
tokens, updated_at)
        VALUES (target_bucket, p_capacity, p_refill_rate, p_capacity,
clock_timestamp())
        ON CONFLICT (bucket) DO NOTHING;

    RETURN QUERY
    WITH refilled AS (
        SELECT b.bucket,
            LEAST(b.capacity, b.tokens + b.refill_rate *
                EXTRACT(EPOCH FROM (clock_timestamp() -
b.updated_at))) AS available
        FROM   rate_limit_buckets b
        WHERE  b.bucket = target_bucket
        FOR UPDATE
    )
    UPDATE rate_limit_buckets u
    SET     tokens      = CASE WHEN r.available >= 1 THEN
r.available - 1 ELSE r.available END,
        updated_at = clock_timestamp()
    FROM   refilled r
    WHERE  u.bucket = r.bucket
    RETURNING (r.available >= 1), u.tokens;

```

```
END;
$$ LANGUAGE plpgsql;
```

The `INSERT ... ON CONFLICT DO NOTHING` lazily creates a bucket at full capacity the first time an address is ever seen — no separate provisioning step. Note `clock_timestamp()`, not `now()`: this is the same gotcha from Chapter 3 — `now()` freezes at transaction start, which would make every refill calculation inside one transaction compute against the same frozen instant instead of real elapsed wall-clock time.

7.5 — Combine host and netblock into one check

```
CREATE OR REPLACE FUNCTION check_rate_limit(check_ip ip4)
RETURNS TABLE (allowed BOOLEAN, reason TEXT, host_tokens NUMERIC,
netblock_tokens NUMERIC) AS $$
DECLARE
    host_result      RECORD;
    netblock_result  RECORD;
BEGIN
    -- Host bucket: tight limit, catches a single misbehaving
    address.
    SELECT * INTO host_result
    FROM   try_consume_bucket(check_ip::ip4r, 5, 1);

    IF NOT host_result.allowed THEN
        RETURN QUERY SELECT FALSE, 'host limit exceeded'::TEXT,
host_result.tokens_remaining, NULL::NUMERIC;
        RETURN;
    END IF;

    -- Netblock bucket: looser limit, catches abuse spread across
    many
    -- hosts in the same /24, each individually within their own
    budget.
    SELECT * INTO netblock_result
    FROM   try_consume_bucket(ip4r_net_prefix(check_ip, 24), 20,
5);

    IF NOT netblock_result.allowed THEN
        RETURN QUERY SELECT FALSE, 'netblock limit exceeded'::TEXT,
host_result.tokens_remaining, netblock_result.tokens_remaining;
        RETURN;
    END IF;

    RETURN QUERY SELECT TRUE, 'ok'::TEXT,
host_result.tokens_remaining, netblock_result.tokens_remaining;
END;
$$ LANGUAGE plpgsql;
```

The host bucket (capacity 5, refills 1/sec) is checked — and, critically, *consumed* — first. Only if the host bucket allows the request does the function go on to check the netblock bucket (capacity 20, refills 5/sec). This ordering is a real design decision, not an accident: it means a request that the netblock bucket ultimately rejects still spent a host token on the way there. The alternative — check both before consuming either — avoids that, at the cost of two extra reads per request. For a security throttle, spending a token on a request you were going to reject anyway is a rounding error; use whichever trade-off matches what you’re actually optimizing for.

7.6 — A single host bursting past its own limit

```
SELECT allowed, reason, round(host_tokens, 2), round(netblock_tokens,
2)
FROM   check_rate_limit('203.0.113.5');
-- run six times in a row
```

allowed	reason	host_tokens	netblock_tokens
t	ok	4.00	19.00
t	ok	3.05	18.27
t	ok	2.09	17.45
t	ok	1.13	16.63
t	ok	0.17	15.83
f	host limit exceeded	0.20	

Five rapid calls spend the host bucket’s burst capacity of 5 (with a sliver of refill accumulating between each real request — these were genuinely separate round trips, not a batch); the sixth is denied, purely on the host bucket. The netblock bucket barely moved. (Your own numbers will differ in the small decimals — they depend on the exact wall-clock gap between your requests — but the allow/allow/allow/allow/allow/deny pattern will not.)

7.7 — Distributed abuse across a /24: the case this whole exercise exists for

Five different hosts in the same netblock, each making five rapid requests — each individually within its own host budget:

```
-- for host in 203.0.113.10 .. 203.0.113.14, 5 calls each:
SELECT allowed, reason, round(host_tokens, 2), round(netblock_tokens,
2)
FROM   check_rate_limit(:host);
```

host	allowed	reason	host_tokens
netblock_tokens			

+-----+-----			
203.0.113.10	t	ok	4.00
	19.00		
203.0.113.10	t	ok	3.04
	18.19		
203.0.113.10	t	ok	2.08
	17.38		
203.0.113.10	t	ok	1.11
	16.57		
203.0.113.10	t	ok	0.15
	15.76		
203.0.113.11	t	ok	4.00
	14.94		
203.0.113.11	t	ok	3.04
	14.13		
203.0.113.11	t	ok	2.07
	13.32		
203.0.113.11	t	ok	1.11
	12.51		
203.0.113.11	t	ok	0.15
	11.70		
203.0.113.12	t	ok	4.00
	10.88		
203.0.113.12	t	ok	3.04
	10.07		
203.0.113.12	t	ok	2.08
	9.27		
203.0.113.12	t	ok	1.11
	8.46		
203.0.113.12	t	ok	0.15
	7.64		
203.0.113.13	t	ok	4.00
	6.82		
203.0.113.13	t	ok	3.04
	6.03		
203.0.113.13	t	ok	2.08
	5.23		
203.0.113.13	t	ok	1.12
	4.42		
203.0.113.13	t	ok	0.16
	3.62		
203.0.113.14	t	ok	4.00
	2.81		
203.0.113.14	t	ok	3.04
	2.02		
203.0.113.14	t	ok	2.08
	1.20		
203.0.113.14	t	ok	1.11
	0.38		
203.0.113.14	f	netblock limit exceeded	0.15
	0.56		

Read the `host_tokens` column top to bottom: **every single request succeeds at the host level** — no address ever exhausts its own bucket, each one’s burst-of-5 fits comfortably within its own limit every time. And yet the 25th request, from the fifth host, is denied — for `netblock limit` exceeded, not a host limit. Twenty-five individually reasonable requests from five individually well-behaved addresses added up to more than the /24 as a whole was ever budgeted for, and the second bucket caught it. This is precisely the pattern a per-host-only rate limiter is structurally blind to.

Summary — What You Should Now Know

Tool	What it does
<code>ip4 / ip4r</code>	Dedicated address and range types — 4 and 8 bytes, versus 10 for <code>inet/cidr</code>
<code>ip6 / ip6r / ipaddress / iprange</code>	The same machinery for IPv6, and family-agnostic polymorphic types
<code>cidr</code>	Stricter than <code>inet</code> — rejects any value with host bits set past the mask
<code>'a-b'::ip4r</code>	An arbitrary range, not required to be CIDR-aligned — something <code>cidr</code> cannot express at all
<code>a >> b / b <=< a</code>	Containment, read from either direction — “does a contain b?” / “is b contained by a?”
<code>a && b</code>	Overlap — true if two ranges share any address, regardless of which contains which
<code>GIST (col) on ip4r</code>	Fast containment/overlap lookups — but only for <code>ip4r</code> , <code>ip4</code> has no GiST opclass
<code>GIST ((col::ip4r)) on ip4</code>	The expression-index fix when your column stores points, not ranges
<code>ip4r_net_prefix(ip, n)</code>	

Tool	What it does
	Native way to compute the containing /n network for an address
<code>network(set_masklen(col::cidr, n))</code>	The same computation via built-in <code>inet/cidr</code> functions
Token bucket, keyed by <code>ip4r</code>	Rate limiting at both host (/32) and netblock (/24) scope from one table, catching distributed abuse a per-host limit alone can't see

The key design insight from this chapter is that a blocklist and an allowlist are not just two tables you check in some order — they are two sources of truth that *will* disagree, and Exercise 5 found a real disagreement in four hand-written rows almost by accident. A security system that doesn't actively check for that disagreement is trusting whichever list happens to get consulted first, silently. && makes finding every such conflict a single query; `is_blocked()` makes resolving it a permanent, enforced policy instead of a hope.

Going further: this chapter's `is_blocked()` resolves conflicts with a fixed rule — allowlist always wins. A more nuanced system might weight by range specificity (Exercise 4's observation that narrower ranges tend to be higher-confidence) or by how recently an entry was added. For very large blocklists — commercial threat-intelligence feeds run to millions of entries — also benchmark GiST against `ip4r`'s B-tree and hash opclasses for the access patterns that are pure equality or pure sorting rather than containment, since GiST is the right tool specifically for range queries, not a strictly faster B-tree replacement. And if this system needs to block traffic in real time rather than just log a verdict, the natural next step is exposing `is_blocked()` at the application's edge — a reverse proxy or API gateway calling it per request — which is exactly the kind of function Chapter 10's PostgREST chapter would expose as an RPC endpoint. The same is true of `check_rate_limit()` from Exercise 7 — a real deployment calls it once per incoming request, before any other work happens, and a single `rate_limit_buckets` table means every application server sees the same limits regardless of which one handled the last request, something an in-process rate limiter can't offer without a shared store of its own. If per-request PL/pgSQL round trips ever become the bottleneck rather than the safeguard, that's the point at which a dedicated in-memory limiter (Redis and similar) starts to earn its keep — but plenty of systems

never reach the traffic level where that trade-off is worth the added moving part.

Chapter 8 — Declarative Partitioning and BRIN Indexes

“A table doesn’t get slow because it’s big. It gets slow because every query has to consider all of it — partitioning doesn’t shrink the table, it shrinks the part any single query has to look at.”

Background

Every chapter so far has reached for an extension — `pg_trgm`, `pgvector`, `ip4r` — to get a data type or index PostgreSQL doesn’t ship with. This chapter is different: declarative partitioning (since PostgreSQL 10) and BRIN indexes (since PostgreSQL 9.5) are both built-in core features. No `CREATE EXTENSION` anywhere in this chapter. What they solve isn’t “how do I store this,” it’s “how do I keep a table that never stops growing from turning every query, every index, and every maintenance operation into a scan of everything that has ever happened.”

Partitioning splits one logical table into many physical ones, each holding a slice of the rows — by date range, by category, by hash — while every query still just says `SELECT ... FROM sensor_readings`. The planner decides, before it reads a single row, which physical partitions could possibly contain a match and skips the rest entirely. That’s **partition pruning**, and it’s the entire value proposition: a query asking for one month of data out of a year doesn’t pay for the other eleven months, not even to rule them out.

BRIN (Block Range INdex) is a different kind of index for a specific shape of data: values that are large in volume but naturally correlated with physical storage order. A B-tree index stores an entry for every single row and costs roughly proportional to the row count. A BRIN index stores one small summary — just a min and max — per range of physical disk pages (128 pages by default), and costs almost nothing regardless of how many rows are in the table. The catch, and it’s a real one: BRIN only helps if physical order and value order actually correlate. For a table where rows are inserted in timestamp order and essentially never updated — exactly what an append-only sensor log looks like — the correlation is close to perfect, and that’s the case this chapter builds.

The Scenario

Portsmouth has quietly wired up 120 sensors across the city — temperature sensors, traffic counters, and air quality monitors — all reporting on a schedule, all day, every day. A year of this is over ten million rows, and it only grows. The data starts life the way most time-series data actually starts: as one plain, unpartitioned table, because nobody provisions a partitioning scheme before they have a problem. This chapter is that problem arriving, and fixing it.

Table	Purpose
sensors	Metadata for all 120 sensors — type, label, neighborhood/street
sensor_readings_staging	~10.5 million raw readings for 2024, as one unpartitioned table
sensor_readings	<i>(built in this chapter)</i> the same data, partitioned by month
sensor_readings_by_type	<i>(built in this chapter)</i> a second copy, partitioned by type <i>and</i> month

sensor_readings_staging is the state the data starts in. Building sensor_readings from it — deciding the scheme, creating the partitions, backfilling — is the exercise, not something handed to you already done.

Exercise Goals

By the end of this chapter you will be able to:

- Design and create a range-partitioned table, and understand why the partition key has to be chosen before the first row is loaded, not after.
- Backfill a partitioned table from an existing unpartitioned one and confirm that PostgreSQL routes every row to the correct partition automatically.
- Read an EXPLAIN plan well enough to tell the difference between a query that touches one partition and one that touches all of them.
- Build a BRIN index and explain, using real numbers, why it can be hundreds of times smaller than a B-tree on the same column — and what it gives up to get there.

- Drop a month of data in milliseconds via `DROP TABLE` on a partition, and see exactly what that avoids compared to `DELETE` on an unpartitioned table.
 - Combine two partitioning strategies — `LIST` over `RANGE` — into one table, and confirm the planner prunes on *both* dimensions at once.
-

Installation

Nothing to install. `PARTITION BY` and `USING BRIN` are both part of core PostgreSQL — if `psql --version` says 10 or later for partitioning, or 9.5 or later for BRIN, you already have everything this chapter uses. (The Portsmouth cluster is PostgreSQL 16, well past both.)

Loading the Data

Run the seed script

```
python data/ch08_seed.py
```

Expected output (the exact seconds will vary with your machine; the row counts will not):

```
Connecting to: dbname=portsmith
Creating schema ...
Inserting 120 sensors ...
Generating and loading readings (this takes a few minutes) ...
  generated January 2024 – 892,800 rows
  generated February 2024 – 835,200 rows
  generated March 2024 – 892,800 rows
  generated April 2024 – 864,000 rows
  generated May 2024 – 892,800 rows
  generated June 2024 – 864,000 rows
  generated July 2024 – 892,800 rows
  generated August 2024 – 892,800 rows
  generated September 2024 – 864,000 rows
  generated October 2024 – 892,800 rows
  generated November 2024 – 864,000 rows
  generated December 2024 – 892,800 rows
Indexing staging table on recorded_at (mirrors a typical pre-
partitioning table) ...
Done – 10,540,800 rows in sensor_readings_staging (131.3s to
generate + load), 120 sensors.
```

Months vary in row count because February is shorter and because temperature/traffic sensors report every 5 minutes while air quality sensors report every 15 — not a bug, just three sensor types on three different schedules, same as a real deployment would have.

Pin the session timezone before you do anything else

```
SET timezone = 'UTC';
```

This matters more than it looks like it should. TIMESTAMPTZ values are stored as absolute instants, but a bound like '2024-01-01' in a CREATE TABLE ... PARTITION OF ... FOR VALUES FROM (...) clause gets interpreted relative to the session's timezone *at the moment the partition is created*. Two people creating “the same” January partition in two different session timezones will get two different absolute boundaries. Run every example in this chapter in a UTC session — this SET at the top of each psql session is enough — and the partition boundaries and the displayed timestamps will both match what's printed below exactly.

Verify the load

Check 1 — the metadata table:

```
\d sensors

      Table "public.sensors"
  Column      | Type      | Collation | Nullable | Default
-----+-----+-----+-----+-----
 id           | integer   |           | not null |
 sensor_type  | text      |           | not null |
 label        | text      |           | not null |
 neighborhood | text      |           | not null |
Indexes:
    "sensors_pkey" PRIMARY KEY, btree (id)
Check constraints:
    "sensors_sensor_type_check" CHECK (sensor_type = ANY
    (ARRAY['temperature'::text, 'traffic'::text,
    'air_quality'::text]))
Referenced by:
    TABLE "sensor_readings_staging" CONSTRAINT
    "sensor_readings_staging_sensor_id_fkey" FOREIGN KEY (sensor_id)
    REFERENCES sensors(id)
```

Check 2 — the staging table structure:

```
\d sensor_readings_staging
```

```

Table
"public.sensor_readings_staging"
  Column      |      Type      | Collation | Nullable
|      Default
+-----+-----+-----+-----+
id            | bigint         |           | not null |
generated always as identity
sensor_id     | integer        |           | not null |
sensor_type   | text           |           | not null |
reading_value | double precision |           | not null |
recorded_at   | timestamp with time zone |           | not null |

```

Indexes:

```
"idx_sensor_readings_staging_recorded_at" btree (recorded_at)
```

Check constraints:

```
"sensor_readings_staging_sensor_type_check" CHECK (sensor_type
= ANY (ARRAY['temperature'::text, 'traffic'::text,
'air_quality'::text]))
```

Foreign-key constraints:

```
"sensor_readings_staging_sensor_id_fkey" FOREIGN KEY
(sensor_id) REFERENCES sensors(id)
```

Note this table already has a B-tree index on recorded_at — a perfectly ordinary choice for “the column I filter on,” and exactly what a lot of real unpartitioned time-series tables look like right up until someone asks why disk usage keeps climbing.

Check 3 — sensor and reading counts by type:

```
SELECT sensor_type, COUNT(*) FROM sensors GROUP BY sensor_type
ORDER BY sensor_type;
```

```

sensor_type | count
+-----+-----+
air_quality |    30
temperature |    50
traffic     |    40
(3 rows)

```

```
SELECT sensor_type, COUNT(*) FROM sensor_readings_staging GROUP BY
sensor_type ORDER BY sensor_type;
```

```

sensor_type | count
+-----+-----+
air_quality | 1054080
temperature | 5270400
traffic     | 4216320
(3 rows)

```

If these match, proceed to the exercises.

Exercises

Exercise 1 — Designing and Creating the Partitioned Table

1.1 — Choose the partition key up front

Partitioning isn't something you bolt onto an existing table — `PARTITION BY` is part of `CREATE TABLE` and can't be added with `ALTER TABLE` afterward. That means the decision has to be made before a single row is loaded into the partitioned structure, based on how the table will actually be queried. `sensor_readings` is queried almost exclusively by time range ("show me last week," "roll up by month"), so `RANGE (recorded_at)`, monthly, is the natural key — coarse enough to keep the partition count manageable, fine enough that most queries touch only one or two partitions.

1.2 — Create the partitioned table and its monthly partitions

```
CREATE TABLE sensor_readings (  
  id          BIGINT GENERATED ALWAYS AS IDENTITY,  
  sensor_id   INTEGER NOT NULL REFERENCES sensors(id),  
  sensor_type TEXT NOT NULL  
              CHECK (sensor_type IN ('temperature',  
'traffic', 'air_quality')),  
  reading_value DOUBLE PRECISION NOT NULL,  
  recorded_at  TIMESTAMPTZ NOT NULL  
) PARTITION BY RANGE (recorded_at);
```

```
CREATE TABLE sensor_readings_2024_01 PARTITION OF sensor_readings  
FOR VALUES FROM ('2024-01-01') TO ('2024-02-01');  
CREATE TABLE sensor_readings_2024_02 PARTITION OF sensor_readings  
FOR VALUES FROM ('2024-02-01') TO ('2024-03-01');  
CREATE TABLE sensor_readings_2024_03 PARTITION OF sensor_readings  
FOR VALUES FROM ('2024-03-01') TO ('2024-04-01');  
CREATE TABLE sensor_readings_2024_04 PARTITION OF sensor_readings  
FOR VALUES FROM ('2024-04-01') TO ('2024-05-01');  
CREATE TABLE sensor_readings_2024_05 PARTITION OF sensor_readings  
FOR VALUES FROM ('2024-05-01') TO ('2024-06-01');  
CREATE TABLE sensor_readings_2024_06 PARTITION OF sensor_readings  
FOR VALUES FROM ('2024-06-01') TO ('2024-07-01');  
CREATE TABLE sensor_readings_2024_07 PARTITION OF sensor_readings  
FOR VALUES FROM ('2024-07-01') TO ('2024-08-01');  
CREATE TABLE sensor_readings_2024_08 PARTITION OF sensor_readings  
FOR VALUES FROM ('2024-08-01') TO ('2024-09-01');  
CREATE TABLE sensor_readings_2024_09 PARTITION OF sensor_readings  
FOR VALUES FROM ('2024-09-01') TO ('2024-10-01');
```



```
CREATE TABLE sensor_readings_2024_10 PARTITION OF sensor_readings
FOR VALUES FROM ('2024-10-01') TO ('2024-11-01');
CREATE TABLE sensor_readings_2024_11 PARTITION OF sensor_readings
FOR VALUES FROM ('2024-11-01') TO ('2024-12-01');
CREATE TABLE sensor_readings_2024_12 PARTITION OF sensor_readings
FOR VALUES FROM ('2024-12-01') TO ('2025-01-01');
CREATE TABLE sensor_readings_default PARTITION OF sensor_readings
DEFAULT;
```

Each `FOR VALUES FROM (...) TO (...)` is a half-open interval — inclusive of the lower bound, exclusive of the upper — so consecutive months tile the year with no gap and no overlap. The last one, `sensor_readings_default`, is not optional in any meaningful sense: without it, any row whose `recorded_at` doesn't fall in one of the twelve defined ranges is simply rejected at insert time. With it, such a row is accepted and routed there instead — silently, unless you go looking. Exercise 2 finds a case where that's exactly what happens.

1.3 — Confirm the structure

```
\d sensor_readings

Partitioned table
"public.sensor_readings"
  Column      |          Type          | Collation | Nullable
|      Default
+-----+-----+-----+-----+
id            | bigint                 |           | not null |
generated always as identity
sensor_id     | integer                |           | not null |
sensor_type   | text                   |           | not null |
reading_value | double precision       |           | not null |
recorded_at   | timestamp with time zone |           | not null |
Partition key: RANGE (recorded_at)
Check constraints:
  "sensor_readings_sensor_type_check" CHECK (sensor_type = ANY
  (ARRAY['temperature'::text, 'traffic'::text,
  'air_quality'::text]))
Foreign-key constraints:
  "sensor_readings_sensor_id_fkey" FOREIGN KEY (sensor_id)
  REFERENCES sensors(id)
Number of partitions: 13 (Use \d+ to list them.)
```

`sensor_readings` itself has no storage of its own — it's a routing target. Every row you insert into it physically lands in exactly one of the thirteen tables underneath.

Exercise 2 — Backfilling and Watching Automatic Partition Routing

2.1 — Backfill from staging

```
INSERT INTO sensor_readings (sensor_id, sensor_type,
reading_value, recorded_at)
SELECT sensor_id, sensor_type, reading_value, recorded_at
FROM   sensor_readings_staging;
```

```
INSERT 0 10540800
Time: 87814.162 ms (01:27.814)
```

One `INSERT ... SELECT`, no destination named beyond `sensor_readings` itself — PostgreSQL evaluates each row’s `recorded_at` against the partition bounds and routes it to the matching child table. Nothing in this statement says which of the thirteen tables any given row ends up in.

2.2 — Confirm the routing, per partition

```
SELECT tableoid::regclass AS partition, COUNT(*)
FROM   sensor_readings
GROUP BY tableoid
ORDER BY 1;
```

partition	count
sensor_readings_2024_01	892800
sensor_readings_2024_02	835200
sensor_readings_2024_03	892800
sensor_readings_2024_04	864000
sensor_readings_2024_05	892800
sensor_readings_2024_06	864000
sensor_readings_2024_07	892800
sensor_readings_2024_08	892800
sensor_readings_2024_09	864000
sensor_readings_2024_10	892800
sensor_readings_2024_11	864000
sensor_readings_2024_12	891648
sensor_readings_default	1152

(13 rows)

`tableoid` is a hidden system column every table has, and on a partitioned table it tells you which physical partition a given row actually lives in — the only reliable way to ask “where did this row end up?” after the fact.

2.3 — The row count that should stop you

Every month has the row count you'd expect from Exercise 0's per-type counts — except December, which is short exactly 1,152 rows, and `sensor_readings_default`, which has exactly 1,152 rows in it. That's not a coincidence:

```
SELECT sensor_id, MIN(recorded_at), MAX(recorded_at), COUNT(*)
FROM   sensor_readings
WHERE  tableoid = 'sensor_readings_default'::regclass
GROUP BY sensor_id;
```

```
sensor_id |          min          |          max          |
count
-----+-----+-----+
+-----+
      17 | 2025-12-27 00:00:00+00 | 2025-12-31 23:55:00+00 |
1152
```

Sensor 17 — a temperature sensor — has a clock bug: for its last four days of reporting, it stamped its readings a full year fast. Nothing about the INSERT failed or warned; the rows are sitting exactly where the schema said out-of-range data should go. This is precisely the scenario a DEFAULT partition exists for, and precisely why “does the default partition have anything in it?” is worth checking after any real backfill, not just this synthetic one — it's the difference between a sensor firmware bug you catch in five minutes and one you catch when someone asks why a report is missing a device.

Exercise 3 — Partition Pruning, Seen in EXPLAIN

3.1 — A date-filtered query touches one partition

```
EXPLAIN (ANALYZE, BUFFERS)
SELECT sensor_type, COUNT(*), round(AVG(reading_value)::numeric,
2)
FROM   sensor_readings
WHERE  recorded_at >= '2024-06-01' AND recorded_at < '2024-07-01'
GROUP BY sensor_type
ORDER BY sensor_type;
```


```
Finalize GroupAggregate (cost=12411.99..12480.46 rows=200
width=72) (actual time=182.595..203.871 rows=3 loops=1)
  Group Key: sensor_readings.sensor_type
  Buffers: shared hit=106 read=7168, temp read=3264 written=3276
  -> Gather Merge (cost=12411.99..12473.96 rows=400 width=72)
    (actual time=152.757..203.835 rows=9 loops=1)
      Workers Planned: 2
      Workers Launched: 2
      -> Partial GroupAggregate (cost=11411.97..11427.77
```

```

rows=200 width=72) (actual time=98.794..149.785 rows=3 loops=3)
    Group Key: sensor_readings.sensor_type
    -> Sort (cost=11411.97..11415.42 rows=1380
width=40) (actual time=93.166..120.544 rows=288000 loops=3)
        Sort Key: sensor_readings.sensor_type
        Sort Method: external merge  Disk: 8720kB
        -> Parallel Seq Scan on
sensor_readings_2024_06 sensor_readings (cost=0.00..11340.00
rows=1380 width=40) (actual time=0.034..37.199 rows=288000
loops=3)
            Filter: ((recorded_at >= '2024-06-01
00:00:00+00':timestamp with time zone) AND (recorded_at <
'2024-07-01 00:00:00+00':timestamp with time zone))
    Planning Time: 6.417 ms
    Execution Time: 205.458 ms

```

The GROUP BY sensor_type on top adds a sort and some parallel-worker bookkeeping that isn't the point here — the line that matters is Parallel Seq Scan on sensor_readings_2024_06, singular. Eleven other monthly partitions and the default partition exist; none of them appear anywhere in this plan. The planner threw them out during planning, before execution ever started, based purely on the WHERE clause matching only June's range.



3.2 — The same query, no date filter, for contrast

```

EXPLAIN
SELECT sensor_type, COUNT(*)
FROM   sensor_readings
GROUP BY sensor_type;

```

```

Finalize GroupAggregate (cost=156197.01..156247.68 rows=200
width=40)
-> Gather Merge (cost=156197.01..156243.68 rows=400 width=40)
    -> Sort (cost=155196.99..155197.49 rows=200 width=40)

```

```

-> Partial HashAggregate
(cost=155187.34..155189.34 rows=200 width=40)
-> Parallel Append (cost=0.00..138351.15
rows=3367238 width=32)
-> Parallel Seq Scan on
sensor_readings_2024_01 sensor_readings_1 ...
-> Parallel Seq Scan on
sensor_readings_2024_03 sensor_readings_3 ...
-> Parallel Seq Scan on
sensor_readings_2024_05 sensor_readings_5 ...
-> Parallel Seq Scan on
sensor_readings_2024_07 sensor_readings_7 ...
-> Parallel Seq Scan on
sensor_readings_2024_08 sensor_readings_8 ...
-> Parallel Seq Scan on
sensor_readings_2024_10 sensor_readings_10 ...
-> Parallel Seq Scan on
sensor_readings_2024_12 sensor_readings_12 ...
-> Parallel Seq Scan on
sensor_readings_2024_04 sensor_readings_4 ...
-> Parallel Seq Scan on
sensor_readings_2024_06 sensor_readings_6 ...
-> Parallel Seq Scan on
sensor_readings_2024_09 sensor_readings_9 ...
-> Parallel Seq Scan on
sensor_readings_2024_11 sensor_readings_11 ...
-> Parallel Seq Scan on
sensor_readings_2024_02 sensor_readings_2 ...
-> Parallel Seq Scan on
sensor_readings_default sensor_readings_13 ...

```

With nothing to prune on, every one of the thirteen partitions shows up as its own Parallel Seq Scan under a Parallel Append. This is the direct, visible contrast: pruning isn't a performance optimization happening invisibly under the hood, it's a specific number of Seq Scan lines disappearing from the plan, and you can watch it happen by adding or removing a date filter on the exact same query.

Exercise 4 — BRIN vs. B-tree, With Real Numbers

4.1 — Why this data is a good fit for BRIN

BRIN's usefulness rests entirely on physical/value correlation — do rows with nearby recorded_at values actually sit near each other on disk? Since this data was generated and loaded in chronological order and never updated, it should be close to perfectly correlated:

```
ANALYZE sensor_readings;
```

```

SELECT tablename, attname, correlation, n_distinct
FROM pg_stats
WHERE tablename = 'sensor_readings_2024_09' AND attname =
'recorded_at';

```

tablename	attname	correlation	n_distinct
-----+-----+-----			
+-----			
sensor_readings_2024_09	recorded_at	0.99979776	8622

correlation ranges from -1 to 1; 0.9998 is as close to “physically sorted by this column” as real data gets. This is the number that decides whether BRIN is a good idea here — not the row count, not the table size.

4.2 — Build both indexes on one partition, compare cost to build and size

```

CREATE INDEX idx_sr_2024_09_brin ON sensor_readings_2024_09 USING
BRIN (recorded_at);

```

```

CREATE INDEX idx_sr_2024_09_btree ON sensor_readings_2024_09 USING
BTREE (recorded_at);

```

```

CREATE INDEX
Time: 122.619 ms
CREATE INDEX
Time: 249.923 ms

```

```

SELECT relname, pg_size_pretty(pg_relation_size(oid)) AS
index_size
FROM pg_class
WHERE relname IN ('idx_sr_2024_09_brin', 'idx_sr_2024_09_btree');

```

relname	index_size
-----+-----	
idx_sr_2024_09_brin	24 kB
idx_sr_2024_09_btree	5800 kB

864,000 rows in this partition. The B-tree costs 5.8 MB — one entry per row, roughly what you’d expect. The BRIN index costs 24 KB: **about 240 times smaller**, because it isn’t storing one entry per row, it’s storing one min/max pair per 128-page block range — a few thousand summaries instead of 864,000 index entries. It also built about twice as fast. Neither number changes much if the table has a hundred million rows instead of 864,000; a B-tree’s cost scales with row count, BRIN’s scales with table *pages*, and pages-per-block-range is fixed.

4.3 — What that size difference costs you at query time

Force each index in turn and run the same narrow range query — one day out of the month:

```
SET enable_seqscan = off;
DROP INDEX idx_sr_2024_09_btree; -- isolate the BRIN index

EXPLAIN (ANALYZE, BUFFERS)
SELECT COUNT(*), round(AVG(reading_value)::numeric, 2)
FROM   sensor_readings_2024_09
WHERE  recorded_at >= '2024-09-15' AND recorded_at < '2024-09-16';
```

```
Aggregate  (cost=7818.29..7818.30 rows=1 width=40) (actual
time=7.350..7.351 rows=1 loops=1)
  Buffers: shared hit=9 read=384
    -> Bitmap Heap Scan on sensor_readings_2024_09
        (cost=19.27..7674.10 rows=28836 width=8) (actual time=0.635..5.708
rows=28800 loops=1)
          Recheck Cond: ((recorded_at >= '2024-09-15
00:00:00+00'::timestamp with time zone) AND (recorded_at <
'2024-09-16 00:00:00+00'::timestamp with time zone))
          Rows Removed by Index Recheck: 17280
          Heap Blocks: lossy=384
          Buffers: shared hit=9 read=384
            -> Bitmap Index Scan on idx_sr_2024_09_brin
                (cost=0.00..12.06 rows=30322 width=0) (actual time=0.060..0.060
rows=3840 loops=1)
                  Index Cond: ((recorded_at >= '2024-09-15
00:00:00+00'::timestamp with time zone) AND (recorded_at <
'2024-09-16 00:00:00+00'::timestamp with time zone))
          Execution Time: 7.493 ms
```

Heap Blocks: lossy=384 and Rows Removed by Index Recheck: 17280 are BRIN being honest about what it is: a **lossy** index. It doesn't know which rows in a matching block range satisfy the condition, only that the range's min/max *could* contain a match — so it hands back every block in range and PostgreSQL rechecks the actual condition row by row once the blocks are read. That recheck cost is real, and it's the trade for the 240x smaller index.

Now the B-tree, same query:

```
DROP INDEX idx_sr_2024_09_brin;
CREATE INDEX idx_sr_2024_09_btree ON sensor_readings_2024_09 USING
BTREE (recorded_at);

EXPLAIN (ANALYZE, BUFFERS)
SELECT COUNT(*), round(AVG(reading_value)::numeric, 2)
FROM   sensor_readings_2024_09
WHERE  recorded_at >= '2024-09-15' AND recorded_at < '2024-09-16';
```

```

Aggregate  (cost=1076.88..1076.89 rows=1 width=40) (actual
time=5.781..5.782 rows=1 loops=1)
  Buffers: shared hit=384 read=27
    -> Index Scan using idx_sr_2024_09_btree on
sensor_readings_2024_09  (cost=0.42..932.69 rows=28836 width=8)
(actual time=0.055..3.552 rows=28800 loops=1)
      Index Cond: ((recorded_at >= '2024-09-15
00:00:00+00'::timestamp with time zone) AND (recorded_at <
'2024-09-16 00:00:00+00'::timestamp with time zone))
    Execution Time: 5.850 ms

```

5.85 ms for the B-tree against 7.49 ms for BRIN — the B-tree is genuinely a bit faster here, exact and recheck-free by construction. But look at what each one costs to get that: 5.8 MB of index versus 24 KB, on *one month of one partition*. At the scale this chapter’s title implies — many partitions, much more than a year of data — a B-tree on every partition’s timestamp column stops being “a bit more disk” and starts being a meaningful fraction of the table’s own size, while BRIN stays essentially free. For an append-only, time-correlated column like this one, that trade is rarely close.

4.4 — Keep the BRIN index, drop the comparison B-tree

```

DROP INDEX idx_sr_2024_09_btree;
CREATE INDEX idx_sr_2024_09_brin ON sensor_readings_2024_09 USING
BRIN (recorded_at);
RESET enable_seqscan;

```

Recreate this same USING BRIN (recorded_at) index on every other monthly partition before moving on — Exercise 5 will drop one of them whole, and Exercise 6 revisits the size question at the whole-table level.

Exercise 5 — Dropping a Partition vs. DELETE

5.1 — Time the unpartitioned way: DELETE

sensor_readings_staging is still sitting there, unpartitioned, exactly as it was before this chapter touched anything. Delete a month’s worth of rows from it the only way an unpartitioned table allows:

```

DELETE FROM sensor_readings_staging
WHERE recorded_at >= '2024-01-01' AND recorded_at < '2024-02-01';

DELETE 892800
Time: 815.293 ms

```

5.2 — Time the partitioned way: DROP TABLE


```
DROP TABLE sensor_readings_2024_01;
```

```
DROP TABLE
```

```
Time: 17.489 ms
```

Same amount of data — the January 2024 slice, 892,800 rows either way — and dropping the partition is roughly **47 times faster** than deleting the equivalent rows. That gap isn't a fluke of this particular run: it's two operations of fundamentally different cost. `DELETE` is an MVCC operation — PostgreSQL has to visit every matching row, mark each one dead, and write a WAL record for each one, transaction-safely, one row at a time. `DROP TABLE` on a partition is a catalog operation — it removes one entry saying “this table is part of that partition set” and unlinks the file. No row is individually touched because no row was individually anything; the whole partition was always one physical object.

5.3 — The part `DELETE` doesn't finish

```
SELECT n_live_tup, n_dead_tup,
pg_size_pretty(pg_relation_size('sensor_readings_staging'))
FROM pg_stat_user_tables
WHERE relname = 'sensor_readings_staging';
```

```
 n_live_tup | n_dead_tup | pg_size_pretty
-----+-----
 9648230 |    892800 | 687 MB
```

892,800 dead tuples, exactly matching what was just deleted — and the table is still 687 MB, unchanged, because `DELETE` doesn't reclaim disk space, it just marks rows as no-longer-visible. Getting that space back needs a `VACUUM` (or `VACUUM FULL` to actually shrink the file), which is its own I/O-heavy pass over the table, usually running sometime later via autovacuum rather than as part of the `DELETE` itself. `DROP TABLE` skipped all of this — there is no bloat to vacuum, because there's no table left to have any.

5.4 — A gotcha worth knowing about before it surprises you

```
SELECT pg_size_pretty(pg_total_relation_size('sensor_readings'));

pg_size_pretty
-----
0 bytes
```

That’s not an error and the data is not gone —
pg_total_relation_size() on a *partitioned* table reports only the parent’s own storage, and a partitioned table’s parent has none; all the bytes live in the child partitions. To get the real total, sum across the children:

```
SELECT pg_size_pretty(SUM(pg_total_relation_size(inhrelid))) AS
total_size
FROM pg_inherits
WHERE inhparent = 'sensor_readings'::regclass;

total_size
-----
629 MB
```

5.5 — Clean up the staging table


It’s done its job — everything from here on works from
sensor_readings:

```
DROP TABLE sensor_readings_staging;
```

Exercise 6 — Combining LIST and RANGE: Subpartitioning

6.1 — Why one partition key isn’t always enough

sensor_readings prunes beautifully on date, but a query that also filters on sensor_type — “just the traffic sensors” — still has to scan a whole month’s worth of every sensor type, because sensor_type isn’t part of the partition key at all. PostgreSQL allows a second level: partition by LIST (sensor_type) at the top, then partition *each of those* by RANGE (recorded_at) underneath — a partition of a partition, usually called subpartitioning. Two routing decisions, made independently, before a row ever reaches physical storage:



The root and the three type branches hold no data of their own — they exist purely to route a row one level further down. Only the nine boxes at the bottom are real tables with real storage, and every one of this chapter's rows lives in exactly one of them, decided first by `sensor_type`, then by which month `recorded_at` falls in.

6.2 — Build the two-level hierarchy

```
CREATE TABLE sensor_readings_by_type (
  id          BIGINT GENERATED ALWAYS AS IDENTITY,
  sensor_id   INTEGER NOT NULL,
  sensor_type TEXT NOT NULL
                      CHECK (sensor_type IN ('temperature',
'traffic', 'air_quality')),
  reading_value DOUBLE PRECISION NOT NULL,
  recorded_at  TIMESTAMPTZ NOT NULL
) PARTITION BY LIST (sensor_type);

CREATE TABLE srbt_temperature PARTITION OF sensor_readings_by_type
  FOR VALUES IN ('temperature') PARTITION BY RANGE
(recorded_at);
CREATE TABLE srbt_traffic PARTITION OF sensor_readings_by_type
  FOR VALUES IN ('traffic') PARTITION BY RANGE (recorded_at);
CREATE TABLE srbt_air_quality PARTITION OF sensor_readings_by_type
  FOR VALUES IN ('air_quality') PARTITION BY RANGE
(recorded_at);

-- Three months of subpartitions under each type branch:
CREATE TABLE srbt_temperature_2024_04 PARTITION OF
srbt_temperature FOR VALUES FROM ('2024-04-01') TO ('2024-05-01');
CREATE TABLE srbt_temperature_2024_05 PARTITION OF
srbt_temperature FOR VALUES FROM ('2024-05-01') TO ('2024-06-01');
CREATE TABLE srbt_temperature_2024_06 PARTITION OF
srbt_temperature FOR VALUES FROM ('2024-06-01') TO ('2024-07-01');

CREATE TABLE srbt_traffic_2024_04 PARTITION OF srbt_traffic FOR
VALUES FROM ('2024-04-01') TO ('2024-05-01');
CREATE TABLE srbt_traffic_2024_05 PARTITION OF srbt_traffic FOR
VALUES FROM ('2024-05-01') TO ('2024-06-01');
CREATE TABLE srbt_traffic_2024_06 PARTITION OF srbt_traffic FOR
VALUES FROM ('2024-06-01') TO ('2024-07-01');

CREATE TABLE srbt_air_quality_2024_04 PARTITION OF
srbt_air_quality FOR VALUES FROM ('2024-04-01') TO ('2024-05-01');
CREATE TABLE srbt_air_quality_2024_05 PARTITION OF
srbt_air_quality FOR VALUES FROM ('2024-05-01') TO ('2024-06-01');
CREATE TABLE srbt_air_quality_2024_06 PARTITION OF
srbt_air_quality FOR VALUES FROM ('2024-06-01') TO ('2024-07-01');
```

Three months only, and only as a working example — building out all twelve for all three types is the same pattern repeated, not a new idea, and this chapter’s disk budget would rather spend on data than on empty demonstration partitions.

6.3 — Backfill from sensor_readings and inspect the tree

```
INSERT INTO sensor_readings_by_type (sensor_id, sensor_type,
reading_value, recorded_at)
SELECT sensor_id, sensor_type, reading_value, recorded_at
FROM sensor_readings
WHERE recorded_at >= '2024-04-01' AND recorded_at < '2024-07-01';
```

```
INSERT 0 2620800
Time: 3389.698 ms (00:03.390)
```

```
SELECT relid::regclass AS relation, parentrelid::regclass AS
parent, isleaf, level
FROM pg_partition_tree('sensor_readings_by_type')
ORDER BY level, 1;
```

	relation	parent	isleaf	
level				
	-----+-----+-----			
	+-----			
	sensor_readings_by_type		f	
0				
	srbt_air_quality	sensor_readings_by_type	f	
1				
	srbt_temperature	sensor_readings_by_type	f	
1				
	srbt_traffic	sensor_readings_by_type	f	
1				
	srbt_air_quality_2024_04	srbt_air_quality	t	
2				
	srbt_air_quality_2024_05	srbt_air_quality	t	
2				
	srbt_air_quality_2024_06	srbt_air_quality	t	
2				
	srbt_temperature_2024_04	srbt_temperature	t	
2				
	srbt_temperature_2024_05	srbt_temperature	t	
2				
	srbt_temperature_2024_06	srbt_temperature	t	
2				
	srbt_traffic_2024_04	srbt_traffic	t	
2				
	srbt_traffic_2024_05	srbt_traffic	t	
2				
	srbt_traffic_2024_06	srbt_traffic	t	
2				
(13 rows)				

`pg_partition_tree()` is the built-in way to see this shape directly instead of inferring it from table names: one root, three LIST branches at level 1, nine RANGE leaves at level 2 — every row actually lives in one of those nine leaf tables, never in an intermediate one.

6.4 — Confirm pruning happens on both dimensions at once

```
EXPLAIN (ANALYZE, BUFFERS)
SELECT COUNT(*), round(AVG(reading_value)::numeric, 2)
FROM   sensor_readings_by_type
WHERE  sensor_type = 'traffic'
AND    recorded_at >= '2024-05-01' AND recorded_at < '2024-06-01';
```

```
Finalize Aggregate  (cost=6794.58..6794.60 rows=1 width=40)
(actual time=37.692..40.669 rows=1 loops=1)
  -> Gather  (cost=6794.47..6794.58 rows=1 width=40) (actual
time=37.612..40.655 rows=2 loops=1)
    Workers Planned: 1
    -> Partial Aggregate  (cost=5794.47..5794.48 rows=1
width=40) (actual time=36.350..36.351 rows=1 loops=2)
      -> Parallel Seq Scan on srbt_traffic_2024_05
sensor_readings_by_type  (cost=0.00..5794.45 rows=4 width=8)
(actual time=0.018..26.204 rows=178560 loops=2)
    Filter: ((recorded_at >= '2024-05-01
00:00:00+00':::timestamp with time zone) AND (recorded_at <
'2024-06-01 00:00:00+00':::timestamp with time zone) AND
(sensor_type = 'traffic'::text))
    Execution Time: 40.738 ms
```

One line, Parallel Seq Scan on `srbt_traffic_2024_05`, out of the nine leaf partitions that exist. The `sensor_type = 'traffic'` predicate pruned two of the three top-level branches before the range predicate even got involved, and the range predicate then pruned two of that branch’s three months. Both partition keys did real work on this single query, independently, in the same plan.

Summary — What You Should Now Know

Tool	What it does
<code>PARTITION BY RANGE (col)</code>	Split a table into physical children by value range on one column — the whole table stays queryable as one name
<code>FOR VALUES FROM (...) TO (...)</code>	Half-open interval defining one range partition’s bounds
<code>... DEFAULT partition</code>	

Tool	What it does
	Catches any row that doesn't match a defined range — without one, out-of-range inserts are rejected outright
<code>tableoid::regclass</code>	Ask which physical partition a given row actually lives in
Partition pruning	The planner drops non-matching partitions from the plan <i>before</i> execution, visible as missing Seq Scan lines in EXPLAIN
<code>USING BRIN (col)</code>	Index storing one min/max summary per block range — tiny and cheap, but only useful when physical and value order correlate
<code>correlation</code> in <code>pg_stats</code>	The number that tells you whether BRIN is actually a good idea for a given column
<code>DROP TABLE</code> on a partition	Metadata-only, near-instant removal of a whole partition's worth of data — no per-row MVCC cost, no bloat to vacuum afterward
<code>pg_total_relation_size()</code> on a partitioned table	Reports the (empty) parent only — sum over <code>pg_inherits</code> for the real total
<code>PARTITION BY LIST</code> over <code>PARTITION BY RANGE</code>	Subpartitioning — two partition keys, pruned independently, in one query plan
<code>pg_partition_tree()</code>	Walk a partition hierarchy — root, branches, and leaves — as a query instead of by table-naming convention

The key design insight from this chapter is that partitioning and indexing solve different halves of the same problem, and BRIN sits exactly at their intersection. Partition pruning throws away whole partitions the planner can prove are irrelevant, cutting the search space by orders of magnitude before a single row is touched. BRIN then indexes what's left almost for free, precisely because append-only, time-ordered data — the same shape of data that makes range partitioning by date the obvious choice — is also the shape that makes a physically-correlated, lossy index a good trade instead of a bad one. Neither technique is really “for time-series data” in the abstract; both are for data whose physical order already tells you something true about its values, and this chapter's sensor readings happen to be exactly that.

Going further: hand-writing twelve CREATE TABLE ... PARTITION OF statements a year, forever, is exactly the kind of recurring maintenance that Chapter 19's pg_cron is suited to automate — schedule a monthly job that creates next month's partition before it's needed and, if a retention policy allows it, drops the oldest one the same way Exercise 5 did by hand. In production this whole lifecycle — creation, retention, even the migration of aging partitions to cheaper storage — is usually handled by the pg_partman extension rather than bespoke scripts; everything in this chapter is what pg_partman is automating under the hood, which is worth knowing before reaching for it, not instead of. sensor_readings also isn't done being useful: Chapter 9's materialized views roll it up into daily summaries without re-scanning raw readings each time, Chapter 11's window functions compute rolling averages and day-over-day deltas across it, and Chapter 16 adds a generated reading_date column to make date-only grouping cheaper still. And a caution in the other direction — BRIN is not a free upgrade for every index. A column with low correlation (a randomly- assigned sensor_id, say, rather than an ever-increasing timestamp) gets none of BRIN's benefit and most of a sequential scan's cost, because every block range ends up spanning the entire value domain. Check correlation before reaching for BRIN, the same way Exercise 4 did here — it's a bet on the physical shape of your data, not a strictly smaller B-tree.

